

ABM ENGINE

PRINCIPLES 4 & 5 + ARCHITECTURE

TIMING • REASONING • DECISION ARCHITECTURE • RULE COMPRESSION

Principle 4: 39 Rules | Principle 5: 35 Rules | Architecture: 16 Rules + 25 Meta-Principles

Decimal Technologies | KSA Banking & Financial Institutions

Date: June 11, 2026

Version: 3.0 — Complete Section 1 Excavation

SECTION 1 COMPLETE TOTALS

Principle 1 (Event-Driven Architecture): 92 rules

Principle 2 (Enterprise Decision System): 139 rules

Principle 3 (Distributed Decision Intelligence): 104 rules

Principle 4 (Timing as Competitive Weapon): 39 rules

Principle 5 (Strategic Reasoning as Core Intelligence): 35 rules

Architecture Layer (Compression + Decision Framework): 25 meta-principles + 16 architecture rules

COMPRESSED TOTAL: 25 Meta-Principles + ~230 Operational Rules + 8-Step Decision Framework

PRINCIPLE 4 — TIMING AS PRIMARY COMPETITIVE WEAPON

REWRITTEN PRINCIPLE

Timing is not a constraint on outreach. It is the first-order determinant of outreach success.

Every account exists in a window lifecycle (CLOSED > OPENING > OPEN > CLOSING > CLOSED_AGAIN).

The system detects window transitions before competitors, prepares during OPENING, executes during OPEN.

A perfect message at the wrong time is worthless. An adequate message at the right time is invaluable.

Initial Pass Rules (14 Rules)

NEW RULE [TIMING-PHILOSOPHY-001]

"Window-Based Timing Model" — 5 states per account: WINDOW_CLOSED (no evaluation, outreach wasted), WINDOW_OPENING (early signals, PREPARE — enrich, identify committee, draft messages, but do NOT outreach yet), WINDOW_OPEN (active evaluation, EXECUTE — every day matters, launch outreach), WINDOW_CLOSING (shortlist formed, URGENT ENTRY — if not already in conversation, attempt with high-value differentiator), WINDOW_CLOSED_AGAIN (decision made, STOP — log outcome, begin cooldown). Every signal interpreted for where it places account in window cycle. KEY INSIGHT: prepare during OPENING, execute during OPEN. Most competitors detect OPEN and start preparing then — Decimal should already be in conversation.

NEW RULE [TIMING-PHILOSOPHY-002]

"Window State Inference" — Signal patterns mapped to states: CLOSED = no signals, stable. OPENING = first change signal (new leader, initiative, regulatory). OPEN = multiple corroborating signals (budget confirmed, hiring, vendor meetings). CLOSING = advanced signals (shortlist, procurement). CLOSED_AGAIN = decision signal. System response per state: CLOSED → monitor. OPENING → prepare. OPEN → execute immediately. CLOSING → urgent entry if viable. CLOSED_AGAIN → stop. WHY: The system that prepares during OPENING and executes during OPEN beats the system that starts preparing during OPEN.

NEW RULE [TIMING-FAIL-001]

"Four Timing Failure Modes": TOO_EARLY (outreach during CLOSED — contact burned, "didn't you hear me last time?"). TOO_LATE (outreach during CLOSING — "we've already chosen," wasted effort). RIGHT_TIME_WRONG_SPEED (window OPEN but 5-day pipeline vs 30-day window — competitors arrive first, frame evaluation). RIGHT_TIME_WRONG_INTENSITY (IMMEDIATE urgency but standard 14-day sequence — prospect met 3 vendors by our Touch 2). Each has different root cause: signal interpretation, detection speed, pipeline velocity, urgency calibration.

NEW RULE [TIMING-METRICS-001]

"Timing Quality Metrics": Window Detection Lead Time (days from OPENING signal to first outreach, target ≤7 P1, ≤14 P2). Golden Window Utilization (% of windows with outreach launched inside, target >80%). Window Stage Accuracy (% correct state inference, target >70%). First-Contact Advantage (% where Decimal was first/second vendor, target >50%). Timing-Attributed Failures (% failed outreach where root cause was timing, target <20%). Signal-to-Outreach Velocity (median days P1 to sent, target ≤5 business days).

NEW RULE [TIMING-PORTFOLIO-001]

"Window Urgency Prioritization": CLOSING window > OPEN window > OPENING window. Within same state, by Effective Opportunity. Account with EO 60 whose window CLOSING in 2 weeks MORE URGENT than EO 80 whose window just OPENED for 3 months. Dynamic priority changes daily as windows progress. System re-evaluates whenever any window state changes.

NEW RULE [TIMING-PORTFOLIO-002]

"Window Collision Management": Too many windows open at once: rank by urgency x opportunity. Accounts that can't be served: perform preparation (enrichment, connector activation) even if outreach can't launch. Accounts that WILL miss window: log "WINDOW_MISSED_CAPACITY" — different from detection failure. Feeds capacity planning.

NEW RULE [TIMING-INTEL-001]

"Contact-Level Timing Patterns": Learn over time: email open time per contact, response latency per persona (CTOs 2-3 days, procurement 5-7, C-suite 7-14), day-of-week patterns. After 5+ data points: suggest optimal send times. MVP: use defaults (Tue-Thu 9-11am). Personalize as patterns emerge.

NEW RULE [TIMING-INTEL-002]

"Event-Relative Timing": Leadership change: outreach Day 7-14 (not Day 1 onboarding, not Day 30+ opinions formed). SAMA announcement: Day 1-3 (urgency highest immediately). Funding: Day 14-21 (week 1-2 celebration, weeks 3-4 "what do we build?"). After negative reply: minimum 45-day gap, but P1 regulatory can compress to 21 days. After sequence completion: 45-day cooldown, but OPENING window during cooldown can shorten to 30 days.

NEW RULE [TIMING-INTEL-003]

"Competitor Timing Intelligence": Track when competitors approach banks. Competitors typically 2-3 weeks after SAMA → Decimal approaches within 1 week. Competitors dominate conference-season → Decimal goes pre/post-conference. Competitor annual sales push (Q1 blitz, year-end discount) → Decimal positions BEFORE.

NEW RULE [TIMING-RELATIONSHIP-001]

"Timing as Signal to Prospect": 2 days after announcement = "paying attention." 3 weeks after = "mass-mailing from list." During Ramadan = "don't understand culture." During crisis = "care about selling not about you." 1 day after competitor loss = "ambulance-chasing." 30 days after = "respect the process." Every timing decision must pass: "Would recipient perceive this as thoughtful and relevant, or tone-deaf?"

NEW RULE [TIMING-RELATIONSHIP-002]

"Timing Blackout Events" beyond Ramadan/Eid: Active M&A; (full blackout per ORG-002). Contact personal difficulty (30-day delay for that contact). Account public failure (14-day delay, no reference to failure unless human-approved). National mourning (government-declared, all outreach halts).

NEW RULE [TIMING-FAIL-002]

"Timing Drift Detection": 3 drift patterns: (1) sequences get slower (review SLAs creep, 48h target becomes 5 days — weekly velocity check). (2) Send windows become rigid (always Tue-Thu 9am despite data suggesting Wednesday 2pm — monthly A/B test). (3) Cooldowns become universal (45 days for everyone despite context-based exceptions — audit exceptions applied, zero = logic broken). Monthly timing audit: are we as fast as we claim?

NEW RULE [TIMING-FAIL-003]

"Timing Cannibalization": Poorly timed touches relative to each other: email at 9am + LinkedIn at 9:05 = "automated system." Touch 2 arrives while prospect composing reply to Touch 1. Three contacts get Touch 1 same day despite 48h stagger. Detection: validate actual send times match planned. Flag if two touches <24h apart. Prevention: pre-launch timing simulation maps ALL touches across contacts/channels on timeline. Conflicts adjusted before sending.

NEW RULE [TIMING-PRINCIPLE-REWRITE]

Original: "Low volume, high relevance." Rewritten: "Timing is not a constraint. It is the primary competitive weapon. The vendor who arrives at the right moment in the buying window frames the entire evaluation. The system detects transitions before competitors, prepares during OPENING, executes during OPEN, and never outreaches into closed windows."

Iteration 1: Timing as Information Asymmetry (5 Rules)

NEW RULE [TIMINGINFO-001]

"Timing Advantage Score": Track per outreach — FIRST MOVER (first vendor, frame evaluation), EARLY MOVER (among first 2-3, influence evaluation), SIMULTANEOUS (no advantage), LATE MOVER (catch-up, respond to others' criteria). Track First Mover Rate: target >60% first/early. Ultimate validation of entire timing engine. Consistently late = engine failing at primary mission.

NEW RULE [TIMINGINFO-002]

"Competitive Timing Discovery": 3 methods: (a) Direct — owner asks "speaking with other vendors?" log answer. (b) Indirect — connector intel. (c) Signal-derived — competitive signals. Every meeting includes timing position. Quarterly: first-mover rate. Below 40% = timing engine outperformed — investigate competitor detection speed, dark intent sources, pipeline processing.

NEW RULE [TIMINGINFO-003]

"Proactive Timing Strategies": (a) THOUGHT LEADERSHIP — publish relevant content before prospects have projects. Known entity when project starts. (b) CONFERENCE PRE-POSITIONING — identify target executives at conferences, arrange connector intros. (c) BUDGET CYCLE SEEDING — Q4 ROI analyses to HIGH-FIT LOW-INTENT accounts, get in budget proposals. (d) REGULATORY PRE-POSITIONING — brief targets on upcoming SAMA implications before mandate official. NOT outreach sequences. POSITIONING activities that create future timing advantage.

NEW RULE [TIMINGINFO-004]

"Positioning vs Outreach — Separate Tracks": OUTREACH: signal-triggered, personalized, conversion-focused. All volume/timing/review rules apply. Counted in daily limits. POSITIONING: proactive, topic-focused, awareness-building. NOT counted in limits. Lighter QC (content review, not 7-agent chain). Max 2 positioning touches per account per quarter. Critical: positioning is GENERIC to topic, outreach is PERSONALIZED to account.

NEW RULE [TIMINGINFO-005]

"Timing Advantage Moat": Sustainable edge from: (a) DARK INTENT (human networks competitors can't replicate), (b) PROCESSING SPEED (24h faster than competitors from signal to sent), (c) ACCUMULATED INTELLIGENCE (second outreach smarter than competitor's first), (d) RELATIONSHIP DEPTH (timely + known > timely + stranger). Track: is advantage improving, stable, declining? First-mover rate declining QoQ = moat eroding, escalate to strategic review.

Iteration 2: Temporal Decision Architecture (5 Rules)

NEW RULE [TEMPORAL-DECISION-001]

"Timing Paradox Rule": For timing-sensitive decisions calculate: COST_OF_ACTING_NOW (risk of premature × impact of TOO_EARLY) vs COST_OF_WAITING (risk of missing window × impact of TOO_LATE + daily information decay). Waiting > Acting → ACT NOW, accept uncertainty. Acting > Waiting → WAIT with maximum wait period. P1: cost of waiting very high, default ACT NOW. P2: can wait 48-72h for better intelligence. P3/P4: can wait for full enrichment.

NEW RULE [TEMPORAL-DECISION-002]

"Window State Confidence": HIGH ≥80% (multiple corroborating signals, act on it). MODERATE 50-79% (some signals, act with lighter touch — positioning or exploratory). LOW <50% (ambiguous, do NOT launch full outreach, investigate first). Display confidence alongside state: "Window: OPENING (62%)". Basis: new CTO + Q1 timing. Missing: budget confirmation."

NEW RULE [TEMPORAL-DECISION-003]

"Cross-Account Timing Coordination": Peer group / shared board accounts: (a) sequence strategically — stronger opportunity first, positive signal becomes social proof for second. (b) Domino Account timing takes priority. (c) NEVER same-signal outreach to two connected accounts on same day — stagger 3+ business days. (d) Referral from Account A to Account B: approach B within 48h while endorsement is fresh.

NEW RULE [TEMPORAL-DECISION-004]

"Internal Readiness Gates": External timing (window OPEN) must align with internal readiness across 6 dimensions: enrichment sufficient (ICS ≥40), message quality (QC passed), reviewer available, account owner assigned, competitive intelligence known, relevant collateral exists. Each not-ready has specific response: emergency enrichment, fix message, queue for reviewer, assign owner, flag blind spot, proceed anyway for P1/P2 but prepare collateral in parallel.

NEW RULE [TEMPORAL-DECISION-005]

"Timing Irreversibility Assessment": LOW (LinkedIn connection — ignorable, recoverable): confidence ≥50% sufficient. MEDIUM (email to Director — they'll delete and forget): confidence ≥65%. HIGH (email to C-suite — they REMEMBER):

confidence $\geq 80\%$. EXTREME (connector introduction — damages both relationships if wrong): confidence $\geq 85\%$ AND Level 2+ reviewer. Higher cost of mistake → higher confidence required.

Iteration 3: Timing Failure Cascades (6 Rules)

NEW RULE [TIMINGCASCADE-001]

"5 Cascade Patterns": (1) Detection→Pipeline→Window: 5-day detection lag + 3-day processing = 8 days into 14-day window. (2) Enrichment→Reasoning→Quality: thin data → generic message → weak first impression. (3) Volume→Queue→Staleness: surge queues 10 accounts → queued messages reference 7-10 day old signals. (4) Review→Delay→Context_Shift: 48h review delay + new signal arrives → message sent without new competitive context. (5) Cooldown→Signal→Missed_Window: P1 signal on Day 20 of 45-day cooldown → window closes during remaining cooldown.

NEW RULE [TIMINGCASCADE-002]

"Cascade Prevention": (1) Stage-by-stage tracking, alert before delay compounds. (2) Quality floor: thin enrichment flagged "LOW_CONTEXT_QUALITY," reviewer acknowledges. (3) Auto-regenerate queued messages if signal ages past 7 days. (4) New signal flags pending review message: "NEW_CONTEXT_AVAILABLE." (5) P1 overrides cooldown (45→21 days regulatory, 45→30 days non-regulatory if fundamentally different signal).

NEW RULE [TIMINGCASCADE-003]

"System-Wide Timing Health": Weekly: (a) messages queued >48h? (<5% of active accounts). (b) Open Golden Windows with no outreach? (0 for P1, <3 for P2). (c) Active cascade chains? (Flag any 2+ sequential delays). (d) Signal-to-outreach velocity vs 30-day average? (Trending up = slowing). (e) Messages regenerated due to staleness? (Trend up = queue problem). Any RED → immediate investigation.

NEW RULE [TIMINGCASCADE-004]

"Review Timing Priority Queue": Not FIFO. Ordered by TIMING URGENCY: (1) Golden Window closes within 48h. (2) CLOSING window accounts. (3) P1 signal-triggered. (4) Queue >4h (approaching SLA). (5) Everything else chronological. Reviewer sees urgency: "Golden Window closes in 36h" or "In queue 3h 45m (SLA: 4h)."

NEW RULE [TIMINGCASCADE-005]

"Temporal Content Alignment": QC includes: TENSE CHECK (signal 14+ days old → no "recent announcement" language). SEASON CHECK (October message for Q1 initiative → "as you plan for 2027"). BUYING STAGE CHECK (POC offer to AWARE-stage = premature). COMPETITIVE TIMING CHECK (competitor just won/launched → acknowledge or avoid).

NEW RULE [TIMINGCASCADE-006]

"Meta-Timing Question": Before any outreach: "Will recipient perceive timing as (a) timely and relevant, (b) random and unprompted, (c) late/tone-deaf/awkward?" Only (a) → proceed. (b) → reconsider, find better trigger. (c) → DO NOT SEND even if QC passes and message is well-written. Synthesis of all timing rules into one judgment call. Timing equivalent of the ethics "Respect Test."

Iteration 4: Timing Learning (5 Rules)

NEW RULE [TIMINGLEARN-001]

"Timing Data Collection": 10 data points per outreach: signal detection timestamp, signal event timestamp, detection lag, pipeline stage timing (each stage), Golden Window position (day X of Y), window state at send, window state actual (validated), day/week/calendar context, engagement timing, competitive position.

NEW RULE [TIMINGLEARN-002]

"Timing Pattern Analysis" quarterly: (a) Optimal Golden Window position (hypothesis: Day 3-7 optimal). (b) Best day/time by persona type. (c) Detection lag vs outcome correlation. (d) Pipeline bottleneck trends. (e) Seasonal pattern validation (does Ramadan actually reduce engagement, or assumption?). Findings feed AILEARN-002 pipeline with minimum sample sizes.

NEW RULE [TIMINGLEARN-003]

"Timing Experimentation": A/B rules: one variable at a time, minimum 20 messages per condition, control for message quality, 4-week minimum duration. TESTABLE: send time within windows, sequence gaps, LinkedIn timing relative to email, cooldown length. NOT TESTABLE: Ramadan blackout, weekend exclusion, C-suite review, contact stagger minimum.

NEW RULE [TIMINGLEARN-004]

"Per-Account Timing Profile": After 2+ cycles: contact response patterns, organizational decision speed, budget cycle specifics, communication channel preferences, historical engagement timing. Part of Cumulative Intelligence. General rules remain DEFAULT, account-specific patterns override with 5+ data points.

NEW RULE [TIMINGLEARN-005]

"Timing Effectiveness Assessment" quarterly: (a) Timing attribution (prospect says "perfect timing" in >30% of meetings?). (b) Window prediction accuracy (>70%). (c) First-mover correlation (do first-movers convert higher? If not, test principle). (d) Timing waste (outreach to retrospectively-closed windows <15%). (e) Opportunity cost of timing rules (delayed opportunities that still converted = rules working; lost = too conservative).

Iteration 5: Timing Philosophy (4 Rules + Rewrite)

NEW RULE [TIMINGPHILOSOPHY-001]

"Timing Primacy": 5 reasons timing is first-order in KSA banking: (a) Buying windows are SHORT (3-9 months, once per 3-5 years). (b) Market is SMALL (25 banks, 200 decision-makers — can't burn through prospects). (c) First mover FRAMES evaluation criteria. (d) Trust is TIME-DEPENDENT (6 months visible > stranger with same message). (e) Regulatory deadlines are ABSOLUTE (SAMA doesn't negotiate). All other outreach qualities are second-order multipliers on timing.

NEW RULE [TIMINGPHILOSOPHY-002]

"Timing-Centric Operating Modes": Three simultaneous: EXECUTE (open windows — maximum speed, full chain, immediate outreach), PREPARE (opening windows — enrichment, committee discovery, connector activation, message pre-drafting), POSITION (closed windows — content sharing, thought leadership, relationship building, monitoring). System directs daily attention to where opportunity exists NOW.

NEW RULE [TIMINGPHILOSOPHY-003]

"Market Entry Timing Calibration": ENTRY PHASE (months 0-6): 5-8 accounts, OPEN windows + warm paths only, avoid cold C-suite, build first reference. Conservative timing thresholds. ESTABLISHMENT (7-18): 15-20 accounts, leverage references, begin cold outreach. SCALING (19+): full 50 accounts, multiple references, brand recognized, cold outreach viable. Timing aggressiveness scales with market position.

NEW RULE [TIMINGPHILOSOPHY-004]

"Window Multiplier": CLOSED ×0.1. OPENING ×0.5. OPEN ×1.5. CLOSING ×1.2. CLOSED_AGAIN ×0.0. Added to Effective Opportunity: $\text{Dynamic_Score} \times (\text{ICS}/100) \times \text{Stage} \times \text{Budget} \times \text{Entrenchment} \times \text{Risk} \times \text{Window_Multiplier}$. Account with perfect scores but CLOSED window = near-zero effective opportunity.

PRINCIPLE 5 — STRATEGIC REASONING AS CORE INTELLIGENCE

REWRITTEN PRINCIPLE

Every outreach is the visible surface of deep strategic reasoning. The chain exists not to produce messages but to UNDERSTAND the prospect's world well enough that the message creates recognition: 'This vendor gets us.' Deep reasoning produces concise messages (5:1 ratio). The chain is willing to recommend silence, self-aware about assumptions and unknowns, and evolutionary.

Initial Pass Rules (10 Rules)

NEW RULE [REASONING-PHILOSOPHY-001]

"Template Trap": System must NEVER degrade to template-filling. Symptoms: (a) messages that apply to any bank by changing name, (b) same opening line in 3+ messages/week, (c) Agent 5 same angle for 5+ consecutive accounts, (d) pain inferences are generic role-based not account-specific. Any symptom → AIQUALITY-001 quality drift alert. Value of chain: UNIQUE strategic thinking per account-contact pair.

NEW RULE [REASONING-PHILOSOPHY-002]

"Optimized Chain Architecture": Agents 1 (Signal), 2 (Research), 3 (Persona) run PARALLEL — they depend on raw data, not each other's outputs. Then sequential: 4 (Pain) → 5 (Strategy) → 6 (Message) → 7 (QC). Reduces chain time 30-40%. Dependency-valid: 1+2+3 are independent analyses that converge at Agent 4.

NEW RULE [REASONING-COHERENCE-001]

"Internal Coherence Checks" at 3 junctures: Check 1 (after 1-3, before 4): signal interpretation + research + persona tell consistent story? If contradiction → COHERENCE_CONFLICT flag for Agent 5 to address. Check 2 (after 5, before 6): strategy aligns with persona? If mismatch → Agent 5 justifies or adjusts. Check 3 (after 6, by 7): message reflects strategy, references correct signal, addresses inferred pain, matches persona preferences. Conflicts presented as tensions to resolve, NOT ignored.

NEW RULE [REASONING-MULTICONTACT-001]

"Account-Level Before Contact-Level": ACCOUNT REASONING (once per cycle): overall situation, 2-3 key themes, cross-contact coordination, unified narrative. CONTACT REASONING (per contact, within account framework): adapt narrative for THIS persona, select role-appropriate proof points and CTA. Account strategy is the CONSTRAINT. No individual message diverges from account narrative.

NEW RULE [REASONING-ACCUMULATION-001]

"Cumulative Account Intelligence": Each cycle builds on previous. First cycle: signals + enrichment (thin). Second: adds engagement learnings (which contacts engaged, what resonated). Third: adds human intelligence (meeting notes, stated priorities, political dynamics). System maintains accumulating layer: confirmed pains, rejected pains, confirmed decision-makers, competitor intel, communication preferences. First message is WEAKEST. Each subsequent is SMARTER.

NEW RULE [REASONING-FAIL-001]

"5 Chain Degradation Patterns": SHALLOW (outputs short, generic — prompts stale), ECHO_CHAMBER (agents agree too readily, no nuance — add Devil's Advocate instruction), ANCHORING_BIAS (over-index on first signal — Agent 5 receives ALL signals, considers which is most relevant), RECENCY_BIAS (newest signal dominates — instruction: strategic importance over recency), CONFIRMATION_BIAS (only supporting evidence — QC checks: does message acknowledge complexity?).

NEW RULE [REASONING-FAIL-002]

"Monthly Reasoning Quality Audit": 5 chains scored on: (a) agent outputs substantive or perfunctory? (b) tensions surfaced? (absence suspicious in complex accounts) (c) message derived from reasoning or could be generated without chain? (d) strategy unique to this account? ("swap the name" test) (e) QC raised substantive concerns? Each 1-5. Average <3.0 = degradation alert. <2.0 = fundamental redesign.

NEW RULE [REASONING-CONTEXT-001]

"Context Priority Hierarchy": TIER 1 MUST INCLUDE: triggering signal, account basics, contact info, Effective Opportunity. TIER 2 SHOULD: AI summary, accumulated intelligence, all active signals, relationship paths, competitive landscape. TIER 3 NICE: full enrichment, historical engagement, strategy notes, peer references. TIER 4 ON REQUEST: full archive, complete contact DB, industry reports, Decimal capability docs. Curated hierarchy > data dump.

NEW RULE [REASONING-CONTEXT-002]

"Context Freshness Enforcement": Before chain runs verify: (a) account summary <90 days old (regenerate if not), (b) contact still in same role (update if not), (c) referenced signal not contradicted by newer signal (include both if contradicted), (d) accumulated intelligence flagged with age (6+ month intelligence marked DATED_CONTEXT).

NEW RULE [REASONING-OUTPUT-001]

"Output Quality Standard" per agent: Agent 1: must state what happened, what it means, confidence, what it doesn't tell us (fail: restating signal without interpretation). Agent 2: 2+ specific strategic priorities with evidence, flag knowledge gaps (fail: "large bank" generics). Agent 3: 2+ role-specific motivations + 1 likely objection (fail: job description not psychological profile). Agent 4: 2+ pains with evidence chains connecting signal→context→role (fail: "probably has manual processes"). Agent 5: ONE primary angle with justification + alternatives considered + why rejected (fail: trying to say everything). Agent 6: specific signal reference, specific Decimal capability, persona tone match, low-friction CTA (fail: swap-the-name generic). Agent 7: each QC check individually + highest risk element identified (fail: "all good").

Iteration 1: Reasoning as Competitive Advantage (5 Rules)

NEW RULE [REASONCOMPETE-001]

"Reasoning Depth as Differentiator": SURFACE INSIGHT (competitors match): "SNB announced transformation. Decimal helps." DEEP INSIGHT (competitors cannot match): "SNB's 2/5 onboarding rating + Bajbair's 30% digital finance KPI (at 25%) = onboarding bottleneck. No-code origination for 60-second onboarding, as built for [Bank], addresses his quarterly gap." Deep connects public data + internal KPI + persona motivation + Decimal capability + proof point. Minimum depth test: "Could a competitor with basic tools produce this?"

NEW RULE [REASONCOMPETE-002]

"Five Reasoning Capabilities": (a) PAIN CHAIN: signal → pressure → operational problem → measurable symptom → Decimal solution. (b) PERSONA-CONTEXTUAL: THIS CTO with THEIR background cares about THIS specific aspect. (c) TIMING-STRATEGIC: this signal + Q1 budget + CTO 90-day window = convergence that won't recur for 12+ months. (d) COMPETITIVE-STRATEGIC: bank chose [Competitor] for core banking → needs complementary origination that integrates. (e) ECOSYSTEM: existing [Backbase] front-end + [Competitor] core → Decimal fills the origination middle layer.

NEW RULE [REASONCOMPETE-003]

"Reasoning Quality Scorecard": 5 dimensions rated 1-5: SPECIFICITY (account-specific vs generic), DEPTH (inference steps from signal to action), NOVELTY (would surprise informed observer?), ACTIONABILITY (clear specific strategy?), VERIFIABILITY (traces to evidence?). Average target ≥ 3.5 . Monthly: score 10 random chains, track trend.

NEW RULE [REASONCOMPETE-004]

"Automated Quality Proxies": (a) SPECIFICITY: named entity density, threshold ≥ 5 per chain. (b) DEPTH: causal connectors ("because," "therefore"), threshold ≥ 3 . (c) NOVELTY: overlap with common insight database, threshold $\leq 20\%$. (d) EVIDENCE: source citations, threshold ≥ 3 . 2+ proxies below threshold → REASONING_QUALITY_FLAG for human check.

NEW RULE [REASONCOMPETE-005]

"Reasoning Cost-Benefit Tiers": MAXIMUM (Domino/Tier A/C-suite): full 7-agent, expanded context, highest model, no shortcuts. STANDARD (Tier B, VP/Director, P1/P2): full chain, standard config. EFFICIENT (Tier C, P3/P4): abbreviated — Agents 1+2+3 parallel → Agent 5 (handles pain as part of strategy) → Agent 6. 40% cheaper, 30% faster. TEMPLATE_PLUS (positioning): pre-approved templates with minor personalization, no chain needed.

Iteration 2: Reasoning Failure Recovery (5 Rules)

NEW RULE [REASONRECOVER-001]

"Strategic Error Classification": WRONG_PAIN (AI infers cost reduction, they need speed-to-market — update enrichment, re-run Agent 4 with "reconsider" instruction). WRONG_ANGLE (leads compliance, they care about CX — re-run Agent 5 with all angles ranked). WRONG_PERSONA_FIT (technical deep-dive to business executive — update persona classification). WRONG_COMPETITIVE_POSITION (replacement when should be complement — feed entrenchment data). WRONG_TONE (casual to formal executive — update formality rules). Each error feeds correction to specific agent.

NEW RULE [REASONRECOVER-002]

"Post-Outreach Correction": Prospect contradicts reasoning = GOLDEN DATA. (1) Log correction. (2) Update Cumulative Intelligence. (3) Recalculate derived data. (4) If prospect still engaged: generate CORRECTED outreach using new information. (5) Feed into agent improvement: why was chain wrong? Insufficient data? Flawed logic? Undetectable error?

NEW RULE [REASONRECOVER-003]

"Reasoning-Messaging Gap": Right analysis but message doesn't resonate = TRANSLATION failure between Agent 5 and Agent 6. Symptoms: opens but no reply, "I see what you mean but that's not how we think about it." Analysis: compare strategy intent to message execution. Common failures: strategy nuanced but message oversimplified, right pain in vendor-speak not prospect-speak, right angle buried after generic opening. Fix: re-run Agent 6 with specific guidance.

NEW RULE [REASONRECOVER-004]

"Pre-Send Sanity Checks": (a) CONSISTENCY: message claims match enrichment? "Onboarding takes too long" but 4.5-star app? Flag. (b) PLAUSIBILITY: inferred pain plausible for segment/maturity? Tier 1 high-maturity bank doesn't struggle with "basic digital capabilities." (c) DIFFERENTIATION: could competitor make same claim? If yes, reasoning too

generic. (d) RECIPROCITY: message offers VALUE before asking for TIME?

NEW RULE [REASONRECOVER-005]

"Reasoning Invalidation Triggers": 5 circumstances voiding previous reasoning: (a) contradicting signal, (b) contact role change, (c) competitive landscape shift, (d) budget status change, (e) account restructuring. Any trigger with pending/active outreach → "REASONING_INVALIDATED" flag. Messages regenerated with current context before sending.

Iteration 3: Architecture Optimization (5 Rules)

NEW RULE [REASONARCH-001]

"Chain Weaknesses & Mitigations": SINGLE_PASS → add self-critique step (Agent 6b reviews own output, +30 seconds, catches 20-30% self-fixable issues). NO_FEEDBACK → add clarification protocol (Agent 5 can request more from Agent 2, one round). EQUAL_WEIGHT → each output carries confidence score, Agent 5 weights accordingly. MONOLITHIC → cost-benefit tiers. NO_ALTERNATIVES → Agent 5 produces top 2 strategies for MAXIMUM tier, human selects.

NEW RULE [REASONARCH-002]

"Multi-Option Generation": MAXIMUM tier only: Agent 5 produces TWO strategies (primary + alternative). Agent 6 generates message for EACH. Agent 7 validates both. Human sees both with strategy rationale, tradeoffs, recommended choice. Selects one or requests third with guidance. Double compute cost trivial relative to Domino Account opportunity value.

NEW RULE [REASONARCH-003]

"Chain Early Termination": Valid aborts: (a) Agent 1: intent = CONTEXTUAL confidence <0.5 → "signal insufficient." (b) Agent 2: discovers critical negative (competitor won, crisis, acquisition) → "account not suitable now." (c) Agent 4: no specific addressable pain → "recommend enrich further." (d) Agent 5: can't construct compelling differentiated strategy → "recommend positioning instead of outreach." Early termination is GOOD JUDGMENT not failure. Healthy rate: 10-20%. <5% = too permissive. >30% = data quality issues.

NEW RULE [REASONARCH-004]

"Multi-Signal Reasoning": Before per-contact chain, SIGNAL SYNTHESIS step: review ALL active signals, identify dominant narrative, strongest lead signal, supporting signals, contradictions. Message uses SYNTHESIZED narrative: "With new technology leadership, regulatory deadlines, and fresh capital, your bank is positioned to modernize origination" — far more compelling than "I saw your new CTO."

NEW RULE [REASONARCH-005]

"Sequential Touch Reasoning": Touch 1: full chain, signal-triggered opener. Touch 2: abbreviated re-reasoning — different angle (value-led insight), same narrative, new proof point. Touch 3: social proof / comparable, peer reference. Touch 4: low-friction close, acknowledge they're busy, easiest possible CTA. Each builds on previous, never repeats. New signals between touches incorporated: "Since I reached out, I also noticed [new signal]..." Dynamically reasoned = feels human, not pre-scheduled.

Iteration 4: Reasoning Governance (5 Rules)

NEW RULE [REASONGOV-001]

"Quarterly Reasoning Audit": 20 chains (5 meetings, 5 positive replies, 5 no response, 5 negative). Score per agent: signal interpretation correct? Research accurate? Persona insightful? Pains validated? Strategy right choice? Message faithful translation? QC substantive? Reveals strong/weak agents. Track trend over quarters.

NEW RULE [REASONGOV-002]

"Human-AI Reasoning Alignment": Compare AI conclusions to senior BD expert looking at same data. ALIGNMENT +2, MINOR DIVERGENCE +1, SIGNIFICANT DIVERGENCE -1, OPPOSITE CONCLUSIONS -2. Aggregate 20 chains: -40 to +40. Target >=+20. Below +10 = serious concern. Below 0 = emergency review of all prompts and configurations.

NEW RULE [REASONGOV-003]

"Reasoning Rules Lifecycle": Agent prompts follow AILEARN-003 (version control). Quality thresholds follow META-001 (review frequency, owner). Chain architecture follows META-004 (cost-benefit). New features follow META-005 (Dual Improvement Test). Self-critique, multi-option, multi-signal synthesis must justify cost after 6 months — retire if not improving outcomes.

NEW RULE [REASONGOV-004]

"Edge Case Handling": (a) Fintech startup (not bank): detect SINGLE_DECISION_MAKER, startup-optimized prompts for Agents 3 and 5. (b) Arabic signal: route to Arabic-capable processing. (c) Dark-intent-only (no public signals): skip Agent 1, Agent 5 receives human note as context. (d) Previous negative contact: Agent 5 receives interaction history, must offer something fundamentally different. Each edge case has defined handling path. Unhandled = graceful termination.

NEW RULE [REASONGOV-005]

"Architecture Evolution Path": 4 potential evolutions: (1) merge Agent 4+5 into "Strategic Analysis" if consistently better combined. (2) promote self-critique to full agent if highly valuable. (3) promote signal synthesis to Agent 0 if essential. (4) add intelligence retrieval agent for cross-account patterns. Change only with 6+ months data, A/B tested, human approved.

Iteration 5: Reasoning Philosophy (4 Rules + Rewrite)

NEW RULE [REASONPHILOSOPHY-001]

"Gets Us Standard": Would recipient think "they understand our specific situation" or "well-researched mass email"? Indicators of "gets us": references specific initiative by name, connects capability to measurable pain, uses prospect's language, acknowledges complexity. Indicators of "mass email": could apply to any bank, vendor-speak, generic claims, disconnected CTA. Chain success measured by "Gets Us" rate, not QC pass rates.

NEW RULE [REASONPHILOSOPHY-002]

"Depth-Brevity Paradox": Reasoning should be LONG (500-1000+ words internal). Message should be SHORT (80-150 words). Ratio $\geq 5:1$. "Your 2/5 onboarding rating costs you the 5% digital share gap you need by 2027. We solved this for [Bank] in 60 days." — 25 words requiring 500+ words of reasoning. Below 3:1 = reasoning too shallow, template-filling-with-extra-steps. Target: 7:1 MAXIMUM, 5:1 STANDARD, 3:1 EFFICIENT.

NEW RULE [REASONPHILOSOPHY-003]

"Value of Silence": "Recommend no outreach" is VALID and VALUABLE output. Correct when: pain doesn't exist, window CLOSED, incumbent too entrenched, stability too low, intelligence insufficient. Track as POSITIVE outcome (good judgment) not negative. Silence Rate: healthy 10-25%. $<5\%$ = too permissive. $>35\%$ = criteria too strict or data quality issues.

NEW RULE [REASONPHILOSOPHY-004]

"Reasoning About Unknowns": Each chain must answer: "What COULD be true that would completely change our strategy?" Examples: "If they've already selected a vendor..." "If the CTO is leaving..." "If SAMA deadline gets extended..." Agent 5 includes KEY UNKNOWN section. Feeds NEXT BEST QUESTION, reviewer context, and meeting prep. Best strategic thinking always acknowledges its own blind spots.

ARCHITECTURE: DECISION FRAMEWORK & RULE COMPRESSION

This section transforms 409+ flat rules into a three-layer decision architecture: 25 immutable Meta-Principles, ~230 operational rules in 7 domains, and an 8-step universal decision framework. This is the layer that makes the system a DECISION ENGINE rather than a RULES ENGINE.

Hidden Layer A: Constraint Hierarchy (4 Rules)

NEW RULE [RULEHIER-001]

"Universal Rule Priority Hierarchy" — When rules conflict, higher tier ALWAYS wins: TIER 0 LEGAL/COMPLIANCE (PDPL, SAMA CSF, opt-outs — NEVER overridable). TIER 1 SAFETY/TRUST (kill switch, hallucination prevention, C-suite review — override only by BD leadership). TIER 2 RELATIONSHIP PROTECTION (fatigue, cooldowns, blockers, cultural blackouts — override by owner + documented). TIER 3 TIMING (Golden Windows, urgency sequences, pipeline velocity — override by owner). TIER 4 REVENUE (scoring, tiers, volume caps, portfolio — override by leadership). TIER 5 OPTIMIZATION (A/B tests, prompts, analytics — override by anyone). Example resolution: P1 signal (Tier 3) + Ramadan (Tier 2) + C-suite (Tier 1) + reviewer unavailable (Tier 3) → Tier 1 wins (review mandatory), Tier 2 wins (Ramadan hours), message queues for reviewer during 10am-2pm.

NEW RULE [RULEHIER-002]

"Same-Tier Resolution": Step 1: EXPECTED VALUE — which preserves more opportunity? Step 2: REVERSIBILITY — if close, choose more reversible (delayed outreach can be sent, sent cannot be unsent). Step 3: DEFAULT TO CAUTION — if still equal, choose conservative (cost of one aggressive mistake > cost of one missed opportunity).

NEW RULE [RULEHIER-003]

"Unresolvable Conflicts": NEVER silently choose. Conflict becomes WORKFLOW ITEM with both options, tradeoff analysis, system's best guess, confidence. Account_owner decides. 4h resolution limit for OPEN windows, 24h for others. If not resolved in time → conservative option auto-activates.

NEW RULE [RULEHIER-004]

"Conflict Resolution Library": Every resolution stored as precedent. 3+ similar resolutions with same outcome = STANDING RULE for auto-resolution. Quarterly: recurring conflicts suggest base rules need restructuring to eliminate conflict source.

Hidden Layer B: Scarcity Economics (3 Rules)

NEW RULE [SCARCITY-001]

"Five Capital Types": ATTENTION (human hours — depleted by reviews, investigations; replenished only by time/scaling). RELATIONSHIP (prospect goodwill — depleted by every outreach; replenished by successful meetings, delivered value). CONNECTOR (introduction willingness — depleted by every request; replenished by reciprocal value). BRAND (market reputation — depleted by bad outreach, failures; replenished by references, thought leadership). SYSTEM (AI compute/APIs — depleted by chain runs; replenished by budget). Track all five continuously.

NEW RULE [SCARCITY-002]

"Expected Return on Attention (EROA)": $EROA = \text{Effective_Opportunity} / \text{Total_Capital_Cost}$. High EROA = pursue. Low EROA = reconsider. Negative EROA = do not pursue. Critical application: FORTRESS account + KPMG connector = HIGH capital cost (burns connector on low-probability), LOW return. EROA negative despite moderate opportunity score. Save connector for higher-probability account.

NEW RULE [SCARCITY-003]

"Connector Capital Management": Per-connector quarterly budgets: ITC 3-4 intros, Backbase 2-3, Tarabut 1-2 (high-value, limited), Geidea 2-3, KPMG 1-2 (reserve for Domino), Mastercard 1/quarter max (nuclear option — CEO access). Before requesting: budget available? Highest-value use? Reciprocal value provided recently? Exceeding budget without reciprocation → connector fatigue → introductions declined → permanent channel degradation.

Hidden Layer C: Market Memory (3 Rules)

NEW RULE [MARKETMEM-001]

"Market Reputation Impact": Every outreach is a reputation event in a 25-bank market. Before C-suite/VP outreach: "If this goes badly, maximum reputational blast radius?" Contacts with board memberships, advisory roles, conference speaking = LARGE blast radius → MAXIMUM reasoning, Level 2+ reviewer, relationship path considered. Never cold outreach to high-blast-radius contact without warm path or exceptional signal confidence.

NEW RULE [MARKETMEM-002]

"Cross-Career Tracking": System checks: "Have we interacted with this PERSON before at any account?" If YES + previous NEGATIVE → flag "PREVIOUS_NEGATIVE_HISTORY at [Old Account]. Strategy must account for prior impression." If YES + previous POSITIVE → flag "PRIOR_ENGAGEMENT. Leverage: 'Great conversation when you were at [Old Account]...' People carry memories across jobs.

NEW RULE [MARKETMEM-003]

"Brand Capital Depletion Alert": 5 signals: declining LinkedIn accept rate QoQ, declining email open rate across accounts, connector reports "mixed things about Decimal," prospect tone negative despite awareness, multiple rejections with similar objections (narrative spreading). 2+ signals → BRAND_CAPITAL_WARNING: pause expansion, investigate narrative, address root cause, invest in repair.

Hidden Layer D: Strategic Patience (3 Rules)

NEW RULE [PATIENCE-001]

"Strategic Readiness Score" (0-100): Base 100. Deductions: no relevant KSA reference (-20), missing critical collateral (-15), no qualified owner with capacity (-25), cold only + brand unknown (-20), product feature gap in critical area (-30), competitive blind spot (-10). ≥70 READY. 40-69 PARTIALLY (engage only if EO >60 AND window OPEN). <40 NOT READY — do NOT engage, invest in readiness. Flag "STRATEGIC_HOLD."

NEW RULE [PATIENCE-002]

"Patience Test": Before engaging any Tier A account: "Are we presenting our BEST possible face, or reaching out because signal is hot but WE aren't ready?" Tier 1 bank seeing Decimal's best = potential partnership. Seeing unprepared version = permanent "not enterprise-ready" classification. Cost of engaging Domino Account before ready EXCEEDS cost of letting one window pass. Another window in 2-3 years. No second first impression.

NEW RULE [PATIENCE-003]

"Readiness Investment Roadmap": For accounts where readiness <70 + high fit: identify specific gap, create plan with timeline ("KSA reference 12mo, SAMA package 3mo, local SI partnership 6mo — readiness 35→80"), track monthly, engage when crosses 70. Strategic patience isn't passive. It's active preparation for future opportunities.

THE 25 META-PRINCIPLES

These are the IMMUTABLE strategic truths that every operational rule must trace back to. If a rule can't be traced to a meta-principle, it doesn't belong in the system.

Intelligence Domain (5 Principles)

- **MP-1:** Intelligence reliability must be measured independently from intelligence content.
- **MP-2:** Intelligence from multiple independent sources is exponentially more reliable than single-source.
- **MP-3:** All intelligence decays over time and must be actively refreshed.
- **MP-4:** The system must detect and react to its own intelligence failures.
- **MP-5:** Signals that cannot be independently verified must be treated with proportional skepticism.

Decision Domain (5 Principles)

- **MP-6:** Enterprise buying is a consensus-building process, not an individual decision.
- **MP-7:** Every opportunity exists in a multi-dimensional space: Fit x Intent x Confidence x Timing x Budget x Maturity x Entrenchment x Risk. No single dimension justifies action.
- **MP-8:** The direction of change predicts outcomes better than current state.
- **MP-9:** The system must model what COULD kill an opportunity, not just what makes it attractive.
- **MP-10:** Every recommendation must include: what, why, alternative, tradeoffs, uncertainty, and what information would improve confidence.

Timing Domain (4 Principles)

- **MP-11:** Timing is the first-order variable. All other qualities are second-order multipliers.
- **MP-12:** Every account exists in a window lifecycle. The system must detect transitions and act accordingly.
- **MP-13:** The cost of a timing mistake is proportional to the irreversibility of the action.
- **MP-14:** Timing communicates respect. The when says as much as the what.

Reasoning Domain (4 Principles)

- **MP-15:** Every outreach must be the product of structured reasoning, not template-filling. Output must be unique to the account-contact pair.
- **MP-16:** Deep reasoning produces concise messages. Reasoning-to-message ratio should be at least 5:1.
- **MP-17:** Recommending silence is as valuable as recommending action.
- **MP-18:** Reasoning must be self-aware: explicit about assumptions, evidence strength, and unknowns.

Human-AI Domain (4 Principles)

- **MP-19:** AI operates in the data domain. Humans operate in the relationship domain. The relationship domain is permanently human.
- **MP-20:** Trust is earned incrementally and lost instantly. Measured by consequence-weighted reliability, not error rate.
- **MP-21:** Every AI output must be explainable, every human decision auditable, accountability traceable to a specific person.
- **MP-22:** Human judgment must be preserved through active capability maintenance, not assumed through exposure.

System Domain (3 Principles)

- **MP-23:** When rules conflict: Legal > Safety > Relationship > Timing > Revenue > Optimization.
- **MP-24:** Every action consumes scarce capital. Actions must justify consumption against expected return.
- **MP-25:** The market has memory. Every interaction creates reputation that persists across time, accounts, and careers.

THE 8-STEP UNIVERSAL DECISION SEQUENCE

Every decision in the system — outreach, scoring, enrichment, prioritization — flows through this sequence. The ~230 operational rules define criteria at each step. The 25 meta-principles define WHY each step exists.

STEP 1 — LEGAL/COMPLIANCE CHECK (Tier 0)

Does this action violate any legal or compliance rule?

If YES → STOP. No exceptions. Not even CEO override.

STEP 2 — SAFETY/TRUST CHECK (Tier 1)

Does this action require human authorization not yet given?

Does it risk trust capital beyond acceptable thresholds?

If blocked → queue for human.

STEP 3 — RELATIONSHIP CHECK (Tier 2)

Would this action damage a relationship? Timing blackout? Fatigue limit? Blocker?

Market memory risk? Cultural sensitivity?

If YES → HOLD.

STEP 4 — READINESS CHECK

Is Decimal ready to engage? (Strategic Readiness Score ≥ 40 ?)

Is intelligence sufficient? (ICS check)

Is reasoning quality adequate?

If NOT READY → PREPARE instead of ACT.

STEP 5 — TIMING CHECK (Tier 3)

Is the window open? Is timing position advantageous?

Would waiting improve position?

CLOSED → MONITOR. OPENING → PREPARE. OPEN → proceed to next step.

STEP 6 — CAPITAL CHECK

Sufficient attention, relationship, connector, and brand capital?

Is EROA positive?

If insufficient → queue behind higher-EROA actions.

STEP 7 — EXECUTE

All gates passed. Execute with appropriate reasoning depth, review level, and monitoring.

STEP 8 — LEARN

Capture outcome data: timing quality, reasoning quality, capital consumed, relationship impact, market impact. Feed into learning loops.

THE MASTER EQUATIONS

EFFECTIVE OPPORTUNITY

$\text{Dynamic_Score} \times (\text{ICS} / 100) \times \text{Stage_Bonus} \times \text{Budget_Modifier} \times \text{Entrenchment_Modifier} \times \text{Risk_Modifier} \times \text{Window_Multiplier}$

DECISION SCORE (drives all action)

$\text{DECISION_SCORE} = \text{EFFECTIVE_OPPORTUNITY} \times (\text{Strategic_Readiness} / 100) / \text{Capital_Cost}$

This single number answers: Should we act on this opportunity RIGHT NOW

given everything we know, everything we don't know, and everything it would cost us?

SECTION 1 — COMPLETE ARCHITECTURE SUMMARY

Layer	Component	Count
META-PRINCIPLES	Immutable strategic truths across 5 domains	25
OPERATIONAL RULES	Actionable rules in 7 domains (Signal, Scoring, Account, Timing, Reasoning, Human-AI, Meta-Gov)	~230
DECISION FRAMEWORK	8-step universal decision sequence	1
MASTER EQUATIONS	Effective Opportunity + Decision Score	2
CONFLICT HIERARCHY	6-tier priority system	1
CAPITAL MODEL	5 scarce capital types + EROA	1

This is no longer an ABM engine specification. This is an Enterprise Decision Intelligence System architecture. The transformation: from 409 flat rules to a three-layer decision architecture with immutable principles, operational rules organized by domain, a conflict hierarchy, a scarcity economics model, and a universal decision sequence.

The Five Rewritten Principles

- **P1:** Event-driven intelligence with self-monitoring governance. Three capacity modes. Intelligence Confidence Score.
- **P2:** Enterprise Decision System Architecture. Consensus-building across political stakeholders. 10-dimension trajectory tracking.
- **P3:** Distributed Decision Intelligence. AI reduces uncertainty at scale. Humans apply contextual judgment. Trust earned incrementally, lost instantly.
- **P4:** Timing is the primary competitive weapon. Window lifecycle model. First-mover advantage through preparation.
- **P5:** Strategic reasoning as core intelligence. Deep understanding enables confident brevity. Willing to recommend silence. Self-aware about unknowns.

END OF SECTION 1 — ALL FIVE PRINCIPLES + ARCHITECTURE COMPLETE

Generated: June 11, 2026 at 17:04 UTC