

# ABM BUSINESS LOGIC BIBLE

## §5 v3 (ARCHITECTURAL) — RESOURCE ALLOCATION OPERATING SYSTEM

13 Iterations · 8 Architectural Layers · 27 New Rules · Phase 1 / Phase 2 Flags · OPEN-Q Resolution

Original §5 framing: *Tier Assignment & Routing*. After 13 iterations of excavation, that framing is no longer adequate. §5 is the engine's **Resource Allocation Operating System**. It determines how scarce organizational capital — human expertise, relationship capital, enrichment budget, founder attention, reviewer capacity, execution bandwidth — is deployed across a portfolio of uncertain opportunities. **Tiers are not labels. They are investment decisions.** Routing is not work distribution. It is strategic matching of opportunities to the humans most capable of converting them. This v3 specifies the eight layers that make that reframe operational, marks each rule Phase 1 (MVP-eligible) or Phase 2 (design-now-activate-later), and resolves the two OPEN-DESIGN-QUESTIONS surfaced in §5 v2.

|                    |                          |                         |                         |
|--------------------|--------------------------|-------------------------|-------------------------|
| Iterations         | 13 (5+5+3)               | Architectural layers    | 8                       |
| New rules          | 27                       | Phase 1 / Phase 2 split | ≈11 / 16                |
| OPEN-Q resolutions | Q4 resolved, Q13 partial | Reframe                 | Routing → Allocation OS |

### THE BEDROCK QUESTION — WHY DO TIERS EXIST?

The surface answer is *prioritization*. The deeper answer is that a tier is not a property of an account; it is a property of **Decimal's willingness to invest**. HOT means more enrichment budget, more SDR attention, more founder visibility, more AI compute, more review effort. WARM means less. COLD means minimal. §5 v3 makes this explicit: every rule below either commits resources, preserves them, allocates them across competition, or decides where they should *not* be spent.

## THE EIGHT LAYERS — §5 v3 ARCHITECTURE MAP

Each layer is a distinct excavation. Together they reframe §5 from a routing module into an allocation operating system. Phase markers indicate which rules are MVP-eligible (P1, green) and which require infrastructure that doesn't yet exist (P2, amber).

| # | Layer                         | What it specifies                                                                                                                                               | Rules | Phase mix                       |
|---|-------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------|-------|---------------------------------|
| 1 | Philosophy (Phil)             | Tier as resource commitment; Opportunity Tier vs Relationship Tier; portfolio-relative tiers; tier $\perp$ timing; routing maximizes Account $\times$ Human Fit | 5     | Mixed                           |
| 2 | Portfolio Capacity (Port)     | Finite HOT capacity; admission $\neq$ eligibility; opportunity cost tracked per promotion                                                                       | 4     | Mostly P1                       |
| 3 | Relationship Capital (Rel)    | Relationship Equity Score independent of opportunity; high-equity accounts can't fall below RELATIONSHIP_MONITOR                                                | 2     | P2                              |
| 4 | Expertise Routing (Route)     | Human Expertise Graph; assignment = Opportunity $\times$ Human Fit; round-robin demoted to tie-breaker                                                          | 3     | P2 (Phase 1 fallback specified) |
| 5 | Strategic Escalation (Esc)    | 5-level ladder (SDR $\rightarrow$ AE $\rightarrow$ Domain Expert $\rightarrow$ Founder $\rightarrow$ Board/Advisor); founder time as capital                    | 3     | Mostly P1                       |
| 6 | Attention Economy (Attn)      | Per-account Attention Budget in measured units; every action consumes budget                                                                                    | 3     | P2                              |
| 7 | Human Cognitive Load (Human)  | Cognitive Load Score per account; capacity in Complexity Units, not account count                                                                               | 3     | P2                              |
| 8 | Strategic Patience (Patience) | Action Readiness Score $\perp$ Opportunity Score; HOT_EXECUTE / HOT_MONITOR / HOT_WAIT / HOT_ESCALATE substates                                                 | 3     | Mixed                           |

### PHASE 1 vs PHASE 2 — BUILD DISCIPLINE

**Phase 1 (P1):** rule ships in the MVP. Doesn't require new data infrastructure beyond what §1–§4 already specify. Example: HOT capacity is just a configurable integer.

**Phase 2 (P2):** rule is specified now, activated when its prerequisite data exists. Example: Cognitive Load Score requires a complexity model trained on actual account data; Human Expertise Graph requires expertise tagging across SDRs. Engineering implements the Phase 1 safe default behind a feature flag, with a clean upgrade path to the Phase 2 rule.

The Phase 1 / Phase 2 discipline matches the maturity-flag pattern already used for modules M0–M19. It lets the spec be architecturally complete without forcing day-one implementation of every layer.

## LAYER 1 — PHILOSOPHY (PHIL)

The bedrock reframe. Tiers are resource commitments, not classifications. Opportunity and relationship are independent dimensions. Tiers compete in a portfolio. Timing is separate from tier. Routing is matching, not distribution.

| Rule ID       | Phase | Specification                                                                                                                                                                                                                                                                                                 |
|---------------|-------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| TIER-PHIL-001 | P1    | A tier represents Decimal's <b>Resource Commitment Level</b>, not merely account attractiveness. HOT commits enrichment budget, SDR attention, founder visibility, AI compute, and review effort. WARM commits less. COLD commits minimal. Every tier-driven action is a capital deployment.                  |
| TIER-PHIL-002 | P2    | Separate <b>Opportunity Tier</b> from <b>Relationship Tier</b>. An account may be Opportunity=WARM and Relationship=STRATEGIC simultaneously; the two scores move independently. Tier demotion on opportunity decay does not erase relationship investment (see Layer 3).                                     |
| TIER-PHIL-003 | P1    | Every tier exists within <b>portfolio capacity</b>. Promotion eligibility (score $\geq 78$ ) is not promotion guarantee. HOT admission is bounded by a configurable HOT capacity (Layer 2); accounts exceeding capacity enter HOT_WAITLIST. Tiers are portfolio-relative, not absolute.                       |
| TIER-PHIL-004 | P1    | Tier and Timing are <b>separate dimensions</b>. HOT + WAIT must be expressible. A high-opportunity account may correctly require deliberate non-action (e.g. new CTO 60-day onboarding); the engine must support tier without forcing immediate execution (see Layer 8).                                      |
| TIER-PHIL-005 | P2    | Routing objective is <b>maximize (Account Potential $\times$ Human Fit)</b>, not round-robin fairness. Humans are not interchangeable resources; the engine treats them as strategic assets whose expertise materially changes conversion probability (see Layer 4). <b>Resolves OPEN-Q4 architecturally.</b> |

## LAYER 2 — PORTFOLIO CAPACITY (PORT)

Universities admit a fraction of qualified applicants. §5 v3 applies the same discipline: qualifying for HOT (score  $\geq 78$ ) is not the same as receiving HOT treatment. The engine must know what its actual capacity is and decide which qualifying accounts get the scarce slots. Every commitment records its opportunity cost.

| Rule ID       | Phase | Specification                                                                                                                                                                                                                                                                                                          |
|---------------|-------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| TIER-PORT-001 | P1    | The engine maintains a configurable <b>HOT capacity</b> derived from SDR count, reviewer count, enrichment budget, and founder availability. Accounts qualifying for HOT that exceed capacity enter <b>HOT_WAITLIST</b> rather than HOT immediately. WAITLIST is a real state, not a queue label.                      |
| TIER-PORT-002 | P1    | HOT admission ranking is <b>Expected Opportunity Value ÷ Expected Resource Cost</b> , not score alone. An 82-score account with low resource cost can outrank a 91-score account with extreme cost. The waitlist re-ranks every recalc; admission is highest-ratio first.                                              |
| TIER-PORT-003 | P1    | Every HOT promotion records its <b>Displaced Opportunity</b> — which waitlisted account would have received resources if this one had not. Displacement is a first-class field for portfolio learning, not just a debug log.                                                                                           |
| TIER-PORT-004 | P2    | The engine tracks <b>Opportunity Cost Score</b> per HOT account = Expected Value of the best alternative opportunity it displaced. Failed accounts are learned-from at the portfolio level: not 'account failed' but 'account failed AND displaced a better opportunity'. This requires a multi-period outcome record. |

### PARTIAL RESOLUTION OF OPEN-Q13

OPEN-Q13 (oscillation prevention in score→tier→enrichment→signal loop) is **partially resolved** by Layer 2. A capacity-gated HOT can't oscillate freely because a promotion must displace an existing HOT account; displacement is a heavyweight decision the engine cannot make casually. The 24h tier-change cooldown remains the Phase 1 safe default for cases where capacity is not yet a binding constraint.

## LAYER 3 — RELATIONSHIP CAPITAL PRESERVATION (REL)

Six months of engagement, multiple meetings, an executive intro, conference relationship — all of it sunk into an account whose composite score has now fallen. The original §5 demotes the account and releases the owner. That is operationally correct on the opportunity axis and strategically wrong on the relationship axis. Layer 3 separates them.

| Rule ID      | Phase | Specification                                                                                                                                                                                                                                                                                                                                                                                                                           |
|--------------|-------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| TIER-REL-001 | P2    | Maintain a <b>Relationship Equity Score</b> per account, independent of the composite score, computed from: meetings held, introductions received, referrals made, executive familiarity, trust history (positive CRM notes), and the durability of warm-path connectors. Equity accumulates; opportunity decays. They move independently.                                                                                              |
| TIER-REL-002 | P2    | Accounts with Relationship Equity above threshold cannot fall below <b>RELATIONSHIP_MONITOR</b> tier regardless of opportunity decay. RELATIONSHIP_MONITOR is a fifth-state alongside HOT / WARM / COLD / WAITLIST: opportunity engagement is paused, but relationship maintenance (occasional value sends, executive check-ins, conference outreach) continues. The relationship asset is preserved against opportunity-axis demotion. |

### WHY THIS IS PHASE 2 (DESIGN-NOW-ACTIVATE-LATER)

Relationship Equity requires data the engine does not yet possess at MVP: structured meeting outcomes, introduction graphs, trust-history scoring. Layer 3 is specified now so that when those streams come online (via INT-SIG SIG-PATH and INT-CRM enrichment), the rules are ready to bind. Phase 1 engineering may store the raw signals (meeting count, intro events) without computing the score, and activate the rule when the model has enough data to be calibrated.

## LAYER 4 — EXPERTISE ROUTING (ROUTE)

OPEN-Q4 surfaced this layer in §5 v2. The bedrock answer: humans are not interchangeable. Assigning a generic SDR to a KSA Open Banking opportunity vs assigning a former banking consultant produces different conversion probability. Routing is not work distribution; it is expertise matching.

| Rule ID        | Phase | Specification                                                                                                                                                                                                                                                                                                                                                        |
|----------------|-------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| TIER-ROUTE-001 | P2    | The engine maintains a <b>Human Expertise Graph</b> across dimensions: geography (KSA / GCC / global), industry (Islamic banking, digital banking, BNPL, embedded finance), product (lending, origination, open banking, KYC), relationship network (who the human knows at target accounts), and language (EN, AR, fluency level).                                  |
| TIER-ROUTE-002 | P2    | <b>Owner Assignment Score = Opportunity Score × Human Fit Score.</b> The Human Fit Score is derived from the Expertise Graph against the account's profile. Highest combined score wins assignment. <b>Resolves OPEN-Q4.</b>                                                                                                                                         |
| TIER-ROUTE-003 | P1    | Round-robin is a <b>tie-breaker only</b>, never the primary assignment logic. When two humans score equally on Opportunity × Human Fit (or when the Human Fit Score is not yet computable in Phase 1), round-robin selects between them. The Phase 1 safe default (pure round-robin everywhere) is upgraded to TIER-ROUTE-002 when the Expertise Graph is populated. |

### RESOLUTION OF OPEN-Q4

The architectural answer to OPEN-Q4 is **expertise-aware routing** (TIER-ROUTE-001/002). The Phase 1 engineering default remains pure round-robin (TIER-ROUTE-003 tie-breaker behavior, applied as the primary algorithm until the Expertise Graph exists). When the Graph is populated, the feature flag flips and Opportunity × Human Fit takes over. OPEN-Q4 is no longer open at the architectural level; it remains a Phase 2 build item.

## LAYER 5 — STRATEGIC ESCALATION (ESC)

Founder time, board introductions, executive outreach — these are scarce strategic capital, not workflow steps. The engine must know when an account is worth that capital and when it is not. The original \$5 implied escalation through tier promotion; v3 makes the escalation ladder explicit and gates each level by cost-vs-incremental-value.

| Rule ID      | Phase | Specification                                                                                                                                                                                                                                                                                                                                                                                                                                                               |
|--------------|-------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| TIER-ESC-001 | P1    | The engine maintains an <b>Escalation Eligibility Score</b> per account derived from: account value (deal-size potential), strategic importance (flagship reference value), relationship path quality (P1–P5 from SIG-PATH), urgency (timing window from M5), and target executive level. Score determines which escalation tiers are available.                                                                                                                            |
| TIER-ESC-002 | P1    | Founder intervention requires Expected Incremental Value > Founder Attention Cost. If the same outcome is achievable with an AE or Domain Expert, the founder is not invoked. Founder time is allocated like venture capital, not like SDR rotation.                                                                                                                                                                                                                        |
| TIER-ESC-003 | P1    | Escalation ladder is explicit and five-level: <b>L1 SDR</b> (standard outreach), <b>L2 Account Executive</b> (qualified discovery), <b>L3 Domain Expert</b> (technical / Islamic-banking / open-banking depth), <b>L4 Founder</b> (strategic accounts, reference banks, high-value handoffs), <b>L5 Board / Advisor</b> (board-level intros, governance-grade relationships). Each level consumes increasing strategic capital and requires the prior level's gate to pass. |

## LAYER 6 — ATTENTION ECONOMY (ATTN)

The deepest reframe: §5 is not allocating accounts, outreach, or SDRs. It is allocating **attention**. Everything below is a measurement of attention in concrete units — SDR-hours, reviewer-minutes, enrichment credits, executive-attention units. Attention is the actual scarce resource the engine governs.

| Rule ID       | Phase | Specification                                                                                                                                                                                                                                                                                                                    |
|---------------|-------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| TIER-ATTN-001 | P2    | Every account receives an <b>Attention Budget</b> measured in concrete units: SDR-hours, reviewer-minutes, enrichment credits, and executive-attention units. The budget scales with tier (HOT receives more, COLD receives minimal) and is bounded by team-level totals.                                                        |
| TIER-ATTN-002 | P2    | Every action consumes attention budget at a defined rate. Examples: manual dossier review = 30 attention units; founder outreach = 100 units; enrichment cycle = 20 units; AI-assisted message review = 5 units. Spend is logged per action; budget overruns require an explicit reallocation.                                   |
| TIER-ATTN-003 | P1    | The stated purpose of tiering is <b>allocation of scarce attention under uncertainty</b> , not account classification. Every §5 rule must be interpretable as an attention-allocation rule; rules that fail this test are not §5 rules. This is the architectural test, not a Phase 2 implementation — it binds spec design now. |

### ATTENTION AS THE UNIT OF ACCOUNT

Layer 6 is the unifying frame underneath every other layer. Layer 2's portfolio capacity is an attention constraint; Layer 3's relationship preservation is attention reserved against future opportunity; Layer 4's expertise routing is attention matched to its highest-value use; Layer 5's escalation is attention deployed up the seniority ladder; Layer 7's cognitive load is attention measured per human; Layer 8's patience is attention withheld when no current deployment improves outcomes. §5 v3 is, in the end, an attention-allocation system.

## LAYER 7 — HUMAN COGNITIVE LOAD (HUMAN)

Humans do not fail because of account count; they fail because of cognitive load. Ten routine accounts and ten politically complex enterprise accounts are not equivalent workloads. Current routing treats them as equal. Layer 7 measures capacity in Complexity Units instead.

| Rule ID        | Phase | Specification                                                                                                                                                                                                                                                                                                                     |
|----------------|-------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| TIER-HUMAN-001 | P2    | Every account receives a <b>Cognitive Load Score</b> derived from: stakeholder count (size of buying committee), buying-center count (decentralized vs centralized), political complexity (SIG-TENSION from §3), signal volatility (rate of new signals), relationship-path complexity (P1–P5 mix), and escalation level (L1–L5). |
| TIER-HUMAN-002 | P2    | Owner capacity is measured in <b>Complexity Units</b> , not account count. A typical SDR capacity is ~100 Complexity Units; a simple account may consume 5 units, a complex KSA bank account may consume 40. Assignment respects unit-budget, not account-count caps.                                                             |
| TIER-HUMAN-003 | P2    | The routing engine optimizes <b>load balance by complexity</b> , not by account volume. Two SDRs with equal account counts can have very different effective workloads; the engine prefers the better-balanced complexity allocation. Burnout prevention is a first-class consideration.                                          |

## LAYER 8 — STRATEGIC PATIENCE (PATIENCE)

Most ABM systems optimize action. Elite systems optimize *omission*. The highest-value action on a new-CTO signal may be to wait 60 days through onboarding rather than reach out on Day 1. The original §5 implies this through M5 timing; Layer 8 makes deliberate non-action a first-class engine output.

| Rule ID           | Phase | Specification                                                                                                                                                                                                                                                                                                                                             |
|-------------------|-------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| TIER-PATIENCE-001 | P1    | Every HOT account carries an <b>Action Readiness Score</b> separate from its Opportunity Score. Opportunity = 90 with Readiness = 25 is not a contradiction; it indicates an immature window. Readiness is derived from M5 Timing inputs (event recency, onboarding curves, regulatory deadline distance) and EPIS coverage.                              |
| TIER-PATIENCE-002 | P1    | HOT has four substates: <b>HOT_EXECUTE</b> (act now), <b>HOT_MONITOR</b> (watch closely, no outreach this cycle), <b>HOT_WAIT</b> (deliberate non-action, scheduled re-evaluation), <b>HOT_ESCALATE</b> (route up the escalation ladder before acting). Same tier, different behavior. The substate is computed from Opportunity × Readiness.             |
| TIER-PATIENCE-003 | P2    | The routing engine may intentionally allocate <b>Monitoring Capital</b> instead of <b>Engagement Capital</b> . Monitoring is a valid attention deployment, not a workflow stall. The engine asks not 'should we engage?' but 'is engagement currently the highest-value use of attention?' — and may answer no while still committing the account to HOT. |

## EXHAUSTION ASSESSMENT & FRONTIER LAYERS

§5 v3 covers thirteen iterations across eight layers. The honest assessment of what is now in spec vs what remains as a frontier for a possible v4:

| Concern                         | §5 v1 | §5 v2 | §5 v3 | Frontier (v4 candidate)                                          |
|---------------------------------|-------|-------|-------|------------------------------------------------------------------|
| Threshold logic                 | ~85%  | ~95%  | ~95%  | — (effectively exhausted)                                        |
| Hysteresis logic                | ~85%  | ~95%  | ~95%  | — (effectively exhausted)                                        |
| Race / ordering / state machine | ~50%  | ~90%  | ~90%  | —                                                                |
| Routing discipline              | ~50%  | ~60%  | ~80%  | Multi-owner collaboration; team-level resource markets           |
| Portfolio theory                | ~10%  | ~25%  | ~75%  | Strategic account selection; portfolio re-balancing              |
| Relationship capital            | —     | —     | ~65%  | Relationship-capital portfolio management; degradation models    |
| Strategic escalation            | ~20%  | ~30%  | ~60%  | Multi-level co-pilot dynamics; founder-led-vs-led-by dynamics    |
| Attention economics             | —     | —     | ~50%  | Cross-team attention markets; pricing models; auction dynamics   |
| Cognitive load / human          | —     | —     | ~50%  | Burnout indicators; recovery cycles; expertise growth            |
| Strategic patience              | —     | —     | ~70%  | Patience-as-strategy vs patience-as-default; window decay models |

### WHAT v3 DELIBERATELY DOES NOT COVER

Three frontier layers remain genuinely unexplored: (1) multi-owner collaboration models (when does an account have two owners, and how is attention shared?), (2) team-level resource markets (do SDRs bid for accounts in an internal market?), and (3) relationship-capital portfolio management (how does Decimal optimize its total relationship portfolio across the KSA banking ecosystem, not just per account?). These would constitute a v4 iteration and are worth pursuing only if the resource-allocation framing of v3 holds up under engineering review. Until then, v3 is the recommended landing point.

## §5 v3 RULE ROLL-UP & CROSS-LINKS

| Layer      | Rules                  | Phase                 | Key cross-links                                                                           |
|------------|------------------------|-----------------------|-------------------------------------------------------------------------------------------|
| 1 Phil     | TIER-PHIL-001..005     | Mixed                 | P4 Constraint Hierarchy (capital), Principle 5 (Strategic Reasoning), §4 score, M5 Timing |
| 2 Port     | TIER-PORT-001..004     | Mostly P1             | §4 score ÷ cost, EXEC Stage 5 Portfolio Allocation, Tier 6 Counterfactual                 |
| 3 Rel      | TIER-REL-001..002      | P2                    | SIG-PATH equity inputs, INT-CRM positive notes, Tier 7 RelationshipEdge                   |
| 4 Route    | TIER-ROUTE-001..003    | P2 / P1 fallback      | ACCESS_STRENGTH (§3), SIG-PATH connector registry, GOV-LAD-04 (notify on assignment)      |
| 5 Esc      | TIER-ESC-001..003      | Mostly P1             | GOV Autonomy Ladder (A3/A4), Scarcity Economics, P3 Calibrated Confidence                 |
| 6 Attn     | TIER-ATTN-001..003     | P2 / P1 architectural | Scarcity Economics (every action priced), OBS-01 logging                                  |
| 7 Human    | TIER-HUMAN-001..003    | P2                    | SIG-TENSION (political complexity), Cognitive Load is dual-use with INT-SIG               |
| 8 Patience | TIER-PATIENCE-001..003 | Mixed                 | M5 Timing & Window, EPIS-COV-02 coverage cap, EXEC stage gating                           |

### §5 v3 COMPLETE — WHAT IT ESTABLISHED

§5 is no longer Tier Assignment & Routing. It is the engine's **Resource Allocation Operating System**: the layer that decides how scarce organizational capital is deployed across uncertain opportunities. Eight layers, 27 new rules, with each rule marked Phase 1 (MVP-eligible) or Phase 2 (design-now-activate-later). OPEN-Q4 resolves architecturally as expertise-aware routing (Phase 2 implementation); OPEN-Q13 partially resolves through portfolio capacity gating, with the 24h cooldown remaining as the Phase 1 fallback.

The architectural test from now on: every §5 rule must be interpretable as an attention-allocation rule (TIER-ATTN-003). Rules that fail this test are not §5 rules and belong elsewhere in the architecture.

### PLACEMENT MAP UPDATE

§5 status: **Complete** (all OPEN-Qs resolved or downgraded to Phase 2 build items with safe defaults). Engineering may build against §5 in Phase 1 mode using the documented safe defaults; Phase 2 rules activate as their prerequisite data infrastructure comes online.

### ON THE HORIZON

§6 Enrichment Engine — the 7-step workflow, buying-committee discovery, email verification states, strategic initiative analysis, AI account summary generation, refresh cadence — bridged to INT-ENR's verified-only, cost-gated, fallback-chained architecture. Or, alternatively, a §5 v4 frontier pass on multi-owner / team markets / relationship portfolio if v3 stands engineering review.