

# ABM BUSINESS LOGIC BIBLE

## §6 v2 (DEEP) — ENRICHMENT ENGINE

3 Architectural Iterations · 13 New Rules · 1 New OPEN-Q · Cross-Section Implications

§6 (the operational summary) specified the 7-step workflow, the verified-only gate, the fallback chain, and the cost-tiered cadence. This deeper pass excavates three architectural questions the original spec didn't reach: what does 'reachable' actually mean (Iteration 1); what does 'strategic priority' actually capture (Iteration 2); and is time-based refresh the right cadence model at all (Iteration 3). Each iteration changes the architecture, not just the rule count. Each touches §4 (scoring) and §7 (personalization) downstream.

|            |                             |                      |        |
|------------|-----------------------------|----------------------|--------|
| Iterations | 3                           | New rules            | 13     |
| New OPEN-Q | 1 (cadence vs event-driven) | Cross-section impact | §4, §7 |

### FRAMING — WHAT IS ENRICHMENT REALLY OPTIMIZING FOR?

§6's original framing: enrichment turns a HOT account into a reachable committee with a grounded dossier. The deeper question this pass attacks: **what is enrichment actually optimizing for, and where is it lying to us?** Three places it currently lies: it confuses delivery-reachability with response-reachability (Iteration 1); it treats publicly observable priorities as the buying drivers (Iteration 2); and it uses a fixed time-based refresh when the underlying reality is event-driven (Iteration 3).

## ITERATION 1 — THE REACHABILITY ILLUSION

§6 step 3 verifies that an email is deliverable. §4 dimension 3 (Reachability) credits 30 points for a verified email + LinkedIn. The hidden assumption: verified-to-deliver = reachable. It isn't. A C-suite executive at a Saudi bank may have a verified email that is filtered by an assistant, auto-archived into a folder no human reads, or routes to a generic intake address. The engine has high confidence in its data, and the data is measuring the wrong thing. This is the EPIS epistemic failure mode — high source reliability, wrong target construct.

### Why does this matter?

Because §4 over-rates accounts whose contacts have verified-but-unresponsive emails. They promote to HOT. They consume enrichment capital, AI compute, reviewer minutes, founder attention (§5 v3 attention budget). They produce zero response. The engine then blames messaging or timing, when the real failure was an unresponsive contact dressed as a reachable one.

### Architectural fix — split Reachability

Reachability is two distinct constructs: (a) **Delivery Reachability** — the email won't bounce; the LinkedIn profile exists; the data layer can address this person. (b) **Response Reachability** — a human at this address is likely to read and possibly respond. They require different evidence and decay at different rates. §4.3 must read Response Reachability, not Delivery Reachability. Delivery is the *precondition*; Response is the *measure*.

| Construct             | Evidence                                                                                                                                                     | Decay                                           | Used by                           |
|-----------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------|-------------------------------------------------|-----------------------------------|
| Delivery Reachability | Verified email (INT-ENR-02), LinkedIn URL valid, no bounce history                                                                                           | Slow — a verified email stays verified for ~30d | §9 deliverability layer; INT-EML  |
| Response Reachability | Response history (any prior reply), seniority-access pattern (assistant-gated yes/no), engagement-on-prior-sequence, peer-response-rate-for-persona-and-bank | Fast — changes with each engagement event       | §4.3 Persona Reachability scoring |

### New rules

| Rule ID      | Specification                                                                                                                                                                                                                                                                                |
|--------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| ENR-REACH-01 | Reachability is decomposed into TWO independent scores per contact: Delivery Reachability (0–1, deliverability precondition) and Response Reachability (0–1, engagement likelihood). They are stored separately, decay at different rates, and serve different consumers.                    |
| ENR-REACH-02 | Delivery Reachability is the precondition: contacts with Delivery Reachability = 0 are excluded from any outreach regardless of Response Reachability.                                                                                                                                       |
| ENR-REACH-03 | Response Reachability is derived from: prior-reply history (highest signal), gatekeeper presence (assistant-filtered address → lower score), engagement-on-prior-sequence (opens/clicks), and peer-response-rate (response rate for the same persona at comparable banks).                   |
| ENR-REACH-04 | §4.3 Persona Reachability dimension reads Response Reachability, not Delivery Reachability. A contact with verified email but Response Reachability ≈ 0 (assistant-filtered, no prior response, no peer signal) contributes minimal points to the score — closing the reachability illusion. |

## ITERATION 2 — THE STRATEGIC INITIATIVE TRAP

ENR-INIT-01 has the AI identify the account's top 3 strategic priorities and rank them. M0 hallucination governance ensures every named priority is grounded in a sourced fact — a public announcement, a hiring pattern, an annual report. §7 Agent 5 then selects an outreach angle from these priorities. The hidden assumption: priorities the engine can *observe* are the priorities that *drive buying*.

### Why does this matter?

Because the observable priorities are the *public* priorities — what the bank announces, hires for, discusses in conferences. The actual buying priorities are often internal and contested. A CTO who publicly champions cloud migration may be privately fighting to keep an on-prem investment alive because his political position depends on it. The SIG-TENSION reasoning stream (§3 reasoning streams) exists precisely to detect this asymmetry, but ENR-INIT-01 collapses it: the engine produces a perfectly-grounded cloud-migration pitch and routes it to a political loser. The message is correct on the public axis and wrong on the buying axis.

### Architectural fix — separate Public from Inferred

The dossier carries TWO initiative layers: **Public Strategic Initiatives** (sourced facts, high confidence, suitable as message context) and **Inferred Buying Drivers** (SIG-HYP hypotheses about who actually wants what, plus SIG-TENSION friction signals; lower confidence, RCM-stamped, used to gate angle selection). Public Initiatives feed the message content; Inferred Buying Drivers gate which angle is actually chosen — explicitly carrying their lower confidence per EPIS-RCM-04 (low-confidence withhold rule).

| Layer                        | Evidence                                                                        | Confidence                            | Role in §7                                              |
|------------------------------|---------------------------------------------------------------------------------|---------------------------------------|---------------------------------------------------------|
| Public Strategic Initiatives | Announcements, hires, annual report, conference speeches                        | High (sourced + corroborated)         | Message context — the topic                             |
| Inferred Buying Drivers      | SIG-HYP hypotheses (who wants this); SIG-TENSION friction signals (who resists) | Lower (causal inference; RCM-stamped) | Angle gating — whether to pitch this topic, and to whom |

### New rules

| Rule ID          | Specification                                                                                                                                                                                                                                                                                                          |
|------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| ENR-INIT-DEEP-01 | The account dossier carries TWO initiative layers: Public Strategic Initiatives (sourced, high-confidence, content layer) and Inferred Buying Drivers (causal inference from SIG-HYP + SIG-TENSION; RCM-stamped; gating layer). They are stored separately and never collapsed.                                        |
| ENR-INIT-DEEP-02 | §7 Agent 5 (Strategy Agent) receives BOTH layers. It selects an outreach topic from Public Initiatives and validates the angle against Inferred Buying Drivers. A Public Initiative that conflicts with the Inferred Buying Driver for the target persona is downgraded or rejected; the conflict is recorded for OBS. |
| ENR-INIT-DEEP-03 | When Inferred Buying Drivers carry confidence below the EPIS-RCM-04 threshold, Strategy Agent must default to safer, persona-generic angles (compliance, scalability) rather than persona-specific contested angles — anti-false-confidence applied to message strategy.                                               |
| ENR-INIT-DEEP-04 | Inferred Buying Drivers are bounded to publicly-observable signals (EPIS-ETH-01); the engine does NOT infer private psychology of a named person. SIG-TENSION conflict detection uses only public statements, hiring patterns, and contradictory initiatives — never speculation.                                      |

## ITERATION 3 — THE STALE-DOSSIER FAILURE MODE

ENR-REF-01 sets HOT refresh at 30 days. Reasonable on average. Catastrophic in specific windows. Two cases break it: (a) a HOT account whose CTO leaves on Day 15 — the dossier is 15 days from its next refresh, the contact section is fundamentally wrong, and §7 generates outreach addressed to a person who no longer holds the role. (b) A HOT account whose enrichment completed at the start of Ramadan when public hiring/news volume is at ~50% — the dossier under-samples activity and the engine pitches into a misread of the account's posture. The 30-day cadence is **time-based when the underlying reality is event-based and contextually-conditioned**.

### Architectural fix — event-driven re-enrichment, with cadence as a floor

Replace fixed cadence with explicit event-driven triggers: a SIG-EXEC executive-move on a committee contact triggers immediate contact-targeted re-enrichment; a SIG-VENDOR contract event triggers vendor-section refresh; entering/exiting Ramadan or Eid windows flags the dossier as low-confidence-period (not automatic refresh, but a marker that any data acquired in that window decays faster). The 30-day rule becomes a **floor** (never exceed 30 days without refresh), not a **cadence** (re-enrich at 30 days regardless of need). Most accounts refresh less often — a Scarcity win. Urgent accounts refresh sooner.

| Trigger                                          | Refresh scope                                                  | Cost shape                              |
|--------------------------------------------------|----------------------------------------------------------------|-----------------------------------------|
| SIG-EXEC executive move on committee contact     | Contact-targeted: re-resolve that contact only                 | Low (one verification + LinkedIn check) |
| SIG-VENDOR contract / RFP / case-study event     | Vendor section + displacement-gap re-analysis                  | Medium                                  |
| SIG-REG SAMA / CMA mandate touching this account | Regulatory section + initiative re-analysis                    | Medium                                  |
| SIG-MOBILITY: connector or P1 path member moves  | SIG-PATH section + ACCESS_STRENGTH re-score                    | Low-medium                              |
| Ramadan / Eid window enter or exit               | Mark dossier with low-confidence-period flag (no auto refresh) | Zero                                    |
| 30-day floor reached without any of the above    | Full refresh (matches V2 cadence behavior)                     | Full enrichment cost                    |

**Per your decision: BOTH cadence and event-driven are specified; the choice is flagged as an open question.**

### New rules

| Rule ID         | Specification                                                                                                                                                                                                                                                                                                                                                                                                                                                                           |
|-----------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| ENR-REF-DEEP-01 | Event-driven re-enrichment triggers: (a) SIG-EXEC executive-move on a committee contact → contact-targeted refresh; (b) SIG-VENDOR contract/RFP/case-study event → vendor section + displacement-gap refresh; (c) SIG-REG mandate touching this account → regulatory section refresh; (d) SIG-MOBILITY on a P1 connector → SIG-PATH + ACCESS_STRENGTH refresh; (e) Ramadan/Eid enter or exit → dossier marked low-confidence-period (no automatic refresh, decay-rate adjustment only). |
| ENR-REF-DEEP-02 | Event-driven refreshes are SCOPED — they refresh only the affected dossier section, not the full account. This is the cost-aware variant: most events trigger low-cost partial refreshes rather than full enrichment cycles.                                                                                                                                                                                                                                                            |
| ENR-REF-DEEP-03 | Ramadan/Eid and other contextually-conditioned periods modify the decay rate of facts acquired during them: enrichment-acquired-during-Ramadan facts decay at 1.5× normal rate, reflecting under-sampling of underlying activity. Implements EPIS-HALF-01 contextual half-life.                                                                                                                                                                                                         |
| ENR-REF-DEEP-04 | 30-day HOT cadence becomes a FLOOR (never exceed 30 days without refresh), not a CADENCE (do not refresh at 30 days if event-driven refresh has occurred within the last 30 days). Cost reduction: most accounts will refresh less often; the urgent ones sooner.                                                                                                                                                                                                                       |

#### OPEN-DESIGN-QUESTION OPEN-Q6.4 — CADENCE vs EVENT-DRIVEN

**Question:** Should event-driven re-enrichment REPLACE the 30-day cadence (ENR-REF-DEEP-04 floor-only) or SUPPLEMENT it (cadence remains, event-driven adds early refresh)?

**Candidate options:** (a) Replace — cadence becomes a floor only; event-driven is the primary refresh discipline; lower total cost, requires reliable event detection. (b) Supplement — cadence stays at 30d; event-driven adds opportunistic early refreshes; higher cost but more conservative; cadence covers event detection misses. (c) Phased — ship Supplement at MVP (belt-and-suspenders), evaluate after 90 days, then decide on Replace based on event-detection reliability data.

**Trade-off:** (a) trusts the event-detection layer and saves cost. (b) is paranoid and pays for redundancy. (c) buys data before committing.

**Dependent rules:** ENR-REF-01 (V2 cadence), ENR-REF-DEEP-01..04. Engineering safe default: ship (c) phased approach — cadence + event-driven together at MVP, with metrics on event-detection coverage; resolve at 90d review.

## CROSS-SECTION IMPACT — WHAT §6 v2 CHANGES IN §4 AND §7

These three iterations don't only deepen §6; they propagate into §4 (scoring) and §7 (personalization). Honest accounting of the downstream changes:

| Iteration                   | Touches                                   | What changes downstream                                                                                                                                                                                                                            |
|-----------------------------|-------------------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| 1 Reachability illusion     | §4.3 Persona Reachability                 | Dimension reads Response Reachability, not Delivery. SCORE-REACH-01... inputs change; the 30-point verified-email-plus-LinkedIn formula remains the Phase 1 default until Response Reachability has enough data.                                   |
| 2 Strategic initiative trap | §7 Agent 5 Strategy + Agent 6 Message Gen | Strategy Agent reads BOTH initiative layers; angle selection is gated by Inferred Buying Drivers; conflicting public-vs-inferred angles are downgraded. §7 must be specified to handle the two-layer input — build this into §7 when we get there. |
| 3 Stale dossier             | §6 ENR-REF-01 + INT-SIG event streams     | Event triggers from §3/INT-SIG fan into §6 refresh pipeline. Re-uses existing SIG-EXEC/VENDOR/REG/MOBILITY streams — no new infrastructure; only new bindings.                                                                                     |

### §4 IS NOW SOFT — ITERATION 1 EFFECT

Iteration 1 mechanically changes what §4.3 measures. The 35/30/20/15 weights remain valid; the question is what dimension 3 (Reachability × 0.20) is reading. §4 should be flagged as **partially superseded** on the Placement Map: it remains operationally correct for Phase 1 (Delivery Reachability fallback) but expects an upgrade to Response Reachability when ENR-REACH-03 inputs have sufficient data. This is a future-edit promise, not a present rollback.

## EXHAUSTION ASSESSMENT — §6 AFTER THE DEEP PASS

| Concern                                | §6 v1 | §6 v2                            | Frontier (v3 candidate)                                                   |
|----------------------------------------|-------|----------------------------------|---------------------------------------------------------------------------|
| 7-step workflow operational concretion | ~95%  | ~95%                             | —                                                                         |
| Verified-only + fallback discipline    | ~95%  | ~95%                             | —                                                                         |
| Refresh cadence                        | ~85%  | ~75%* (now contested; OPEN-Q6.4) | Resolve cadence vs event-driven empirically                               |
| Reachability semantics                 | ~50%  | ~80%                             | Response Reachability calibration; peer-response-rate model               |
| Initiative analysis                    | ~70%  | ~80%                             | Inferred Buying Driver confidence model; SIG-TENSION integration depth    |
| Cost economics                         | ~60%  | ~70%                             | OPEN-Q6.2 budget structure; cross-section attention budget reconciliation |
| Failure modes                          | ~70%  | ~70%                             | HOT_UNENRICHED recovery playbooks; contact-disappearance handling         |

\*§6 v2 reduces this from 85% to 75% because the deep pass opened the cadence question; this is honest, not regression. A higher number that papers over an open question is false confidence.

## §6 v2 RULE ROLL-UP & PLACEMENT STATUS

| Iteration               | Rules                | Phase                                       | Downstream binding                      |
|-------------------------|----------------------|---------------------------------------------|-----------------------------------------|
| 1 Reachability illusion | ENR-REACH-01..04     | P1 architecture, P2 data                    | §4.3 (scoring); §9 deliverability layer |
| 2 Initiative trap       | ENR-INIT-DEEP-01..04 | P2 (requires SIG-HYP/TENSION at confidence) | §7 Agent 5 + Agent 6 (when built)       |
| 3 Stale dossier         | ENR-REF-DEEP-01..04  | P1 (re-uses INT-SIG event streams)          | §3 INT-SIG (event fan-out)              |
| OPEN-Q6.4               | —                    | Phased default until resolved               | Affects ENR-REF-01 vs ENR-REF-DEEP-04   |

### §6 v2 COMPLETE — WHAT IT ESTABLISHED

Three architectural iterations changed what §6 is doing rather than just adding rules. Iteration 1 split Reachability into Delivery and Response — the engine stops confusing 'won't bounce' with 'might respond'. Iteration 2 split initiative analysis into Public Strategic Initiatives (sourced, high-confidence, message content) and Inferred Buying Drivers (causal inference, RCM-stamped, angle gating) — the engine stops pitching to political losers on grounded-but-wrong angles. Iteration 3 turned refresh from a time-based cadence into an event-driven trigger model with cadence as a floor — most accounts refresh less often, urgent ones refresh sooner. 13 new rules.

§6 v2 propagates into §4 (Reachability semantics) and §7 (two-layer initiative input). §4 should be marked partially superseded on the Placement Map (the dimension reads a richer construct than before; the weights stand). §7 must be specified to handle the two-layer input when we get there.

### PLACEMENT MAP UPDATE

§6 status: **Complete-with-Opens** (OPEN-Q6.1–6.3 from §6 v1 still open; OPEN-Q6.4 added in this pass). Phase 1 engineering implements the safe defaults; resolutions are A4 governance actions per TIER-OPEN-DISCIPLINE-02. §4 status: **Complete-with-future-edit** — the dimension semantics will read richer constructs once Iteration 1 Phase 2 data is available.