

# ABM BUSINESS LOGIC BIBLE

## §7 v2 — TRUST & INFLUENCE CONSTRUCTION SYSTEM

10 Iterations · 7 Architectural Layers · 27 New Rules · 3 Guardrails Made Explicit

§7 v1 specified the seven-agent personalization architecture. §7 v2 reframes the engine itself. **§7 is not a message-generation engine. It is a trust-and-reputation construction system.** The purpose of personalization is not to maximize replies; it is to convert intelligence into credible, context-aware interactions that compound trust, reinforce a coherent market narrative, manage political risk, and grow Decimal's long-term reputation within the banking ecosystem. Messages are the visible artifacts of that process, not the process itself.

|            |    |                      |              |
|------------|----|----------------------|--------------|
| Iterations | 10 | Architectural layers | 7            |
| New rules  | 27 | Guardrails added     | 3 (explicit) |

### THE OPTIMIZATION OBJECTIVE SHIFT — THE DEEPEST DISCOVERY

Up to §7 v1, the implicit objective was reply rate. §7 v2 makes the real objective explicit: a multi-term value function combining **Trust Growth**, **Narrative Coherence**, **Stakeholder Consensus Progress**, and **Reputation Yield** — with conversion as only one term among several. Every layer below either feeds this objective or constrains it. AI-REP-003 makes this the formal optimization target for Agent 5.

### THREE GUARDRAILS MADE EXPLICIT IN THIS PASS

**Political sensitivity stays sourced-only.** AI-POL-001 binds strictly to SIG-TENSION's public, sourced friction signals — the engine never free-reasons about 'is this politically dangerous' from inference alone.

**Stakeholder-specific objectives are reassurance, not manipulation.** AI-INF-002 defines 'Procurement Neutralization' and similar objectives as reducing legitimate friction through clarity, never as engineering around a stakeholder's legitimate concerns.

**The reputation-over-conversion trade-off is threshold-bound, not vibes-based.** AI-REP-006 requires a measurable Reputation Yield delta to exceed a configured floor before the engine may deprioritize conversion — otherwise Agent 5 defaults to pursuing the close.

## THE SEVEN LAYERS — §7 v2 ARCHITECTURE MAP

| # | Layer                                 | From iter. | What it adds                                                                           |
|---|---------------------------------------|------------|----------------------------------------------------------------------------------------|
| 1 | Narrative Coherence (Narr)            | 1, 9       | Messages accumulate into a market story; narrative stage changes framing               |
| 2 | Message Portfolio / Sequence (Seq)    | 2          | A message has a Portfolio Role; not every message asks for a meeting                   |
| 3 | Trust Accumulation & Transfer (Trust) | 3, 8       | Trust is measurable, compounds, and is transferable via introducers                    |
| 4 | Political Risk (Pol)                  | 4          | Truth isn't enough; sourced-only sensitivity rating; complements two-layer initiatives |
| 5 | Reputation Compounding (Rep)          | 5, 10      | The market is a memory network; every outreach has a Reputation Yield                  |
| 6 | Executive Attention (Attn)            | 6          | The bottleneck is cognitive bandwidth, not persuasion; attention-to-value ratio        |
| 7 | Stakeholder Influence Network (Inf)   | 7          | Persuasion is distributed across the committee with different objectives per role      |

## LAYER 1 — NARRATIVE COHERENCE (NARR)

Individually-correct messages can collectively confuse. A CTO hearing 'API modernization' in Month 1, a CDO hearing 'AI transformation' in Month 2, and Procurement hearing 'cost reduction' in Month 3 leaves the bank unsure what Decimal stands for. Messages are not independent; they accumulate into reputation and narrative. Iteration 9 adds: narratives themselves have a lifecycle — Emerging → Growing → Established → Inevitable — and framing must match the stage.

| Rule ID    | Phase | Specification                                                                                                                                                                                                                                                                                         |
|------------|-------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| AI-NARR-01 | P1    | Every account is assigned an <b>Active Narrative</b> (e.g. API Modernization, Digital Origination, AI Banking, Open Banking) derived from the dossier's Public Strategic Initiatives (§6 v2 ENR-INIT-DEEP-01). The narrative is the account-level story Decimal is telling, not a per-message choice. |
| AI-NARR-02 | P1    | Agent 5 may vary the specific angle within a message but must remain inside the account's Active Narrative. A pivot to a different narrative requires an explicit account-level narrative change, not a per-message improvisation.                                                                    |
| AI-NARR-03 | P1    | Narrative drift — measured as the semantic distance between consecutive message angles for the same account — above a configured threshold triggers a review flag. Consistency is what creates trust; drift is observable and gated, not silently tolerated.                                          |
| AI-NARR-04 | P2    | Every strategy angle is tagged with a <b>Narrative Stage</b>: Emerging, Growing, Established, or Inevitable, derived from the strength and recency of the account's and the market's signal activity around that narrative (SIG-NEWS/EVENT volume + SIG-REG urgency).                                 |
| AI-NARR-05 | P2    | Message framing changes by Narrative Stage: Emerging → education; Growing → differentiation; Established → execution; Inevitable → urgency. Agent 5 reads the stage and selects the matching frame; a mismatch (urgency framing on an Emerging narrative) is rejected by QC.                          |
| AI-NARR-06 | P2    | Agent 5's objective prioritizes Narrative Adoption progress before Feature Awareness — the strategy brief states which stage-transition this message is intended to advance, not merely which feature it highlights.                                                                                  |

## LAYER 2 — MESSAGE PORTFOLIO / SEQUENCE (SEQ)

Executives receive sequences, not isolated messages. A perfect message can be harmful at the wrong stage of the relationship — asking for a meeting in the first touch reads as presumptuous; never asking reads as aimless. Optimization must happen at sequence level, not message level.

| Rule ID   | Phase | Specification                                                                                                                                                                                                                                                             |
|-----------|-------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| AI-SEQ-01 | P1    | Agent 5 reads the contact's <b>Interaction Stage</b>: First Contact, Familiar, Engaged, or Active Discussion — derived from §10 engagement history (touch count, opens, replies) and contact_state (V2 §2.2).                                                             |
| AI-SEQ-02 | P1    | Every message is assigned a <b>Portfolio Role</b>: Awareness, Credibility, Curiosity, Proof, or Commitment. Not every message should ask for a meeting — an Awareness-role message ends in a low-friction curiosity hook; a Commitment-role message ends in a direct CTA. |
| AI-SEQ-03 | P1    | Portfolio Role is determined jointly by Interaction Stage and the sequence touch number (V2 §9.2 Touch 1–4): Touch 1 defaults to Awareness/Credibility; Touch 4 defaults to Proof/Commitment. Agent 5 may override the default with explicit justification.               |
| AI-SEQ-04 | P1    | Message optimization occurs at the sequence level: Agent 5's strategy brief states not just this message's angle but its role in the planned arc of touches, so a later message doesn't repeat an earlier message's role.                                                 |

## LAYER 3 — TRUST ACCUMULATION & TRANSFER (TRUST)

Trust is not binary and not solely earned through direct interaction. It accumulates through repeated correctness (Iteration 3) and can be borrowed through introducers and ecosystem position (Iteration 8 — a Mastercard or KPMG introduction partially transfers trust before any direct interaction occurs).

| Rule ID     | Phase | Specification                                                                                                                                                                                                                                                                                                              |
|-------------|-------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| AI-TRUST-01 | P2    | Every contact carries a <b>Trust Accumulation Score</b> built from: personalization accuracy (did prior claims hold up?), response quality, relevance ratings, and engagement depth over time. The score is append-only and visible in the dossier.                                                                        |
| AI-TRUST-02 | P1    | Aggressive CTAs (direct meeting asks, pricing discussions) are forbidden until the Trust Accumulation Score crosses a configured threshold OR the path carries high Trust Transfer Potential (AI-TRUST-04). Cold trust gets low-friction asks ('worth comparing notes?'); high trust gets direct asks ('15-minute call?'). |
| AI-TRUST-03 | P2    | Agent 5 optimizes Trust Growth before Meeting Acquisition when Trust Accumulation Score is below threshold: the strategy brief's stated objective for early-stage contacts is explicitly 'build trust', not 'book meeting', even though the message still carries a CTA.                                                   |
| AI-TRUST-04 | P2    | Every resolved SIG-PATH carries a <b>Trust Transfer Potential</b> score derived from: relationship strength of the connector, introducer credibility, and the connector's ecosystem position (per §3 SIG-PATH connector registry and §5 v3 TIER-ROUTE-001 Expertise Graph).                                                |
| AI-TRUST-05 | P2    | High Trust Transfer Potential paths (warm introductions through credible connectors) permit more direct CTAs and more ambitious asks earlier in the sequence than the Trust Accumulation Score alone would allow — borrowed trust substitutes partially for earned trust.                                                  |
| AI-TRUST-06 | P2    | Agent 5's strategy brief must explicitly name the trust source being leveraged when one exists (e.g. 'warm path via [connector]') — not merely cite personalization facts. Trust sources are a distinct input from content facts.                                                                                          |

## LAYER 4 — POLITICAL RISK (POL)

Truth is not enough. A factually grounded claim — 'we know cloud migration is your priority' — can be politically dangerous if compliance is privately resisting it. This complements (and is bound by) the two-layer initiative model from §6 v2: SIG-TENSION already detects organizational friction from public signals only; Layer 4 wires that detection into message-level risk rating.

| Rule ID   | Phase | Specification                                                                                                                                                                                                                                                                                                                                                                                                                                                                        |
|-----------|-------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| AI-POL-01 | P1    | Every personalization claim derived from an Inferred Buying Driver (§6 v2 ENR-INIT-DEEP-01) receives a <b>Political Sensitivity Rating</b> : Low, Medium, or High. The rating is computed strictly from SIG-TENSION's sourced friction signals (public statements, contradictory initiatives, budget conflicts) — NEVER from open-ended inference about internal politics. A claim with no SIG-TENSION basis defaults to Low sensitivity, not High; the engine does not invent risk. |
| AI-POL-02 | P1    | High-sensitivity claims require human review regardless of the account's tier or the §8 QC-HUMAN-004 post-MVP relaxation status. A High-sensitivity message never auto-sends, even after the 90-day relaxation criterion (§7 OPEN-Q7.1) is met for that segment.                                                                                                                                                                                                                     |
| AI-POL-03 | P1    | When political risk is elevated (Medium or High), Agent 5 prefers Public Strategic Initiatives over Inferred Buying Drivers as the message's grounding — the safer, sourced, less politically-loaded angle is chosen, consistent with §6 v2 ENR-INIT-DEEP-03's anti-false-confidence default.                                                                                                                                                                                        |
| AI-POL-04 | P1    | Political Sensitivity Rating is recomputed whenever the underlying SIG-TENSION signal changes; a rating is never assigned once and treated as permanent.                                                                                                                                                                                                                                                                                                                             |

## LAYER 5 — REPUTATION COMPOUNDING (REP)

The market is a memory network. Executives, consultants, and board members move between banks (SIG-MOBILITY, §3). One irrelevant outreach is small damage; a hundred is remembered. The deepest asset the engine builds is not meetings — it is Decimal's standing across the ecosystem.

| Rule ID   | Phase | Specification                                                                                                                                                                                                                                                                                                                                                                                                                                          |
|-----------|-------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| AI-REP-01 | P2    | Every outreach action updates a <b>Reputation Impact Score</b> : Positive, Neutral, or Negative, derived from engagement quality (reply sentiment, silence after high-personalization effort, explicit negative feedback).                                                                                                                                                                                                                             |
| AI-REP-02 | P2    | Contact-level Reputation Impact rolls up into <b>Account Reputation</b> ; Account Reputation rolls up into <b>Market Reputation</b> (an ecosystem-wide rolling score). This uses the same roll-up pattern as §4's score architecture (append-only, never overwritten).                                                                                                                                                                                 |
| AI-REP-03 | P1    | Agent 5's <b>formal optimization objective</b> is Expected Opportunity Value combined with Expected Reputation Value — not conversion probability alone. This is the engine's deepest objective-function change in §7 v2; every other layer in this document feeds into this combined objective.                                                                                                                                                       |
| AI-REP-04 | P2    | Every interaction is assigned a <b>Reputation Yield</b> estimate — expected future value generated through ecosystem memory, independent of this interaction's immediate conversion outcome.                                                                                                                                                                                                                                                           |
| AI-REP-05 | P1    | Account-level success metrics (§13 Analytics) report BOTH Immediate Outcome (reply, meeting, opportunity) and Reputation Impact, never immediate outcome alone.                                                                                                                                                                                                                                                                                        |
| AI-REP-06 | P1    | The engine may deprioritize a higher-conversion action in favor of a higher-reputation action ONLY when the Reputation Yield delta exceeds a configured floor (guardrail, not a vibe). Absent that explicit threshold being met, Agent 5 defaults to pursuing the highest-conversion action available — reputation is a constraint that can override conversion, but only when the trade-off is measurable and material, not as a default disposition. |

## LAYER 6 — EXECUTIVE ATTENTION (ATTN)

A KSA bank CTO receives hundreds of information inputs daily — vendor outreach, internal updates, board requests, regulatory asks, consulting decks. The real bottleneck is not persuasion quality; it is earning cognitive bandwidth. The engine is not competing for replies — it is competing for mindshare.

| Rule ID    | Phase | Specification                                                                                                                                                                                                                                                                                                                                                               |
|------------|-------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| AI-ATTN-01 | P1    | Every drafted message receives an <b>Attention Cost</b> estimate — the mental effort required to process it, derived from length, structural complexity, and number of distinct ideas presented.                                                                                                                                                                            |
| AI-ATTN-02 | P1    | Agent 5 optimizes <b>Attention-to-Value Ratio</b> , not message length minimization alone. A 50-word high-relevance message with low cognitive effort beats a 300-word high-relevance message with high effort — this formalizes V2's existing 150/80-word caps as a ratio objective, not just a ceiling.                                                                   |
| AI-ATTN-03 | P1    | Different personas carry different <b>Attention Budgets</b> : C-suite (very low — shortest, most distilled framing); Director (moderate); Manager (higher — more detail tolerated). Agent 6's message complexity adapts to the target's seniority-derived attention budget, extending AI-AG6-07's seniority-matched framing rule with an explicit cognitive-cost dimension. |
| AI-ATTN-04 | P2    | Attention Cost is tracked against actual engagement outcomes (open-to-reply conversion by Attention Cost bracket) to calibrate the cost model over time — requires the Tier 6 outcome loop.                                                                                                                                                                                 |

## LAYER 7 — STAKEHOLDER INFLUENCE NETWORK (INF)

No bank buys through one person (Principle 2). If the CTO is enthusiastic but Compliance, Procurement, and Operations are not, the deal dies. Persuasion must be distributed across the committee, with each stakeholder receiving a message aligned to their specific concern — never the same pitch copy-pasted across roles.

| Rule ID   | Phase | Specification                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                            |
|-----------|-------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| AI-INF-01 | P1    | Every message belongs to a named <b>Stakeholder Influence Strategy</b> : Champion Building, Consensus Building, Executive Sponsorship, Procurement Neutralization, or Risk Reduction. The strategy is set per-contact, not per-account.                                                                                                                                                                                                                                                                                  |
| AI-INF-02 | P1    | Stakeholder-specific objectives mean <b>reducing legitimate friction through clarity and reassurance</b> — never engineering around a stakeholder's legitimate concerns. 'Procurement Neutralization' means addressing real cost/process concerns with concrete, honest answers, not minimizing or routing around procurement's role. A message strategy that asks Agent 6 to obscure a real constraint is rejected by QC-007 (compliance check).                                                                        |
| AI-INF-03 | P1    | Messages to different stakeholders at the same account may pursue different objectives: CTO receives enthusiasm-building content; Compliance receives risk-reduction content (sourced, specific, never dismissive); Procurement receives friction-reduction content (clear process, pricing transparency). Same opportunity, different persuasion objective per role — this is the operational form of multi-threaded outreach already required by §4 (Reachability multi-thread bonus) and §6 (3–7 committee contacts). |
| AI-INF-04 | P2    | Account-level success tracking includes a <b>Committee Influence Progress</b> metric — how many buying-committee members have moved from UNCONTACTED/SEQUENCED toward QUALIFIED — reported alongside (not instead of) reply rate.                                                                                                                                                                                                                                                                                        |
| AI-INF-05 | P1    | A Stakeholder Influence Strategy assignment must be consistent with the Political Sensitivity Rating (Layer 4): a Procurement Neutralization message cannot rely on a High-sensitivity Inferred Buying Driver claim; it defaults to Public Strategic Initiative grounding per AI-POL-03.                                                                                                                                                                                                                                 |

## §7 v2 RULE ROLL-UP & PLACEMENT STATUS

| Layer   | Rules           | Phase   | Cross-section binding                                                |
|---------|-----------------|---------|----------------------------------------------------------------------|
| 1 Narr  | AI-NARR-01..06  | P1 / P2 | §6 v2 Public Initiatives, SIG-NEWS/EVENT volume                      |
| 2 Seq   | AI-SEQ-01..04   | P1      | V2 §9.2 touch cadence, §10 contact_state                             |
| 3 Trust | AI-TRUST-01..06 | P1 / P2 | SIG-PATH connector registry, §5 v3 Expertise Graph                   |
| 4 Pol   | AI-POL-01..04   | P1      | SIG-TENSION (sourced-only), §6 v2 two-layer initiatives, EPIS-ETH-02 |
| 5 Rep   | AI-REP-01..06   | P1 / P2 | §4 score append-only pattern, §13 Analytics, SIG-MOBILITY            |
| 6 Attn  | AI-ATTN-01..04  | P1 / P2 | AI-AG6-07 seniority framing, V2 §7.2 length caps                     |
| 7 Inf   | AI-INF-01..05   | P1 / P2 | §4 multi-thread bonus, §6 committee discovery, QC-007 compliance     |

### §7 v2 COMPLETE — WHAT IT ESTABLISHED

§7 is no longer a personalization engine. It is a **Trust and Influence Construction System**: it strategically allocates attention, builds stakeholder consensus, transfers trust through networks, accelerates narrative adoption, and compounds Decimal's reputation across the banking ecosystem. Messages are the visible artifacts of that process. Seven layers, 27 new rules, with three explicit guardrails: political sensitivity stays sourced-only (never free-inferred), stakeholder-specific objectives mean reassurance not manipulation, and the reputation-over-conversion trade-off requires a measurable threshold, not a default disposition.

Agent 5's formal optimization objective changes: Expected Opportunity Value combined with Expected Reputation Value, not conversion probability alone (AI-REP-03). This is the single deepest change §7 v2 makes — every other layer feeds this objective.

### PLACEMENT MAP UPDATE

§7 status: **Complete-with-Opens** (OPEN-Q7.1, OPEN-Q7.2 from §7 v1 carry forward unchanged — this pass did not touch them). The exhaustion assessment after this pass: Multi-Agent Design ~99%, Grounding/Hallucination ~99%, Strategy Selection ~95%, Narrative/Trust/Reputation theory ~70–85%. Remaining frontier (enterprise persuasion theory, executive attention calibration, reputation economics calibration) is increasingly empirical — requiring outcome data — rather than architectural.

### ON THE HORIZON

§8 Quality Control Engine — the 7 QC checks, regeneration limits, human-review SLA, and the post-MVP relaxation rule. §8 must now also validate against §7 v2's new constraints: narrative-drift flags (AI-NARR-03), Political Sensitivity Rating gating (AI-POL-02 forces human review regardless of relaxation status), and Stakeholder Influence Strategy consistency (AI-INF-05).