

ABM BUSINESS LOGIC BIBLE

§8 v2 — TRUST PRESERVATION & IRREVERSIBILITY MANAGEMENT

SYSTEM

5 Iterations · 5 Architectural Layers · 14 New Rules · 1 Guardrail Bound to Irreversibility

§8 v1 specified QC as the last gate before send. §8 v2 reframes the engine itself. **§8 is not a message-quality checker. It is the Trust Preservation and Irreversibility Management System of the ABM Engine.** Its purpose is not to detect message defects — it is to prevent irreversible damage to relationships, reputation, compliance posture, market memory, and confidence in the platform itself. Everything before §8 is revisable: a score can be recalculated, a dossier can be re-enriched, a strategy can be re-generated. The moment a message sends, none of that is true.

Iterations	5	Architectural layers	5
New rules	14	Guardrail	False-negative tolerance excludes hard gates

THE GUARDRAIL — FALSE-NEGATIVE TOLERANCE DOES NOT APPLY TO HARD GATES

Iteration 2 correctly identifies that over-blocking has a cost. But this tolerance applies ONLY to the seven mechanical checks. It does NOT apply to the three §8 v1 hard gates (narrative drift, political sensitivity, stakeholder-strategy). A blocked-but-safe message is recoverable (delay, regenerate, resend). A sent-but-politically-dangerous message is not — exactly the irreversibility Iteration 1 identifies. The hard gates remain strictness-maximizing regardless of false-negative cost.

THE FIVE LAYERS — §8 v2 ARCHITECTURE MAP

#	Layer	From iter.	What it adds
1	Irreversibility Mapping (Phil)	1	Every QC rule maps to a named Irreversible Risk Category; severity = expected damage, not frequency
2	False Positive / Negative Balance (Risk)	2	QC has two costs; optimize total expected risk — bounded by the guardrail
3	Reviewer Governance (Human)	3	Human reviewers are a decision system too; calibration tracked, blind-audited
4	Systemic Failure Detection (Sys)	4	A failure cluster is a module-health signal, not N independent problems
5	Trust Preservation (Trust)	5	The deepest asset protected is institutional trust, not message throughput

LAYER 1 — IRREVERSIBILITY MAPPING (PHIL)

Everything before §8 can be changed. Once a message sends, reality changes: executive perception, market memory (SIG-MOBILITY), the relationship. QC is not a quality function. It is an **irreversibility management function**.

Rule ID	Specification
QC-PHIL-01	Every QC rule (QC-RULE-001 through 010) maps to one or more Irreversible Risk Categories: Reputation, Relationship, Compliance, Market Memory, Political. A rule with no mapped category is incompletely specified.
QC-PHIL-02	QC severity = Probability × Magnitude × Reversibility Factor (0 for irreversible, approaching 1 for recoverable) — NOT raw frequency. A rare catastrophic failure (e.g. QC-RULE-009) gets highest protection despite low trigger rate.
QC-PHIL-03	Mapping: QC-001/002→Reputation; QC-004/005→Relationship; QC-006→Reputation; QC-007→Compliance; QC-RULE-008→Relationship+Market Memory; QC-RULE-009→Political+Compliance; QC-RULE-010→Compliance+Relationship.

LAYER 2 — FALSE POSITIVE / NEGATIVE BALANCE (RISK)

Rule ID	Specification
QC-RISK-01	Each of the seven mechanical checks carries an explicit False Positive Cost (good message blocked, measured against §4 timing-window decay) and False Negative Cost (bad message sent, measured against Layer 1). Both are visible in QC configuration.
QC-RISK-02	QC optimization target for QC-RULE-001..007 is Expected Total Risk = (FP Rate × FP Cost) + (FN Rate × FN Cost), NOT maximum strictness. Thresholds may loosen where FP Cost is empirically shown to dominate.
QC-RISK-03	GUARDRAIL: applies ONLY to QC-RULE-001..007. Does NOT apply to QC-RULE-008/009/010, which remain strictness-maximizing regardless of false-negative cost.

LAYER 3 — REVIEWER GOVERNANCE (HUMAN)

Human review has been treated as inherently safer than AI. It isn't, automatically — fatigue, haste, inconsistency, bias. This layer extends GOV-OVR-02 specifically to QC reviewers.

Rule ID	Specification
QC-HUMAN-08	Every QC reviewer carries a Reviewer Calibration Score: agreement rate with eventual outcomes, consistency across similar cases, and override-of-AI accuracy.
QC-HUMAN-09	Approval accuracy, override accuracy, and rejection accuracy tracked as THREE SEPARATE metrics per reviewer — a reviewer may excel at compliance issues but miss narrative-drift nuance.
QC-HUMAN-10	Periodic blind review audits: the same message routed to 2+ reviewers independently. Disagreement rate tracked; persistent disagreement on a check type triggers a calibration session, not blame.
QC-HUMAN-11	Reviewer Calibration Score feeds the same per-subsystem trust model as AI components (EPIS-TRUST-01): a degrading reviewer is flagged exactly as a degrading AI subsystem is flagged.

LAYER 4 — SYSTEMIC FAILURE DETECTION (SYS)

Rule ID	Specification
QC-SYS-01	§8 tracks Failure Clusters: a statistically significant spike in any check type within a rolling window, segmented by account, segment, persona, or agent-version.
QC-SYS-02	A Failure Cluster breaching threshold triggers a System Health Investigation routed to engineering/prompt-ops, NOT per-message regeneration. Regenerating 200 failing messages treats symptoms; investigation finds the shared root cause.
QC-SYS-03	Failure-cluster telemetry is the engine's early-warning system: §8 sees every output from every upstream agent, so a rising hallucination rate is detectable here first. Binds to FAIL-LEARN-01 as a QC-specific instance of that broader pattern.

LAYER 5 — TRUST PRESERVATION (TRUST)

The deepest layer. Every upstream module exists because trust exists. Repeated bad outreach collapses institutional trust on multiple fronts: reviewers stop trusting automation, founders stop trusting recommendations, the market stops trusting Decimal (§7 v2 Layer 5). QC's deepest job is protecting this multi-layered trust, not message throughput.

Rule ID	Specification
QC-TRUST-01	Every QC failure is categorized by Trust Impact: Low, Medium, High, Critical — distinct from Irreversible Risk Category. Trust Impact measures damage to confidence IN THE SYSTEM, not the specific relationship.
QC-TRUST-02	§8 tracks a Trust Preservation Score per upstream module (Agent 5, Agent 6, §6 Enrichment, §4 Scoring, Narrative layer): rate of QC intervention weighted by Trust Impact. A declining score is the system's early signal that humans are about to start distrusting that module.
QC-TRUST-03	§8's success metric explicitly includes Trust Preserved, not only Messages Approved. §13 Analytics KPIs are extended to report Trust Preservation Score per module alongside QC pass rate — a high pass rate with declining trust score is a warning, not a success.

§8 v2 RULE ROLL-UP & PLACEMENT STATUS

Layer	Rules	Cross-section binding
1 Phil	QC-PHIL-01..03	Maps every QC-RULE to a named risk category; foundational reframe
2 Risk	QC-RISK-01..03	§4 timing-window decay cost; guardrail excludes hard gates
3 Human	QC-HUMAN-08..11	GOV-OVR-02, EPIS-TRUST-01 per-subsystem trust
4 Sys	QC-SYS-01..03	FAIL-LEARN-01 Failure Pattern Record, OBS telemetry
5 Trust	QC-TRUST-01..03	§7 v2 AI-REP layer, §13 Analytics KPI extension

§8 v2 COMPLETE — WHAT IT ESTABLISHED

§8 is no longer a message-quality checker. It is the engine's Trust Preservation and Irreversibility Management System: a risk-allocation layer, a reviewer-governance layer, an upstream health-monitoring layer, and the final institutional defense of trust itself. 14 new rules, with one explicit guardrail: false-negative tolerance applies only to the seven mechanical checks, never to the three hard gates that protect against genuinely irreversible damage.

The deepest discovery: QC's success metric is not messages approved. It is trust preserved — per module, per reviewer, and institutionally.

PLACEMENT MAP UPDATE

§8 status remains Complete-with-Opens (OPEN-Q8.1, OPEN-Q8.2 unchanged). Exhaustion: QC Checks ~99%, Regeneration ~99%, Human Review Workflow ~95%, Narrative/Political controls ~95%, Reviewer Governance ~75%, Failure Analytics ~80%, Trust Preservation ~75%, System Health Monitoring ~70%. Remaining frontier (reviewer psychology, organizational trust calibration, failure pattern prediction, market reputation risk propagation, human-AI disagreement theory) is increasingly empirical rather than architectural.

ON THE HORIZON

§9 Outreach Execution Engine — channel strategy, V2 sequence design, deliverability rules — bridged to INT-EML/INT-LIN and §7 v2's Portfolio Role + Trust-gated CTA discipline. §9 inherits §8 v2's Trust Preservation lens: every send either builds or spends the trust §8 v2 just made explicit.