ABM BUSINESS LOGIC BIBLE

§9 v2 — DYNAMIC RELATIONSHIP DEPLOYMENT SYSTEM

10 Iterations · 10 Architectural Layers · 27 New Rules · 4 Integration Decisions · 1 New OPEN-Q

§9 v1 specified channel strategy, sequence design, and deliverability. §9 v2 reframes the engine itself. **§9 is not a message-sending module. It is a Dynamic Relationship Deployment System operating in a competitive, adaptive environment.** Its purpose is not to maximize replies or meetings, but to allocate trust, attention, reciprocity, and relationship capital across stakeholder networks — accounting for executive behavior, committee dynamics, competitor reactions, portfolio constraints, and long-term reputation effects. Every interaction is an investment whose value compounds over time, not an isolated attempt to generate an immediate response.

Iterations	10	**Layers**	10
New rules	27	**Integration decisions**	4 explicit

FOUR INTEGRATION DECISIONS MADE BEFORE BUILDING

Multi-Threaded Outreach vs the full-pause rule: genuinely unresolved — flagged as OPEN-Q9.3, not decided here.

Competitive Awareness Probability stays sourced-only: bound strictly to SIG-VENDOR/SIG-NEWS signals, never speculation about an incumbent's internal strategy (mirrors EPIS-ETH-01/02).

Portfolio Optimization reuses §5 v3, not a parallel mechanism: OUT-PORT extends TIER-PORT-001/002 to the sequence-execution layer rather than inventing a second ranking formula.

Influence Signal and Momentum feed Trust Accumulation, not a competing score: one trust ledger with richer inputs (silent engagement, rate-of-change), not multiple parallel trust numbers.

THE TEN LAYERS — §9 v2 ARCHITECTURE MAP

#	Layer	From iter.	What it adds
1	Attention Competition (Attn)	1	The real competitor is existing cognitive load; timing must read board/budget/audit cycles
2	Message Interference (Inter)	2	Channels can reinforce or dilute each other; Sequence Memory tracks what's been seen
3	Response Optionality (Resp)	3	Response ≠ Impact; silent influence signals tracked alongside reply rate
4	Reciprocity Engine (Recip)	4	High-friction asks gated by cumulative Value Contribution, not just trust level
5	Market Memory (Mem)	5	Every interaction leaves residue that propagates through executive mobility
6	Executive Behavior (Beh)	6	Engagement Archetypes — silence means different things for different people
7	Multi-Threaded Outreach (Mthread)	7	Stakeholder State Map per committee member — relationship to full-pause rule is OPEN (Q9.3)
8	Competitive Response (Comp)	8	Outreach changes incumbent behavior too — sourced-only competitive awareness
9	Portfolio Optimization (Port)	9	Sequences compete for scarce resources — extends §5 v3 TIER-PORT, not a new mechanism
10	Relationship Compounding (CompRel)	10	Trust compounds non-linearly; Momentum (rate) matters alongside level

LAYER 1 — ATTENTION COMPETITION (ATTN)

The primary competitor is not another vendor — it's existing cognitive load. A perfect message sent during board week, budget review, or a regulatory audit is ignored not because it's bad but because attention is unavailable. KSA business-hours timing (V2 §17) is necessary but insufficient.

Rule ID	Phase	Specification
OUT-ATTN-01	P2	Every planned send receives an Attention Availability Score derived from: board-cycle proximity, reporting/budget-cycle timing, regulatory-period activity (SIG-REG), major-initiative launches, and executive-travel signals where publicly sourced (conference attendance, public schedule).
OUT-ATTN-02	P2	Send timing optimizes against Attention Availability in addition to KSA business-hours windows (V2 §17-001/002, §9.1 send windows). A touch otherwise due in a low-availability window is held and rescheduled to the next high-availability slot within the sequence's timing tolerance.
OUT-ATTN-03	P1	Attention Availability inputs must be sourced from public signals only (SIG-REG deadlines, SIG-EVENT conference dates, SIG-NEWS earnings dates) — the engine never infers private calendar information about a named individual.

LAYER 2 — MESSAGE INTERFERENCE (INTER)

Rule ID	Phase	Specification
OUT-INTER-01	P1	Every touch reads Sequence Memory — the full record of what this contact has already seen, across BOTH channels — before Agent 5/6 generate the next touch. This extends §7 v2 AI-NARR-01's Active Narrative to the per-contact touch history specifically.
OUT-INTER-02	P1	Every touch receives a Reinforcement Score: does it strengthen the Active Narrative the contact has already been shown, or dilute it? A touch scoring as dilution is rejected and re-queued to Agent 5, using the same hard-gate mechanism as §8 QC-RULE-008 (narrative drift) — message interference is a specific, cross-channel instance of narrative drift.
OUT-INTER-03	P1	Reinforcement scoring applies ACROSS channels: an email's narrative and a LinkedIn touch's narrative for the same contact must score as reinforcing, not merely each channel internally consistent with itself.

LAYER 3 — RESPONSE OPTIONALITY (RESP)

Many successful messages generate no reply. An executive who reads, saves, discusses internally, and returns three weeks later had real influence exerted on them — classifying this as failure (V2's implicit reply-rate framing) is wrong.

Rule ID	Phase	Specification
OUT-RESP-01	P2	Engagement Signal (opens, clicks, replies) is tracked separately from Influence Signal (profile views, website visits, content consumption, colleague forwarding/engagement — where observable via §10 tracking infrastructure).
OUT-RESP-02	P2	Sequence success is measured by Influence Progress, not reply rate alone. A sequence with zero replies but rising Influence Signal is NOT classified as failed; it remains active per its normal cadence rather than being prematurely abandoned.
OUT-RESP-03	P1	Influence Signal feeds the Trust Accumulation Score (§7 v2 AI-TRUST-01) as an additional input category, NOT as a separate competing trust number — one trust ledger, richer inputs (Integration Decision 4).

LAYER 4 — RECIPROCITY ENGINE (RECIP)

Rule ID	Phase	Specification
OUT-RECIP-01	P1	Every touch receives a Value Contribution Score: did this touch give the recipient something useful (an insight, a relevant benchmark, a specific observation) independent of any ask?
OUT-RECIP-02	P1	Before a high-friction ask (direct meeting, pricing discussion), Cumulative Value Contribution across the sequence must exceed a configured threshold. This is the SAME gate as §7 v2 AI-TRUST-02 (Trust Accumulation Score threshold for aggressive CTAs) — Value Contribution is one of the inputs that raises Trust Accumulation, not a separate gate to satisfy independently.

LAYER 5 — MARKET MEMORY (MEM)

What persists after outreach is not emails — it's memories. A good interaction with an SNB executive who later moves to Riyad carries the perception forward. SIG-MOBILITY (§3) already tracks the movement; this layer makes the outreach-specific residue explicit.

Rule ID	Phase	Specification
OUT-MEM-01	P2	Every outreach interaction updates a Relationship Memory Record on the contact: Positive, Neutral, or Negative, derived from engagement quality — the same evidence basis as §7 v2 AI-REP-01 (Reputation Impact Score) applied at the individual-contact level rather than account/market level.
OUT-MEM-02	P2	Relationship Memory propagates when SIG-MOBILITY (§3) detects the contact moving to a new account: the Memory Record transfers with the person, becoming a Trust Transfer Potential input (§7 v2 AI-TRUST-04) at the new account rather than starting cold.
OUT-MEM-03	P1	The engine's account-level success metric explicitly includes Lifetime Relationship Value alongside Sequence Conversion Rate (extends §7 v2 AI-REP-05's dual-metric reporting to the per-sequence level).

LAYER 6 — EXECUTIVE BEHAVIOR (BEH)

Archetype	Behavior pattern	Silence interpretation
Explorer	Engages with new ideas; responds early	Silence = likely disinterest (unusual for this archetype)
Skeptic	Requires evidence; responds late, if at all	Silence = normal; needs more proof points before any reply
Delegator	Rarely responds personally; forwards internally	Silence = normal; success criterion shifts to referral/forward, not direct reply
Protector	Avoids vendor interaction; minimizes risk exposure	Silence = expected default state; low-friction, risk-reduction framing required
Builder	Engages deeply once genuinely interested	Silence = ambiguous; watch Influence Signal (Layer 3) rather than reply alone

Rule ID	Phase	Specification
OUT-BEH-01	P2	Each stakeholder receives an Engagement Archetype (Explorer / Skeptic / Delegator / Protector / Builder) inferred from observable response patterns over the relationship history — NOT from personality inference (EPIS-ETH-01 binds: this is behavioral-pattern classification from actual interaction data, never psychological profiling).
OUT-BEH-02	P2	Sequence interpretation adapts by archetype: a non-response from a Delegator is NOT treated as a negative signal the way a non-response from an Explorer is. §10 sentiment/engagement classification reads the archetype before assigning a positive/neutral/negative interpretation to silence.
OUT-BEH-03	P2	Optimal cadence varies by archetype: Skeptics receive more proof-heavy touches spaced further apart; Explorers receive earlier, more direct touches. Agent 5's sequence-level strategy brief reads the archetype as an input.

LAYER 7 — MULTI-THREADED OUTREACH (MTHREAD)

Accounts don't progress linearly through one champion — they progress through stakeholder networks. CTO engaged, Compliance silent, Procurement resistant are three simultaneous, interacting conversations. V2 itself flags full-account-pause-on-reply as ‘correct for MVP’ but anticipates richer modeling later.

Rule ID	Phase	Specification
OUT-MTHREAD-01	P2	A Stakeholder State Map is maintained per active committee member: unaware / aware / interested / supportive / resistant, derived from each contact's individual engagement history.
OUT-MTHREAD-02	P2	Sequence-level strategy decisions (Agent 5) read the Stakeholder State Map for the whole committee, not just the individual contact's state, when generating any single contact's next touch.
OUT-MTHREAD-03	P2	Account readiness for escalation (§11 Human Handoff) depends on collective stakeholder alignment across the State Map, not single-champion engagement alone.

OPEN-DESIGN-QUESTION OPEN-Q9.3 — MULTI-THREADED OUTREACH vs THE FULL-PAUSE RULE

Question: OUT-SEQ-03 (§9 v1) requires that a reply on ANY contact immediately pauses ALL sequences for the ENTIRE account. OUT-MTHREAD-01..03 implies continued, differentiated outreach to OTHER committee members even while one contact is engaged. These are in direct tension. How should they relate?

Candidate options: (a) Stakeholder State Map ships fully specified but stays DORMANT — tracked for visibility and future analytics, but OUT-SEQ-03's full pause remains the only behavior that acts on it, until a deliberate governance decision (A4) lifts the pause rule for specific account tiers. (b) Observational-only at Phase 1 — track state, but the engine never branches behavior on it even internally; pure telemetry. (c) Skip entirely until MVP graduates, build nothing until this is revisited.

Status: Genuinely unresolved — not defaulted. This question requires a deliberate product decision about whether multi-threaded differentiated outreach is even a desired MVP-adjacent behavior, not just an engineering default. Recommended path if forced to pick: (a), since it lets the data accumulate without changing live behavior, but this is flagged for explicit decision, not silently assumed.

LAYER 8 — COMPETITIVE RESPONSE (COMP)

Outreach changes competitive behavior too: when an incumbent senses an evaluation is happening, it may discount pricing or feed a counter-narrative to the executive. This must be modeled from sourced signals only — never speculation about an incumbent's internal strategy.

Rule ID	Phase	Specification
OUT-COMP-01	**P1**	Every opportunity tracks a Competitive Awareness Probability derived STRICTLY from sourced signals: SIG-VENDOR contract/pricing-activity changes, SIG-NEWS coverage of the incumbent, and observable account behavior (e.g. a sudden incumbent renewal announcement). The engine NEVER infers awareness from unsourced speculation about the incumbent's internal knowledge or strategy (mirrors EPIS-ETH-01/02 + §7 v2 AI-POL-01's sourced-only discipline).
OUT-COMP-02	**P1**	When incumbent entrenchment is high (per §6 v3 SIG-VSAT displacement-probability inputs) AND Competitive Awareness Probability is elevated, Agent 5's strategy shifts toward more proof points and differentiation, and away from direct incumbent-attack framing.
OUT-COMP-03	**P1**	Major sourced competitor moves (SIG-VENDOR contract events, pricing news) become direct inputs to the active sequence's strategy brief, not merely account-dossier updates consumed at the next enrichment cycle — a live sequence can react within its current cadence.

LAYER 9 — PORTFOLIO OPTIMIZATION (PORT)

Rule ID	Phase	Specification
OUT-PORT-01	**P1**	Every active sequence inherits a Portfolio Priority computed via the SAME formula as §5 v3 TIER-PORT-002 (Expected Opportunity Value ÷ Expected Resource Cost) applied to the sequence's specific remaining-touch cost, NOT a new parallel scoring mechanism (Integration Decision 3).
OUT-PORT-02	**P1**	Execution resources (the shared daily send-budget per OUT-SEQ-07, SDR follow-up time, founder attention) are allocated portfolio-wide using Portfolio Priority — resolving §9 v1's OPEN-Q9.2 directly via this rule rather than a separate allocation scheme.
OUT-PORT-03	**P1**	When execution capacity is constrained, the highest Portfolio Priority sequence is served first, NOT the oldest or the one closest to its next scheduled touch — consistent with §5 v3's displacement discipline (TIER-PORT-003): a lower-priority sequence delayed is logged as a displaced opportunity.

LAYER 10 — RELATIONSHIP COMPOUNDING (COMPREL)

Trust doesn't grow linearly. Interaction 1 (minimal trust) → Interaction 4 (trusted source) is not four equal steps; the rate of change matters as much as the current level. A contact at Trust=50 rising rapidly deserves continued investment differently than a contact at Trust=60 that's declining.

Rule ID	Phase	Specification
OUT-COMPREL-01	**P2**	Every contact maintains a Relationship Momentum value — the rate of change of Trust Accumulation Score over the recent interaction window — stored alongside (not instead of) the Trust Accumulation Score itself (Integration Decision 4: one ledger, an added rate-of-change field).
OUT-COMPREL-02	**P2**	Sequence continuation decisions read BOTH level and momentum: a moderate trust level with rising momentum favors continued investment; a higher trust level with declining momentum triggers a pause-and-reassess (this generalizes §9 v1's OUT-SEQ-08, which already pauses on a Trust Accumulation drop — momentum makes the trend explicit rather than reactive only to an already-realized drop).
OUT-COMPREL-03	**P1**	The engine's optimization objective for sequence-level decisions is Lifetime Relationship Value, consistent with §7 v2 AI-REP-03's account-level Expected Reputation Value objective — extended down to the individual contact/sequence level so the two objectives are the same value function at different granularities, not two different goals.

§9 v2 RULE ROLL-UP & PLACEMENT STATUS

Layer	Rules	Phase	Cross-section binding
1 Attn	OUT-ATTN-01..03	P1 / P2	SIG-REG/EVENT/NEWS sourced timing signals
2 Inter	OUT-INTER-01..03	P1	§7 v2 AI-NARR-01/03, §8 QC-RULE-008 hard-gate mechanism
3 Resp	OUT-RESP-01..03	P1 / P2	§10 tracking infra, §7 v2 AI-TRUST-01 (one ledger)
4 Recip	OUT-RECIP-01..02	P1	§7 v2 AI-TRUST-02 (same gate, not a parallel one)
5 Mem	OUT-MEM-01..03	P1 / P2	SIG-MOBILITY, §7 v2 AI-REP-01/05, AI-TRUST-04
6 Beh	OUT-BEH-01..03	P2	EPIS-ETH-01 (behavioral pattern, not personality), §10 sentiment classification
7 Mthread	OUT-MTHREAD-01..03	P2 (status: OPEN)	OUT-SEQ-03 tension — OPEN-Q9.3, §11 escalation readiness
8 Comp	OUT-COMP-01..03	P1	SIG-VENDOR/NEWS sourced-only, §6 v3 SIG-VSAT, EPIS-ETH-01/02
9 Port	OUT-PORT-01..03	P1	§5 v3 TIER-PORT-001/002/003 (same mechanism), resolves §9 v1 OPEN-Q9.2
10 CompRel	OUT-COMPREL-01..03	P1 / P2	§7 v2 AI-TRUST/AI-REP-03, §9 v1 OUT-SEQ-08 (generalizes it)

§9 v2 COMPLETE — WHAT IT ESTABLISHED

§9 is no longer a message-sending module. It is a **Dynamic Relationship Deployment System** operating in a competitive, adaptive environment: it reads attention availability beyond business hours, prevents cross-channel narrative interference, tracks influence distinct from reply, gates high-friction asks on earned reciprocity, propagates relationship memory through executive mobility, adapts to behavioral archetypes, tracks (but does not yet act differently on) stakeholder-network state, watches sourced competitive signals, allocates execution resources portfolio-wide using the same mechanism §5 v3 already established, and optimizes for compounding relationship value rather than campaign conversion. 27 new rules across 10 layers, with one genuinely unresolved design question.

Four integration decisions kept this from becoming five redundant subsystems: sourced-only competitive inference, reuse of §5 v3's portfolio mechanism, one trust ledger instead of parallel scores, and explicit OPEN flagging of the multi-threaded/full-pause tension rather than a silent override.

PLACEMENT MAP UPDATE

§9 status: **Complete-with-Opens** (OPEN-Q9.1, OPEN-Q9.2 from §9 v1 — Q9.2 is now effectively resolved by OUT-PORT-02 and can be closed — plus the new OPEN-Q9.3 on multi-threaded outreach, which is a genuine product decision, not an engineering default).

ON THE HORIZON

§10 Engagement Tracking Engine — engagement events, sentiment classification, and engagement scoring — now also the home for §9 v2's Influence Signal tracking (Layer 3), Engagement Archetype inference (Layer 6), and the observational Stakeholder State Map (Layer 7) pending OPEN-Q9.3's resolution.