

ABM BUSINESS LOGIC BIBLE

§10 — ENGAGEMENT TRACKING ENGINE (Module 8)

V2 Section 10 · Where the World Talks Back · Events + Sentiment + Scoring + Influence

§9 sent. §10 listens. Every open, click, reply, LinkedIn accept, and meeting booking is captured, classified, and used to trigger downstream workflows. This is the first module where the engine stops acting and starts learning from reality's response. §10 also now carries three inheritances from §9 v2: Influence Signal tracking (silent engagement distinct from reply), Engagement Archetype inference (silence means different things for different people), and the observational Stakeholder State Map (currently dormant pending OPEN-Q9.3).

Event types	8 (V2) + Influence Signal	5-Whys chains	5 (root-caused)
New rules	31	Open questions	2 flagged

WHY §10 IS THE ENGINE'S ONLY SENSE ORGAN FOR THE TRUTH

Every rule from §1 through §9 is a hypothesis about what will work. §10 is where those hypotheses meet reality. A scoring weight, a strategy angle, a sequence cadence — all of it is provisional until §10 reports back what actually happened. Get §10 wrong (misclassify sentiment, miss influence signals, conflate silence types) and every upstream learning loop (§13 Analytics, §8 v2 Trust Preservation, §9 v2 Relationship Compounding) learns from corrupted data.

10.0 ENGAGEMENT EVENTS & CLASSIFICATION

Event	Source	Intent level	Triggers
Email Open	Email tool webhook	LOW — awareness only	Update engagement score; log; no handoff
Email Click	Email tool webhook	MEDIUM — active interest	Update score; flag contact 'interested'
Email Reply (positive)	Email + sentiment analysis	HIGH — active engagement	PAUSE all sequences for account; trigger human handoff; create CRM task
Email Reply (negative/unsub)	Email + sentiment analysis	DISQUALIFICATION	STOP all sequences for this contact permanently
Email Reply (neutral/auto)	Email + sentiment analysis	NONE — informational	Log only; continue sequence as planned
LinkedIn Connection Accepted	LinkedIn webhook	MEDIUM — social engagement	Trigger LinkedIn DM (Step 2); update score
LinkedIn Connection Ignored (7+d)	Timer	LOW — no interest	Mark 'connection_pending'; stop LinkedIn cycle
LinkedIn Message Reply	LinkedIn webhook	HIGH — active engagement	PAUSE all sequences for account; trigger handoff
Meeting Booked	Calendar/CRM	CRITICAL — conversion	account_state → MEETING_BOOKED; AI meeting brief; notify sales

Rule ID	Specification
ENG-EVT-01	All nine V2 event types and their triggers transcribed as specified (V2 §10.1) bind exactly as listed above; this table is the canonical event→trigger map referenced by EXEC Stage 10.
ENG-EVT-02	Every engagement event is captured with: event_type, contact_id, account_id, channel, timestamp, raw_payload (verbatim from the source webhook). Append-only — never overwritten, matching the §2.3 SIGNAL entity discipline.

5-WHYS — WHY THIS ENGAGEMENT ARCHITECTURE

Why #1: Why does ANY positive reply pause the ENTIRE account, not just that contact's sequence?

Because the engine is account-centric (V2 Principle 2): if Contact A replies positively, hitting Contact B simultaneously with automated outreach looks uncoordinated to an organization where people talk to each other. The full-account pause is the conservative default that protects the relationship as a whole — exactly the tension §9 v2 OPEN-Q9.3 flagged as still requiring a deliberate decision before any differentiated behavior replaces it.

Why #2: Why must sentiment classification have a confidence threshold (0.7) with human escalation below it?

Because misclassifying a reply's sentiment is worse than not classifying it at all: a negative reply misclassified as neutral continues an unwanted sequence (relationship damage); a positive reply misclassified as neutral misses a hot opportunity (opportunity cost). Both failure directions are expensive, and AI confidence below 0.7 isn't reliable enough to risk either — human classification is cheap insurance against the costlier error.

Why #3: Why distinguish a referral reply as its own category rather than folding it into 'positive'?

Because a referral changes who the next action targets: it isn't 'continue this conversation,' it's 'create a new contact and start fresh with a different, warmer context (the introduction).' Treating it as generic positive sentiment would miss the chance to capture the referred contact and properly attribute the introduction's trust transfer (§7 v2 AI-TRUST-04).

Why #4: Why cap email-open scoring at 3 opens rather than scoring every open?

Because opens past the third are overwhelmingly an artifact of email client preview panes and mobile notification previews re-triggering open pixels, not genuine re-engagement. Uncapped open-scoring would let an email client's polling behavior inflate engagement score with zero behavioral signal behind it — exactly the kind of measurement error §6 v3 Layer 4 (Information Quality vs Quantity) warns against.

Why #5: Why must Influence Signal (§9 v2 Layer 3) be tracked here rather than only at §13 Analytics?

Because §10 is the only module positioned to observe the raw events (website visits, content consumption, profile views) at the moment they happen; §13 only sees aggregated rollups. If Influence Signal isn't captured at §10, the granular, time-stamped evidence needed to feed Trust Accumulation (§7 v2 AI-TRUST-01) and to correctly interpret silence per Engagement Archetype (§9 v2 OUT-BEH-02) is permanently lost.

Root cause: engagement tracking must be account-centric in its pause behavior, confidence-gated in its sentiment classification, type-differentiated for referrals, artifact-resistant in its scoring, and the capture point — not just the rollout point — for every signal that downstream trust and archetype models depend on.

10.1 SENTIMENT CLASSIFICATION RULES

Rule ID	Specification
ENG-SENT-01	POSITIVE indicators: 'interested', 'let's talk', 'send me more', 'schedule', 'meeting', 'call', 'demo', 'tell me more', a question about Decimal's capabilities, forwarding to a colleague. Source: V2 ENG-SENT-001.
ENG-SENT-02	NEGATIVE indicators: 'not interested', 'unsubscribe', 'remove me', 'stop emailing', 'we already have a vendor', 'no budget', 'bad timing', explicit rejection. Source: V2 ENG-SENT-002.
ENG-SENT-03	NEUTRAL indicators: out-of-office auto-reply, 'let me check internally', 'not the right person — try [name]', informational response without commitment. Source: V2 ENG-SENT-003.
ENG-SENT-04	REFERRAL reply ('you should talk to [name]') is POSITIVE and handled distinctly: capture the referred contact's name, create a new contact record, generate personalized outreach for the referral mentioning the introduction. Source: V2 ENG-SENT-004.
ENG-SENT-05	AI classifies sentiment first. If confidence < 0.7, escalate to human classification; human classification overrides AI. Source: V2 ENG-SENT-005.
ENG-SENT-06	Sentiment classification reads the contact's Engagement Archetype (§9 v2 OUT-BEH-01) before finalizing interpretation of an ambiguous or absent reply: a non-response from a Delegator is classified differently than the identical non-response from an Explorer (§9 v2 OUT-BEH-02 binds here operationally).
ENG-SENT-07	A referral's introducer receives a Trust Transfer Potential update (§7 v2 AI-TRUST-04): the act of referring is itself evidence of relationship strength, independent of the referred contact's eventual response.

10.2 ENGAGEMENT SCORING

Event	Points	Notes
Email Open	+2 per open (max 3 counted)	Additional opens beyond 3 add 0 — prevents preview-pane inflation
Email Click	+5 per click	Each unique link click counted once
Email Positive Reply	+25	Counted once per sequence
LinkedIn Accept	+10	Once per contact
LinkedIn Reply	+25	Counted once
Meeting Booked	+50	The ultimate conversion event
Email Negative Reply	−30	Penalizes contact and account engagement score

Rule ID	Specification
ENG-SCORE-01	Point values and caps as listed above (V2 §10.3). Account engagement score = sum of all contact engagement scores at that account (V2 ENG-SCORE-001).
ENG-SCORE-02	Engagement score feeds back into §13 Loop 1 (Signal Quality Feedback) and §4 SCORE-RECAL (weight recalibration) as one of the correlation inputs against actual meetings booked.
ENG-SCORE-03	Engagement score is a SEPARATE field from Trust Accumulation Score (§7 v2 AI-TRUST-01) and Influence Signal (§9 v2 Layer 3): engagement score is point-based and event-driven; Trust Accumulation is a calibrated 0–100 confidence measure; Influence Signal is evidence of impact without direct engagement. All three are visible on the contact record without being collapsed into one number.

10.3 INFLUENCE SIGNAL TRACKING (from §9 v2 Layer 3)

Many successful messages generate no reply. §10 is where the silent evidence of impact — profile views, website visits, content consumption, colleague forwarding — must actually be captured, since §10 is the only module positioned to see these raw events as they happen.

Rule ID	Specification
ENG-INFL-01	§10 captures Influence Signal events where observable: LinkedIn profile views of Decimal team members by the contact, website visits matched to the contact (via tracked links in outreach), content downloads/views, and colleague engagement (a forwarded email opened by a new, previously-uncontacted person at the same account).
ENG-INFL-02	Each Influence Signal event is logged with the same rigor as an Engagement event (ENG-EVT-02): type, contact_id, account_id, timestamp, source. Influence Signal events do NOT pause sequences (only HIGH-intent Engagement events do, per ENG-EVT-01) — they are evidence, not triggers.
ENG-INFL-03	Influence Signal feeds Trust Accumulation Score (§7 v2 AI-TRUST-01) as an additional input category, consistent with §9 v2 Integration Decision 4 (one trust ledger, richer inputs) — it does not create a parallel score.
ENG-INFL-04	A sequence with zero replies but positive, accumulating Influence Signal is explicitly NOT classified as a failed sequence (§9 v2 OUT-RESP-02) and is not prematurely halted by §9's standard cadence-completion logic.

10.4 ENGAGEMENT ARCHETYPE INFERENCE (from §9 v2 Layer 6)

Rule ID	Specification
ENG-ARCH-01	§10 is the inference point for Engagement Archetype (Explorer / Skeptic / Delegator / Protector / Builder, §9 v2 OUT-BEH-01): the archetype is derived from the accumulated pattern of engagement events on this contact (response latency, response channel preference, forwarding behavior, reply length/depth) over the relationship history, not assigned from a single event.
ENG-ARCH-02	Archetype inference requires a minimum evidence threshold (at least 2 prior touches with some engagement signal) before an archetype is assigned; before that threshold, the contact is UNCLASSIFIED and sentiment interpretation defaults to the standard rules (10.1) without archetype adjustment.
ENG-ARCH-03	Archetype classification is re-evaluated on every new engagement event — it is not fixed at first inference. A contact behaving as a Skeptic early (slow, evidence-seeking) who later engages deeply is reclassified toward Builder, and sentiment interpretation adjusts accordingly.
ENG-ARCH-04	Archetype inference is bounded by EPIS-ETH-01: it classifies OBSERVABLE BEHAVIORAL PATTERNS from actual interaction data, never personality or psychological traits. The archetype label describes a response pattern, not a person.

10.5 STAKEHOLDER STATE MAP — OBSERVATIONAL ONLY (from §9 v2 Layer 7, pending OPEN-Q9.3)

STATUS: DORMANT / TELEMETRY-ONLY PENDING OPEN-Q9.3

Per §9 v2's explicit flag, the Stakeholder State Map's relationship to the full-account-pause rule (OUT-SEQ-03) is an unresolved product decision, not an engineering default. §10 implements ONLY the observational half: it tracks and stores per-contact state, but does not yet branch any sequence behavior on the committee-level view. This keeps the data accumulating for the eventual decision without quietly overriding the existing safety-first pause rule.

Rule ID	Specification
ENG-STATE-01	§10 updates each active committee member's Stakeholder State (unaware / aware / interested / supportive / resistant, §9 v2 OUT-MTHREAD-01) on every engagement event for that contact, using the same event→state inference logic as sentiment classification (10.1).
ENG-STATE-02	The Stakeholder State Map is visible in the account dossier and consumed by §11 Human Handoff (as context for the handoff package) and §13 Analytics (Committee Influence Progress metric, §7 v2 AI-INF-04) — but is NOT read by §9's sequence-execution logic for any contact other than the one whose state just changed, pending OPEN-Q9.3.
ENG-STATE-02B	Engagement events on Contact A do not alter the sequence behavior for Contact B at §10/§9, even though both contacts' states are visible on the same map — the map is read-only context for humans and analytics until OPEN-Q9.3 is resolved.

10.6 OPEN-DESIGN-QUESTIONS — §10

OPEN-Q10.1 — INFLUENCE SIGNAL OBSERVABILITY LIMITS

Question: ENG-INFL-01 lists LinkedIn profile views, website visits, content engagement, and colleague forwarding as Influence Signal sources. Several of these (especially LinkedIn profile views of Decimal team members, and matching an anonymous website visit to a specific contact) depend on tooling/tracking infrastructure that may not exist at MVP. Which sources are realistically available at launch vs. require new tracking investment?

Candidate options: (a) Ship only what's free/already available (e.g. LinkedIn accept + reply, which are Engagement not Influence signals) and treat true Influence Signal as fully Phase 2. (b) Invest in basic tracked-link website-visit matching at MVP (low-cost, common martech pattern) while deferring LinkedIn profile-view and colleague-forwarding detection. (c) Build all four sources at MVP.

Safe default: (b). Tracked-link website-visit matching is a well-understood, low-cost addition; the other two sources are deferred to Phase 2 without blocking the overall Influence Signal concept from existing in the dossier (it simply has fewer populated inputs at first).

OPEN-Q10.2 — ARCHETYPE RECLASSIFICATION SENSITIVITY

Question: ENG-ARCH-03 reclassifies archetype on every new engagement event. How sensitive should reclassification be — does ONE atypical event (e.g. a Skeptic's single fast reply) flip the archetype, or does it require a sustained pattern shift?

Candidate options: (a) Any new event recomputes the archetype from the full history each time — simple but may flicker on noisy single events. (b) Archetype changes require 2+ consecutive events consistent with the new pattern (mirrors §9 v1's OPEN-Q9.1 sustained-decline logic). (c) Archetype is weighted/probabilistic (e.g. 70% Skeptic / 30% Builder) rather than a single discrete label, avoiding hard flips entirely.

Safe default: (b). Consistent with the sustained-pattern discipline already chosen for §9's trust-drop pause logic; keeps archetype behavior predictable rather than flickering on noise.

§10 RULE ROLL-UP & CROSS-LINKS

Group	Rules	Key cross-links
Events + classification	ENG-EVT-01..02	EXEC Stage 10, §2.3 append-only discipline
Sentiment	ENG-SENT-01..07	§9 v2 OUT-BEH-02 archetype-aware interpretation, §7 v2 AI-TRUST-04 referral
Engagement scoring	ENG-SCORE-01..03	§13 Loop 1, §4 SCORE-RECAL, kept distinct from Trust/Influence
Influence Signal	ENG-INFL-01..04	§9 v2 Layer 3, §7 v2 AI-TRUST-01 (one ledger)
Archetype inference	ENG-ARCH-01..04	§9 v2 Layer 6, EPIS-ETH-01
Stakeholder State Map	ENG-STATE-01..02B	§9 v2 OPEN-Q9.3 (dormant pending resolution), §11 handoff context
Open questions	2 flagged	A4 governance per TIER-OPEN-DISCIPLINE-02

§10 COMPLETE — WHAT IT ESTABLISHED

Engagement Tracking is now fully specified at V2's operational level and is the capture point — not merely the rollup point — for every signal that §7 v2's Trust Accumulation, §9 v2's Engagement Archetypes, and §8 v2's Trust Preservation depend on. The nine V2 event types, the five sentiment categories (including the distinct referral path), and the seven-row scoring table are transcribed in full. Three inheritances from §9 v2 are operationalized: Influence Signal capture, Archetype inference with a minimum-evidence floor and continuous reclassification, and the Stakeholder State Map shipped strictly as observational telemetry — visible to humans and analytics but not yet wired into sequence behavior, honoring OPEN-Q9.3's unresolved status rather than quietly deciding it. 31 new rules; 2 new open questions on observability limits and reclassification sensitivity.

§10 is the engine's only sense organ for ground truth. Everything it misclassifies corrupts every downstream learning loop.

PLACEMENT MAP UPDATE

§10 status: **Complete-with-Opens** (OPEN-Q10.1 Influence Signal observability limits, OPEN-Q10.2 archetype reclassification sensitivity). §9's OPEN-Q9.3 remains open and is now also a dependency of §10's Stakeholder State Map scope.

ON THE HORIZON

§11 Human Handoff Engine — handoff trigger conditions, the handoff package (account summary, signal timeline, engagement history, persona profile, pain inference, suggested angles), and the workflow steps — now enriched by §10's Stakeholder State Map as handoff context and the full Trust/Influence/Archetype picture per contact.