

ABM BUSINESS LOGIC BIBLE

§11 — HUMAN HANDOFF ENGINE (Module 9)

V2 Section 11 · Where the Engine Hands Reality to a Person · Triggers + Package + Workflow

§10 confirmed that reality voted positively. §11 is where the engine steps back and a human steps in. This handoff must be seamless — the human must receive complete context so the prospect never feels like they're starting over. V2 specifies trigger conditions, the handoff package contents, and the workflow steps. §11 transcribes them operationally and substantially enriches the package with everything built since §7: GOV's autonomy tiers, the Trust/Influence/Archetype picture per contact, the (dormant) Stakeholder Truth Map, and §10 v2's Cause Attribution and Learning Artifacts.

V2 rule groups	3	5-Whys chains	5 (root-caused)
New rules	30	Open questions	2 flagged

WHY §11 IS WHERE TRUST IS REALIZED OR LOST IN ONE CONVERSATION

Every architectural discipline in this Bible exists to produce one outcome well: a human picks up a relationship the engine has nurtured and continues it without friction. If the handoff package is incomplete, the human re-asks questions the prospect already answered, or pitches an angle the engine already learned doesn't work — spending in one conversation the trust that §7 v2, §8 v2, and §10 v2 spent dozens of rules protecting. §11 is the conversion point where all that accumulated discipline either pays off or is wasted.

11.0 HANDOFF TRIGGER CONDITIONS

Trigger	Priority	Assigned to	Response SLA
Positive email reply	URGENT	Account owner (SDR)	Respond within 2 hours, business hours
Meeting request	CRITICAL	Account owner (AE)	Confirm within 1 hour
Pricing question	HIGH	Account Executive	Respond within 4 hours
Security / compliance question	HIGH	Technical team + AE	Respond within 24 hours (needs research)
Referral to colleague	MEDIUM	Account owner (SDR)	New outreach within 24 hours
LinkedIn positive reply	URGENT	Account owner	Respond within 2 hours, business hours

Rule ID	Specification
HAND-TRIG-01	Trigger conditions, priority, assignment, and SLA exactly as specified in V2 §11.1. Source events are §10 v1 ENG-EVT-01's HIGH-intent and CRITICAL-intent rows.
HAND-TRIG-02	Trigger priority is adjusted by §10 v2 ENG-EPIST-01 Intent Confidence: a positive reply with low Intent Confidence (e.g. AI sentiment classification below the 0.7 threshold per §10 v1 ENG-SENT-05) is held for human sentiment confirmation BEFORE the handoff fires, rather than firing an URGENT handoff on an unconfirmed read.
HAND-TRIG-03	A Security/Compliance question is automatically cross-referenced against §7 v2 AI-POL-01 (Political Sensitivity Rating): if the question relates to a topic already flagged Medium/High sensitivity, the Technical team + AE assignment is supplemented with an explicit sensitivity flag in the handoff package, since the 24-hour research window must account for that context.

11.1 THE HANDOFF PACKAGE — V2 BASELINE

V2 specifies nine package components. Each is now sourced from a specific, named upstream module rather than assumed to simply 'exist' — making the package's completeness auditable.

Package component	V2 description	Sourced from
Account Summary	AI-generated brief: organization, priorities, digital maturity	§6 ENR-SUM-01 (500–800 word grounded summary)
Signal Timeline	All signals, reverse chronological, with interpretations	§3 INT-SIG signal stream + §7 Agent 1 interpretations
Engagement History	Every message sent to every contact, with opens/clicks/replies	§10 v1 ENG-EVT log, full append-only record
Reply Content	The actual reply text that triggered the handoff	§10 v1 ENG-EVT-02 raw_payload
Persona Profile	AI-generated persona analysis for the replying contact	§7 Agent 3 output
Pain Inference	Top 3 inferred pains from the AI reasoning chain	§7 Agent 4 output
Suggested Discussion Angles	AI-generated talking points for the first human conversation	§7 Agent 5 strategy brief
Likely Objections	AI-predicted objections with suggested responses	§7 Agent 4 + Agent 5 combined
Comparable Wins	Relevant Decimal case studies or client wins	§6 v3 ENR-PORT-03 Shared Intelligence Assets

Rule ID	Specification
HAND-PKG-01	All nine V2 package components are mandatory; a handoff package missing any component is incomplete and the handoff is flagged for completion before the SLA clock is treated as satisfied (the notification still fires immediately per HAND-WF below, but completeness is tracked separately).
HAND-PKG-02	Every package component traces to its named source module (table above); the package generator does not synthesize content outside what those modules have already produced — the handoff package is an assembly and presentation layer, not a new reasoning step, consistent with M0 hallucination governance applying end-to-end.

11.2 THE HANDOFF PACKAGE — EXTENDED COMPONENTS (POST-§7 ARCHITECTURE)

Five additional components reflect everything built since V2 was written. None replace the original nine; all are additive context that a 2024-era V2 package could not have included.

Rule ID	Specification
HAND-PKG-03	Trust & Influence Profile (new): the contact's Trust Accumulation Score, Relationship Momentum (§9 v2 OUT-COMPREL-01), and any logged Influence Signal events (§9 v2 OUT-RESP-01, §10 v1 ENG-INFL) — so the human rep knows whether this reply is the START of trust-building or the CONFIRMATION of trust already earned, and adjusts their opening accordingly.
HAND-PKG-04	Engagement Archetype (new): the replying contact's inferred Archetype (§10 v1 ENG-ARCH) where classified above the minimum-evidence floor, e.g. 'Skeptic — expect to need more proof points before this converts; do not over-pitch in the first call.'
HAND-PKG-05	Stakeholder Truth Map snapshot (new, observational): where §9 v2/§10 v2's Stakeholder Truth Map has data for OTHER committee members at this account (CTO supportive, Compliance resistant, Procurement silent), it is included as human-readable context — this is exactly the use case ENG-STATE-02 (§10 v1) anticipated: visible to humans immediately, even while dormant for engine-side sequence logic pending OPEN-Q9.3.
HAND-PKG-06	Cause Attribution summary (new): §10 v2 ENG-CAUSE-01's distribution for THIS specific engagement outcome — e.g. 'attributed 60% to outreach, 25% to a SIG-REG regulatory deadline, 15% counterfactual baseline' — so the human rep understands what actually likely drove this reply, not just that a reply happened.
HAND-PKG-07	Political Sensitivity flags (new): any Medium/High Political Sensitivity Ratings (§7 v2 AI-POL-01) active on this account's narrative or recent messaging, so the human conversation doesn't innocently reopen a topic the engine had deliberately steered around.

11.3 HANDOFF WORKFLOW STEPS

Rule ID	Specification
HAND-WF-01	Step 1: PAUSE all automated sequences for the entire account (all contacts, all channels), immediate and automatic. Source: V2 HAND-001. This is the same mechanism as §9 v1 OUT-SEQ-03 / §10 v1 ENG-EVT-01 — restated here as the first handoff-workflow step, not a separate pause.
HAND-WF-02	Step 2: account_state → ENGAGED. Source: V2 HAND-002.
HAND-WF-03	Step 3: Generate the full handoff package (§11.1 + §11.2). Source: V2 HAND-003, extended.
HAND-WF-04	Step 4: Slack notification to account_owner: account name, contact who replied, reply text, link to full brief. Source: V2 HAND-004, bound to INT-SLK's grounded-brief contract (§4 v2 INT-SLK-02).
HAND-WF-05	Step 5: Create CRM task in HubSpot: 'Follow up with [CONTACT] at [ACCOUNT] — positive reply received [DATETIME]'. Source: V2 HAND-005, bound to INT-CRM's authority-split discipline (§4 v2 INT-CRM-02): the CRM task is ABM-authored but the deal-stage progression that may follow is CRM-authoritative.
HAND-WF-06	Step 6: start a 48-hour acknowledgment timer. If account_owner does not acknowledge within 48 hours, escalate to sales manager. Source: V2 HAND-006.
HAND-WF-07	Step 6 escalation is a GOV-governed action: an unacknowledged handoff escalating past 48 hours is logged as a Reviewer/Owner Calibration signal (§8 v2 QC-HUMAN-08's per-person calibration tracking, extended from QC reviewers to account owners) — repeated late acknowledgments by the same owner feed the same per-subsystem trust model (EPIS-TRUST-01) as any other degrading subsystem.
HAND-WF-08	Steps 1–6 execute atomically at the orchestration layer (INT-ORC idempotency, EXEC-03): a retry of the handoff workflow does not re-pause an already-paused account, re-create a duplicate CRM task, or re-send a duplicate Slack notification.

5-WHYS — WHY THIS HANDOFF ARCHITECTURE

Why #1: Why must the entire package be assembled from already-existing module outputs rather than generated fresh at handoff time?

Because generating fresh content at handoff time would mean a NEW reasoning step happening under time pressure (the 1–2 hour SLA), with no QC layer reviewing it (§8's gates are upstream, not at §11). Assembling from already-produced, already-QC'd outputs means the handoff package inherits every grounding guarantee the architecture already built, instead of re-exposing the engine to hallucination risk at the worst possible moment — right before a human's first real conversation with the prospect.

Why #2: Why does a low-Intent-Confidence positive reply get held for confirmation rather than triggering the handoff immediately?

Because an URGENT 2-hour SLA commits real human attention and signals real urgency to the account_owner; firing that signal on a misclassified reply (e.g. an out-of-office auto-reply misread as positive) burns the account owner's trust in the handoff system itself — exactly the per-subsystem trust erosion EPIS-TRUST-01 and §8 v2's Trust Preservation Score exist to prevent. A brief confirmation delay is cheap; a false-urgent handoff is not.

Why #3: Why include the Cause Attribution summary in the package when the human rep just wants to talk to an interested prospect?

Because the OPENING of the human conversation differs enormously depending on cause: if a SIG-REG deadline likely drove the reply (not the outreach itself), the rep should lead with regulatory urgency, not with Decimal's specific email. If the reply is heavily counterfactual-weighted (the prospect was probably going to reach out regardless), the rep should not over-credit the outreach narrative in their opening framing. Getting the opening wrong wastes the warmth §10 v2 worked to correctly attribute.

Why #4: Why does an unacknowledged handoff feed a Calibration Score on the ACCOUNT OWNER, the same mechanism used for QC reviewers?

Because §8 v2 already established that human decision-makers in this system are a governed decision system, not an automatic safety net — that discipline doesn't stop at QC reviewers. An account owner who systematically misses the 48-hour acknowledgment window is a degrading subsystem exactly like an AI module with a rising hallucination rate, and the architecture should treat the two symmetrically rather than auditing AI rigorously while assuming human reliability by default.

Why #5: Why does the Stakeholder Truth Map appear in the package as human-readable context when it's explicitly NOT used by the engine's own sequence logic?

Because OPEN-Q9.3 gates whether the ENGINE acts differently per stakeholder — it says nothing about whether a HUMAN should see the same information. A rep walking into a call blind to 'Compliance is resistant' when the engine has already observed that pattern wastes information the architecture worked to capture. Observational data has immediate value to humans even while remaining inert for automated decision-making; withholding it from the package until OPEN-Q9.3 resolves would be over-applying a caution meant for a different layer of the system.

Root cause: the handoff must assemble only already-grounded upstream outputs, confirm ambiguous high-stakes triggers before committing human attention, surface causal context so the human's opening matches reality rather than a generic assumption, hold human decision-makers to the same governed reliability standard as AI subsystems, and make observational intelligence available to humans immediately even where it remains gated from automated action.

11.4 OPEN-DESIGN-QUESTIONS — §11

OPEN-Q11.1 — LOW-INTENT-CONFIDENCE CONFIRMATION MECHANISM

Question: HAND-TRIG-02 holds a low-Intent-Confidence positive reply for confirmation before firing the URGENT handoff. Who performs that confirmation, and how fast — is it the same AI sentiment re-check at higher scrutiny, a lightweight human glance (faster than full QC review), or does it route through §8's existing QC-HUMAN human-review layer entirely?

Candidate options: (a) A second-pass AI re-check with a stricter confidence bar, only escalating to human if still ambiguous. (b) A lightweight, fast human glance distinct from §8's QC review (different purpose: confirm sentiment, not approve a message). (c) Route through the existing §8 QC-HUMAN reviewer queue, treating it as just another review item with its own SLA.

Safe default: (b). A dedicated lightweight confirmation step keeps the URGENT handoff SLA meaningful (a full QC-queue routing could absorb the 2-hour budget before the account owner even sees it) while still adding a human check before committing real urgency signaling.

OPEN-Q11.2 — PACKAGE COMPLETENESS vs SLA TENSION

Question: HAND-PKG-01 separates the notification SLA (fires immediately) from package completeness (tracked separately). What happens operationally if a component is genuinely missing when the SLA clock expires — does the account owner get a partial package and a follow-up, or does the SLA effectively pause until the package is complete?

Candidate options: (a) SLA is strict regardless — notify on time with whatever's ready, backfill missing components asynchronously. (b) SLA pauses for genuinely critical missing components (e.g. no Account Summary at all) but proceeds with cosmetic gaps (e.g. missing Comparable Wins). (c) SLA never pauses; incompleteness is purely a quality metric tracked in §13 Analytics, never a workflow blocker.

Safe default: (a). The 2-hour SLA exists because timing matters competitively (§5 v3 Patience layer, §9 v2 Attention Competition); a notification delayed waiting for a Comparable Wins case study could cost more than the gap itself. Notify on time, backfill, track completeness as a quality metric.

§11 RULE ROLL-UP & CROSS-LINKS

Group	Rules	Key cross-links
Trigger conditions	HAND-TRIG-01..03	§10 v2 ENG-EPIST-01 (Intent Confidence gate), §7 v2 AI-POL-01
Package — V2 baseline	HAND-PKG-01..02	§6 ENR-SUM, §7 Agents 1–5, §10 v1 ENG-EVT, M0 grounding
Package — extended	HAND-PKG-03..07	§7 v2 AI-TRUST/AI-NARR, §9 v2/§10 v2 Stakeholder Truth Map, §10 v2 ENG-CAUSE
Workflow steps	HAND-WF-01..08	§9 v1 OUT-SEQ-03, INT-SLK, INT-CRM-02, EPIS-TRUST-01, EXEC-03 idempotency
Open questions	2 flagged	A4 governance per TIER-OPEN-DISCIPLINE-02

§11 COMPLETE — WHAT IT ESTABLISHED

Human Handoff is now operationally specified at V2's level of concretion and carries the full weight of everything built since §7: the package's original nine V2 components are each traced to a named source module rather than assumed, and five new components — Trust & Influence Profile, Engagement Archetype, Stakeholder Truth Map snapshot, Cause Attribution summary, and Political Sensitivity flags — give the human rep context no 2024-era V2 package could have included. The trigger layer now confirms ambiguous high-stakes signals before committing real urgency, and the workflow holds account owners to the same governed-reliability standard §8 v2 established for QC reviewers. 30 new rules; 2 open questions flagged with safe defaults.

§11 is the conversion point: every rule built across ten and a half sections exists to make this one handoff land well. §12 (CRM Synchronization) is where the conversation's outcome starts flowing back into the system of record the sales team actually lives in.

PLACEMENT MAP UPDATE

§11 status: Complete-with-Opens (OPEN-Q11.1 confirmation mechanism, OPEN-Q11.2 SLA/completeness tension). Both resolve to safe, low-risk defaults that don't block §12.

ON THE HORIZON

§12 CRM Synchronization Engine — push/pull data, sync rules — architecturally already substantially specified by §4 v2's INT-CRM domain (idempotency, authority split, conflict resolution); §12 transcribes V2's field-level detail and reconciles it against INT-CRM.