ABM BUSINESS LOGIC BIBLE

§11 v2 — HUMAN-AI DECISION PARTNERSHIP ENGINE

10 Iterations · 10 Architectural Layers · 27 New Rules · 5 Integration Decisions

§11 v1 specified trigger conditions, the handoff package, and the workflow. §11 v2 reframes the engine itself. **§11 is not a notification workflow. It is the Authority Transfer and Judgment Continuity Engine — and beyond that, the point where the system becomes a Human-AI Decision Partnership, not software with a human fallback.** The handoff exists because uncertainty exceeded automation's authority: the engine is effectively saying 'I have reduced uncertainty as much as I responsibly can; a human must now assume responsibility.' Success is measured not by package delivery but by Combined Decision Quality — and by whether the human can continue the relationship without losing accumulated context, trust, or judgment.

Iterations	10	**Layers**	10
New rules	27	**Integration decisions**	5 explicit

FIVE INTEGRATION DECISIONS — NOT FIVE NEW PARALLEL MECHANISMS

'Why automation stopped' IS GOV's Autonomy Ladder, rendered human-readable. HAND-AUTH reuses the specific GOV-LAD tier/trigger that fired (trust threshold, political sensitivity, compliance complexity are already named GOV-ESC-01 triggers), not a second taxonomy.

Confidence + Contrarian View ARE EPIS-RCM-03's falsifiability requirement, applied to a recommendation a human is about to act on — the Confidence Level is the existing RCM stamp; the Contrarian View is 'what would change our mind,' asked explicitly before the human acts rather than only inside the engine.

Reverse Handoff enters through the EXISTING signal pipeline: human-discovered context becomes a MANUAL-source signal (V2 §3.1 already lists Manual Input as a signal source) with a full RCM stamp — not a new ingestion path.

Override capture and Decision Memory are ONE Decision Record artifact, not two mechanisms — logged on override, retained as searchable Tier 6 Strategic Memory.

Expertise Allocation extends §5 v3 TIER-ROUTE (the existing Human Expertise Graph) to support multiple simultaneous Support Roles, not a second expertise model.

THE TEN LAYERS — §11 v2 ARCHITECTURE MAP

#	Layer	From iter.	What it adds
1	Context Compression (Comp)	1	60-second executive summary; rank by Decision Importance, not richness
2	Judgment Transfer (Judge)	2	Transfer the reasoning chain and rejected alternatives, not just the conclusion
3	Override & Decision Memory (Learn/Mem)	3, 9	One Decision Record artifact — captured on override, retained as Tier 6 memory
4	Accountability Transfer (Acc)	4	Explicit Decision Ownership Boundary: AI Recommendation / Human Decision / Human Execution
5	Authority Transition (Auth)	5	Why automation stopped, rendered from GOV's existing tier/trigger taxonomy
6	Automation Bias (Bias)	6	Contrarian View required — EPIS-RCM-03 applied before the human acts, not just inside the engine
7	Reverse Handoff (Know)	7	Human-discovered context re-enters as a MANUAL-source signal with full RCM stamp
8	Expertise Allocation (Exp)	8	Multi-role Support Team extends §5 v3's existing Expertise Graph
9	(merged into Layer 3)	9	—
10	Cognitive Division of Labor (CDL)	10	Per-decision-type AI-led/Human-led/Joint classification, one level above GOV's per-action tiers

LAYER 1 — CONTEXT COMPRESSION (COMP)

By §11, the account dossier has accumulated signals, scoring, enrichment, trust, influence, archetypes, narratives, and stakeholder states. A human has limited time, attention, and working memory. A 20-page package gets 2 pages read. The challenge isn't information transfer — it's compression.

Rule ID	Phase	Specification
HAND-COMP-01	**P1**	Every handoff package leads with an Executive Summary Layer: a maximum 60-second read (~150 words) that sits ABOVE the full §11 v1 nine-component package, not instead of it. The full package remains available for the rep who wants depth; the summary is what most reps will actually read first.
HAND-COMP-02	**P1**	Every insight in the package is ranked by Decision Importance (does this change what the rep should say or do in the next 10 minutes?), not by informational richness or recency. The richest, most recent enrichment fact does not automatically rank above an older, less detailed fact if the older fact has higher decision relevance — reuses §6 v3 Layer 4's Information Quality vs Quantity discipline (ENR-QUAL-01), applied to package ordering rather than enrichment acquisition.
HAND-COMP-03	**P1**	The Executive Summary explicitly answers: 'if the rep remembers only three things, what are they?' — capped at exactly three items, forcing real prioritization rather than a longer 'top 5' that defeats the compression purpose.

LAYER 2 — JUDGMENT TRANSFER (JUDGE)

Rule ID	Phase	Specification
HAND-JUDGE-01	**P1**	Every suggested discussion angle (§11 v1 HAND-PKG component, sourced from §7 Agent 5) includes its Reasoning Chain in the package: not 'lead with regulatory urgency' alone, but 'lead with regulatory urgency because SIG-REG contributed the largest share of this account's §10 v2 ENG-CAUSE attribution for the triggering engagement.' The reasoning, not just the conclusion, transfers.
HAND-JUDGE-02	**P1**	Every strategy recommendation in the package includes its Rejected Alternatives — the other angles Agent 5 considered and why each scored lower — so the human understands the tradeoff space, not just the winning choice. This is the package-facing view of the same comparison Agent 5 already performs internally; no new reasoning step, only added transparency.

LAYER 3 — OVERRIDE & DECISION MEMORY (LEARN / MEM)

When a rep ignores the engine's CTO-route suggestion and talks to Procurement instead, that deviation is training data — currently lost. And whether the rep followed the suggestion or not, the rationale for what was actually decided is exactly what organizations forget fastest. One Decision Record artifact serves both needs.

Rule ID	Phase	Specification
HAND-DEC-01	P1	Every §11 handoff that reaches a material next-action decision (which stakeholder to approach, which angle to use, whether to escalate) generates a Decision Record: options considered, selected option, rationale, assumptions, and — if the human's choice differs from the package recommendation — an explicit Override Reason selected from a defined taxonomy (local knowledge, relationship context, political consideration, timing concern, other-with-required-free-text).
HAND-DEC-02	P1	Decision Records are submitted to Tier 6 Strategic Memory (§2, specifically CounterfactualRecord and OrganizationalDNA) as §10 v2 ENG-MEM-01 already established for engagement Learning Artifacts — ONE memory pipeline, with Decision Records as the §11-specific artifact type flowing into it, not a second memory system.
HAND-DEC-03	P2	Override Reasons aggregated across many Decision Records become a feedback signal for §7 v2 Agent 5 strategy calibration: a recommendation pattern that gets overridden for 'relationship context' reasons repeatedly at a specific persona/segment combination indicates Agent 5's strategy model is under-weighting a factor the human reps consistently see.
HAND-DEC-04	P1	Future reps and analysts can query Decision Records for a given account or persona type to see WHY prior decisions were made, not only what happened — institutional memory stores reasoning, retrievable independent of whether the original rep is still at the company.

LAYER 4 — ACCOUNTABILITY TRANSFER (ACC)

Engine recommends, human follows the recommendation, deal is lost. Who was wrong? Without an explicit boundary, responsibility is ambiguous exactly when a postmortem most needs clarity.

Rule ID	Phase	Specification
HAND-ACC-01	**P1**	Every handoff records an explicit Decision Ownership Boundary with three separated stages: AI Recommendation (what §7 Agent 5 suggested and why), Human Decision (what the human actually decided, per the Layer 3 Decision Record), and Human Execution (how the decision was carried out in the actual conversation/email/call). Each stage is independently attributable.
HAND-ACC-02	**P1**	Post-outcome reviews (§13 Analytics) classify failures into one of four categories: reasoning failure (the AI's underlying analysis was wrong — e.g. a hallucinated or miscalibrated fact), recommendation failure (the analysis was right but the strategic recommendation built on it was poor), execution failure (the recommendation was sound but human execution diverged or fell short), or external reality failure (everything upstream was reasonable; the outcome was driven by something genuinely unforeseeable — the §6 v3 ENR-BUDGET-01 'Currently Unknowable' category surfacing downstream).
HAND-ACC-03	**P1**	Failure classification feeds the correct calibration target: a reasoning failure updates §7/§6 model calibration; a recommendation failure updates §7 v2 Agent 5 strategy weighting; an execution failure feeds §8 v2 QC-HUMAN reviewer/owner Calibration Score (the same mechanism §11 v1 HAND-WF-07 already extended to account owners); an external reality failure updates nothing — it is logged and explicitly NOT used to penalize any component, consistent with EPIS-RCM-05's anti-fabrication discipline (don't manufacture a lesson from pure noise).

LAYER 5 — AUTHORITY TRANSITION (AUTH)

The deepest discovery from the first five iterations: §11 is not a workflow, it is an authority transition. The engine is saying 'I have reduced uncertainty as much as I responsibly can.' This IS GOV's Autonomy Ladder, viewed from the moment it hands off rather than the moment it acts.

Rule ID	Phase	Specification
HAND-AUTH-01	**P1**	Every handoff package states explicitly, in plain language, WHY automation stopped — rendered from the specific GOV Autonomy Tier and trigger that fired (GOV-ESC-01/02), not a new taxonomy: 'Escalated because: Political Sensitivity Rating is High on this account's current narrative (§7 v2 AI-POL-02 permanent gate)' or 'Escalated because: Trust Accumulation Score crossed the threshold for direct pricing discussion, which is always Human-led (§11 v2 Layer 10).'
HAND-AUTH-02	**P1**	The handoff package includes Remaining Unknowns — sourced directly from the account's §6 v3 Residual Uncertainty Register (ENR-BUDGET-03) — not just Knowns. A human walking in confident about everything the engine knows, blind to what it has explicitly flagged as unknowable, is exactly the overconfidence EPIS-RCM-04 exists to prevent, now extended to the human-facing side of the boundary.

LAYER 6 — AUTOMATION BIAS (BIAS)

Rule ID	Phase	Specification
HAND-BIAS-01	**P1**	Every recommendation displays its EPIS-RCM Confidence Level and Key Assumptions directly alongside the recommendation text — not buried in an appendix. 'System is usually right' should not be separable from seeing exactly how confident the system actually is on THIS recommendation.
HAND-BIAS-02	**P1**	High-confidence recommendations (above a configured RCM threshold) REQUIRE a Contrarian View in the package: 'what might make this recommendation wrong?' — this is EPIS-RCM-03's falsifiability requirement ('what would change our mind') made mandatory and human-facing specifically because high confidence is precisely when automation bias is most dangerous.
HAND-BIAS-03	**P1**	The package's design objective is explicitly to STIMULATE judgment, not replace it — this is an architectural test applied to every §11 v2 rule going forward, parallel to §5 v3's attention-allocation test and §10 v2's belief-correction test: a proposed §11 rule that makes the human's review faster but less critical fails this test.

LAYER 7 — REVERSE HANDOFF (KNOW)

The package is half the equation. A rep who personally knows the CTO, or learns of a budget freeze in conversation, holds intelligence the engine cannot independently acquire. Currently this either stays in the rep's head or gets buried as unstructured CRM free text — lost to the system that most needs it.

Rule ID	Phase	Specification
HAND-KNOW-01	**P1**	After any handoff interaction, the human can submit Missing Context Feedback (hidden sponsor identified, budget freeze mentioned, political concern surfaced, relationship detail) through a structured input — not a CRM free-text field that nothing downstream reads.
HAND-KNOW-02	**P1**	Missing Context Feedback enters the engine through the EXISTING signal pipeline: it is created as a MANUAL-source signal (V2 §3.1's existing Manual Input source type) with the human as the named source, NOT a new ingestion mechanism. §3 SIG-DET normalization, dedup, and freshness rules apply to it exactly as they apply to any other signal.
HAND-KNOW-03	**P1**	Human-sourced signals carry a full RCM stamp (EPIS-RCM-01) like any other signal: source reliability (calibrated per-reporter over time, similar in spirit to §8 v2's Reviewer Calibration Score), confidence, and a decay profile appropriate to the content (a 'budget freeze mentioned' fact decays faster than a 'longtime personal relationship with the CTO' fact, per §6 v3 Layer 5's multi-layer decay categories).

LAYER 8 — EXPERTISE ALLOCATION (EXP)

A single owner is sometimes the wrong unit. A strong relationship-builder fielding a deep technical-compliance discussion needs a subject-matter expert alongside them, not instead of them. This extends §5 v3's existing Expertise Graph (TIER-ROUTE-001) rather than building a second routing model.

Rule ID	Phase	Specification
HAND-EXP-01	P2	Every escalated engagement receives an Expertise Requirement Profile: tags drawn from the same dimension set as §5 v3's Human Expertise Graph (geography, industry, product, relationship network, language) plus discussion-specific tags (technical, regulatory, procurement, executive, integration) inferred from the triggering reply content and §7 v2 AI-POL Political Sensitivity context.
HAND-EXP-02	P2	Where the Expertise Requirement Profile indicates a need beyond the assigned owner's profile, the system recommends Support Roles — named individuals or role types to loop in — using the same Opportunity × Human Fit scoring as §5 v3 TIER-ROUTE-002, applied per-requirement rather than producing a single primary-owner match.
HAND-EXP-03	P2	Complex opportunities are explicitly modeled as multi-human collaboration objects: the handoff package and CRM task (§12) can carry more than one assigned human, each tagged with their specific role (primary owner, technical support, regulatory support), rather than forcing a single-owner field to awkwardly represent a team.

LAYER 9 — COGNITIVE DIVISION OF LABOR (CDL)

Rule ID	Phase	Specification
HAND-CDL-01	P1	Every decision TYPE (not individual action — that's GOV's Autonomy Ladder) in the ABM pipeline is classified as AI-led, Human-led, or Joint. Examples: signal aggregation → AI-led; pricing negotiation → Human-led; opportunity strategy → Joint. CDL sits one level above GOV: it classifies decision CATEGORIES, which then resolve to specific GOV Autonomy Tiers for the individual actions within them.
HAND-CDL-02	P1	The handoff package explicitly identifies Human Judgment Areas for this specific engagement — the points where §7 v2's EPIS confidence is insufficient for the engine to have a strong view, flagged so the human knows where their judgment is the binding constraint, not a formality.
HAND-CDL-03	P1	§13 Analytics' success metric for handoff quality is explicitly Combined Decision Quality — the joint outcome of AI recommendation quality and human decision quality together — never AI quality or human quality scored in isolation. This is the §11-level instance of §8 v2's Trust Preservation principle: the system succeeds or fails as a partnership, not as two separately-graded components.

§11 v2 RULE ROLL-UP & PLACEMENT STATUS

Layer	Rules	Phase	Cross-section binding
1 Comp	HAND-COMP-01..03	P1	§6 v3 ENR-QUAL-01 (decision relevance over richness)
2 Judge	HAND-JUDGE-01..02	P1	§7 Agent 5 (transparency, not new reasoning)
3 Learn/Mem	HAND-DEC-01..04	P1 / P2	§10 v2 ENG-MEM-01 (one Tier 6 pipeline), §7 v2 Agent 5 calibration feedback
4 Acc	HAND-ACC-01..03	P1	§8 v2 QC-HUMAN-08 Calibration Score, EPIS-RCM-05 anti-fabrication
5 Auth	HAND-AUTH-01..02	P1	GOV-LAD/ESC (reused, not duplicated), §6 v3 ENR-BUDGET-03 Residual Uncertainty Register
6 Bias	HAND-BIAS-01..03	P1	EPIS-RCM-03 (reused), the §11 architectural test
7 Know	HAND-KNOW-01..03	P1	V2 §3.1 Manual Input source (reused), EPIS-RCM-01, §6 v3 multi-layer decay
8 Exp	HAND-EXP-01..03	P2	§5 v3 TIER-ROUTE-001/002 (extended, not duplicated)
9 CDL	HAND-CDL-01..03	P1	GOV Autonomy Ladder (one level up), §8 v2 Trust Preservation (Combined Quality)

§11 v2 COMPLETE — WHAT IT ESTABLISHED

§11 is no longer a notification workflow. It is the **Human-AI Decision Partnership Engine**: it compresses context to a three-item executive summary, transfers reasoning and rejected alternatives rather than bare conclusions, captures every override and decision as one Decision Record feeding Tier 6 memory, draws an explicit Decision Ownership Boundary so post-outcome reviews can correctly attribute failure, renders GOV's existing autonomy logic as a human-readable 'why automation stopped,' forces a Contrarian View on high-confidence recommendations to fight automation bias, lets human-discovered intelligence re-enter through the existing signal pipeline, extends the existing Expertise Graph to multi-role teams, and classifies whole decision categories as AI-led/Human-led/Joint one level above GOV's per-action tiers. 27 new rules, with five integration decisions ensuring none of this duplicates GOV, EPIS, Tier 6, or §5 v3's expertise routing — it extends them.

The success metric for §11 is now explicit: Combined Decision Quality, not package delivery and not AI quality in isolation. The handoff is not the end of automation — it is the beginning of collaborative intelligence.

PLACEMENT MAP UPDATE

§11 status: Complete (OPEN-Q11.1, OPEN-Q11.2 from §11 v1 remain open but unaffected by this pass). No new open questions — every iteration in this pass resolved to an integration with existing architecture rather than a genuinely new undecided question.

ON THE HORIZON

§12 CRM Synchronization Engine — where, per the standard assessment pattern, hidden production failures (authority conflicts, duplicate realities, sync loops, ownership conflicts, data drift) typically emerge. §12 inherits §11 v2's multi-human Support Roles (Layer 8) and Decision Records (Layer 3) as new CRM sync payloads that V2's original field-level design did not anticipate.