ABM BUSINESS LOGIC BIBLE

§13 v2 — ORGANIZATIONAL LEARNING GOVERNANCE ENGINE

10 Iterations · 10 Architectural Layers · 28 New Rules · 4 Integration Decisions · 1 Reconciliation

§13 v1 specified KPIs and feedback loops. §13 v2 reframes the engine itself, and protects everything built before it. **§13 is not a measurement system. It is the Organizational Learning Governance Engine.** Dashboards, charts, and KPIs are artifacts — the actual product is better organizational judgment. Every KPI is a gaming vector the moment it becomes a target; every module-level improvement is a candidate system-level harm; the accounts never pursued teach as much as the ones that converted; and the cumulative effect of many small, individually-approved learning changes can quietly turn this into a different engine than the one specified — unless the Bible's core principles are made explicit, non-negotiable constraints that learning may optimize within but never rewrite.

Iterations	10	**Layers**	10
New rules	28	**Integration decisions**	4 explicit

FOUR INTEGRATION DECISIONS + ONE RECONCILIATION

Goodhart's Law + Local-vs-Global are ONE check, not two review processes: every optimization proposal tests both 'does this distort its own KPI' and 'does this hurt a different KPI' in a single pass against the System Health Dashboard.

Survivorship Bias + Counterfactual Analytics are ONE registry (non-pursued accounts and missed opportunities are the same dataset), bound to §5 v3 TIER-PORT-003's existing Displaced Opportunity tracking, not a parallel mechanism.

Model Drift is EPIS-HALF applied at the model level, not the fact level — Predictive Stability Score is the model-validity analog of the semantic half-life already tracked per fact.

Learning Governance's 'Core Principles' bind explicitly to the actual Section 1 Meta-Principles and named load-bearing rules (AI-POL-02, vendor exclusivity, EPIS-RCM-05) — not an abstract new list.

RECONCILIATION — Exploration Budget vs trust-first discipline: exploration earns a budget ALLOCATION, never an exemption from any hard gate. Exploratory outreach still passes QC-RULE-008/009/010 and AI-POL-02 exactly like exploitation outreach.

THE TEN LAYERS — §13 v2 ARCHITECTURE MAP

#	Layer	From iter.	What it adds
1	Goodhart's Law + Local-vs-Global (Good/Sys)	1, 2	One System Health Guardrail check per optimization proposal
2	Survivorship + Counterfactual (Surv/CF)	3, 6	One Opportunity Miss Registry, bound to §5 v3 displacement tracking
3	Learning Latency (Lat)	4	KPIs classified by horizon; don't judge an 18-month bet on a 30-day window
4	Judgment as the Output (Judge)	5	Every KPI must answer: which decision does this improve?
5	Blind Spot Detection (Blind)	7	Find missing variables, not just optimize existing ones
6	Model Drift (Drift)	8	Predictive Stability Score — EPIS-HALF applied to whole models
7	Exploration vs Exploitation (Exp)	9	Reserved budget for the unknown — never exempt from safety gates
8	Learning Governance (Gov)	10	Learning Change Ledger + quarterly Review Board
9	Constitutional Constraints (Gov-2)	10	Section 1 Meta-Principles named as non-negotiable, explicitly
10	(synthesis)	5+10	Analytics succeeds when judgment improves, not when dashboards fill up

LAYER 1 — GOODHART'S LAW + LOCAL-vs-GLOBAL OPTIMUM (GOOD / SYS)

The single biggest danger in §13: when a measure becomes a target, it stops being a good measure. A provocative message can raise Reply Rate while trust and reputation quietly fall — the KPI rises while the system gets worse. And a module-level win (Loop 2 finds a higher-reply-rate CTA) can be a system-level loss (trust and decision quality fall) even when the local metric looks like a clean improvement.

Rule ID	Phase	Specification
ANL-GUARD-01	**P1**	Every KPI (V2's original nine plus §13 v1's five new ones) has a declared Anti-Gaming Companion Metric: Reply Rate ↔ Trust Preservation Score; Meeting Rate ↔ Reputation Yield; Conversion Rate ↔ Cause-Attributed Conversion. A KPI cannot be reported, in any dashboard or review, without its companion alongside it.
ANL-GUARD-02	**P1**	Every optimization proposal (from any §13 v1 feedback loop) is evaluated in ONE pass against the System Health Dashboard (Trust Preservation, Reputation Yield, Combined Decision Quality, Reality Convergence): it must show its Intended Improvement on its target metric AND its Possible Distortion on every System Health Dashboard metric. A single check, not two separate review processes.
ANL-GUARD-03	**P1**	A proposal is REJECTED if any System Health Dashboard metric would worsen, even when its target local metric improves — global health overrides local gain, with no exception requiring a separate appeal process; this is the binding rule, not a guideline.
ANL-GUARD-04	**P2**	A quarterly KPI Gaming Investigation retrospectively checks whether any APPROVED change, in hindsight, distorted behavior in a way the pre-approval review missed — catching gaming that wasn't visible until enough data accumulated.

LAYER 2 — SURVIVORSHIP BIAS + COUNTERFACTUAL ANALYTICS (SURV / CF)

Of 100 accounts, analytics studies the 10 that became opportunities. The other 90 — deprioritized, waitlisted, never enriched, rejected at escalation — are invisible to every feedback loop. Whether they were correctly screened out or genuinely missed is unknown, because nothing currently asks.

Rule ID	Phase	Specification
ANL-MISS-01	**P1**	The Opportunity Miss Registry extends §5 v3 TIER-PORT-003's existing Displaced-Opportunity tracking (which already records what a HOT promotion displaced) to ALSO record: accounts that never reached HOT, HOT_WAITLIST accounts that timed out, rejected escalations (§11 v2), and abandoned sequences — one registry, not a new parallel dataset alongside the existing displacement record.
ANL-MISS-02	**P1**	Quarterly review explicitly compares Pursued vs Non-Pursued account populations on observable downstream outcomes (did a non-pursued account get won by a competitor per V2 EDGE-SIG-003; did a waitlisted account's opportunity window close) — the question is asked, not only 'what did we pursue' but 'what did we not, and what happened to it.'
ANL-MISS-02B	**P1**	Feedback loops (§13 v1 Loops 1–5) read Action Outcomes AND Non-Action Outcomes from the Opportunity Miss Registry — a scoring-weight or routing rule that systematically deprioritizes a pattern that, in hindsight, the Registry shows often converts elsewhere is a calibration signal exactly as valid as a pursued-account outcome.

LAYER 3 — LEARNING LATENCY (LAT)

Rule ID	Phase	Specification
ANL-HORIZON-01	**P1**	Every KPI is classified by Learning Horizon: Short (Reply Rate, Email Open Rate — evaluable within days), Medium (Meeting Rate, Signal-to-Meeting Conversion — weeks to a quarter), or Long (Reputation Yield, Reality Convergence, the Loop 5 trust-first hypothesis itself — 6–18 months).
ANL-HORIZON-02	**P1**	A feedback loop only evaluates a metric's success/failure AFTER its declared Learning Horizon has elapsed since the relevant change — §13 v1 Loop 5 (90-day cadence) does not get judged on whether its FIRST 90-day check shows dramatic results; per ANL-LOOP5-01's own design, it needs the Long-horizon window to say anything meaningful.
ANL-HORIZON-03	**P1**	Long-horizon strategies (anything gated by §7 v2 AI-REP-06's reputation-over-conversion trade-off) are explicitly NOT penalized using Short-horizon windows — a quarter of lower Reply Rate in service of a trust-building strategy is not, on its own, evidence the strategy failed.

LAYER 4 — JUDGMENT AS THE REAL OUTPUT (JUDGE)

Rule ID	Phase	Specification
ANL-JUDGE-01	**P1**	Every KPI, in its design documentation, must answer: what decision does this improve? A KPI with no identifiable decision it informs is flagged as a vanity metric and removed from standing dashboards (it may still be logged for historical record, per the append-only discipline, but is not a tracked KPI).
ANL-JUDGE-02	**P1**	Every feedback loop's proposal (§13 v1 Loops 1–5) must identify which FUTURE decision becomes better because of the proposed learning — this is a required field in the A4 governance proposal, not implicit.
ANL-JUDGE-03	**P1**	§13's own success is measured by improvement in organizational judgment quality (operationalized as: Combined Decision Quality trending up, per §11 v2 HAND-CDL-03) — NOT by dashboard completeness, report cadence, or the sheer number of KPIs tracked.

LAYER 5 — BLIND SPOT DETECTION (BLIND)

Existing loops optimize known dimensions — signals, personas, messages, trust. None of them can discover a variable the system isn't modeling at all. Repeated unexplained outcomes are the signal that something is missing, not that an existing model needs more tuning.

Rule ID	Phase	Specification
ANL-BLIND-01	**P2**	An Unexplained Outcome Dataset retains outcomes (wins or losses) whose variance from model-predicted likelihood exceeds a configured threshold across every known explanatory factor (signals, trust, persona, narrative, Cause Attribution) — outcomes the existing model genuinely cannot account for, not merely low-confidence ones.
ANL-BLIND-02	**P2**	Sustained accumulation of Unexplained Outcomes (not a single anomaly) triggers a Blind Spot Investigation — a distinct, human-led process from a standard Failure Cluster investigation (§8 v2 QC-SYS-02), because the goal is finding a MISSING variable, not diagnosing a known mechanism's malfunction.
ANL-BLIND-03	**P2**	A Blind Spot Investigation's success criterion is a NEW explanatory factor proposed for the architecture (e.g. 'we are not currently modeling X') — not a tuning adjustment to an existing rule. Its output, if validated, becomes a candidate addition to the §2 entity model or a new INT-SIG sub-stream, following the same governance path as any other architectural change.

LAYER 6 — MODEL DRIFT (DRIFT)

A model can be correct and later become wrong — not because it broke, but because regulations changed, competitors emerged, or buying behavior evolved. This is EPIS-HALF's semantic-half-life concept, applied to a whole model's validity rather than a single fact's freshness.

Rule ID	Phase	Specification
ANL-DRIFT-01	P2	Every major model (§4 scoring weights, §7 v2 Agent 5 strategy calibration, §6 v3 marginal-value estimates) receives a Predictive Stability Score: how well the model's historic assumptions still explain CURRENT outcomes, recomputed on each feedback-loop cycle — the model-level analog of EPIS-HALF-01's per-fact decay curve.
ANL-DRIFT-02	P1	Loops monitor Performance Decay (the trend in Predictive Stability Score over time), not only the current performance level — a model that is still performing acceptably but visibly decaying is flagged before it crosses a failure threshold, not only after.
ANL-DRIFT-03	P1	Rapid decay triggers a Revalidation Review — explicitly human-led, NOT automatic retraining. This matches GOV-MOD-01/02 throughout the Bible: a model whose assumptions may no longer hold requires a human to determine why before any automated recalibration is permitted to run.

LAYER 7 — EXPLORATION vs EXPLOITATION (EXP)

A learning system that only reinforces what already works (target more CTOs because CTOs convert best) eventually goes blind to emerging buyers, new personas, new narratives — the classic exploitation trap. But every exploration dollar is still real outreach to a real Saudi bank executive, so exploration cannot mean lower safety standards.

Rule ID	Phase	Specification
ANL-EXPL-01	P2	A reserved Exploration Budget (target 10–15% of outreach volume) is allocated to new personas, new narratives, new industries, or new outreach patterns not yet validated by the mature strategy models.
ANL-EXPL-02	P1	RECONCILIATION: Exploration Budget outreach passes through EVERY hard gate exactly as exploitation outreach does — QC-RULE-008 (narrative drift), QC-RULE-009 (political sensitivity, permanently gated per AI-POL-02), QC-RULE-010 (stakeholder-strategy consistency), and the full §8 v1/v2 QC chain. The budget allocates WHERE outreach goes; it grants no exemption from HOW safely it must be executed.
ANL-EXPL-03	P2	Exploration outcomes are evaluated separately from mature-strategy outcomes — a new narrative pattern is not penalized for underperforming a five-touch-validated mature pattern on its first attempt; §13 tracks a Learning Yield metric alongside conversion yield, crediting an exploration attempt that produces useful negative knowledge (a confirmed non-viable pattern) as valuable, not wasted.

LAYER 8 — LEARNING GOVERNANCE (GOV)

Every loop changes scoring, prompts, priorities, thresholds. Each individual change is reviewed and A4-approved. But many small approved changes can collectively produce a different engine than the one this Bible specified — and nothing currently asks what the cumulative drift adds up to.

Rule ID	Phase	Specification
ANL-LCHANGE-01	P1	A Learning Change Ledger tracks every approved and rejected learning-driven change (weight adjustments, prompt updates, threshold modifications) across all five feedback loops, with cumulative-impact annotation — not just the individual change log each loop already keeps, but one consolidated cross-loop record.
ANL-LCHANGE-02	P1	A quarterly Learning Review Board (human, senior — the same standing implied by §13 v1 OPEN-Q13.2's named-senior-approver discipline) evaluates the Ledger's cumulative pattern and asks explicitly: 'what has the engine become?' — not merely 'what changed this quarter,' since no single change need look alarming for the sum of a year's changes to represent real architectural drift.

LAYER 9 — CONSTITUTIONAL CONSTRAINTS (GOV-2)

Rule ID	Phase	Specification
ANL-CONST-01	P1	The following are explicitly Constitutional Constraints: learning loops may optimize WITHIN them, never propose changes that weaken or rewrite them. Named set: the 25 Section-1 Meta-Principles; AI-POL-02's permanent political-sensitivity human-review gate; SIG-VENDOR-01's per-account vendor exclusivity; EPIS-RCM-05's anti-fabrication rule; the §5 v3 A0 hard-gates (sanctioned-entity, opt-out); and account-centricity (V2 Principle 2).
ANL-CONST-02	P1	Any feedback-loop proposal that would loosen, bypass, or reinterpret a Constitutional Constraint is automatically rejected at the Learning Review Board (Layer 8) regardless of its projected metric improvement — this is the single hardest stop in §13, harder even than ANL-GUARD-03's System Health rejection, because Constitutional Constraints are not metrics to be traded off against each other at all.
ANL-CONST-03	P2	The Constitutional Constraints list itself can only be amended by an explicit, named governance process outside the standard A4 cycle (consistent with §5 v3's escalation-ladder principle that the stakes of a decision should set the seniority required to make it) — it is not itself subject to quarterly Learning Review Board revision as a matter of course.

§13 v2 RULE ROLL-UP & PLACEMENT STATUS

Layer	Rules	Phase	Cross-section binding
1 Good/Sys	ANL-GUARD-01..04	P1 / P2	§13 v1's existing KPI set as the companion-metric pairs
2 Surv/CF	ANL-MISS-01..02B	P1	§5 v3 TIER-PORT-003 (extended, not duplicated)
3 Lat	ANL-HORIZON-01..03	P1	§13 v1 Loop 5, §7 v2 AI-REP-06
4 Judge	ANL-JUDGE-01..03	P1	§11 v2 HAND-CDL-03 (Combined Decision Quality)
5 Blind	ANL-BLIND-01..03	P2	§2 entity-model governance path, INT-SIG sub-stream addition path
6 Drift	ANL-DRIFT-01..03	P1 / P2	EPIS-HALF-01 (model-level analog), GOV-MOD-01/02
7 Exp	ANL-EXPL-01..03	P1 / P2	§8 hard gates (no exemption), §7 v2 AI-POL-02 (permanent)
8 Gov	ANL-LCHANGE-01..02	P1	§13 v1 OPEN-Q13.2 (senior approver pattern)
9 Gov-2	ANL-CONST-01..03	P1 / P2	Section 1 Meta-Principles, AI-POL-02, SIG-VENDOR-01, EPIS-RCM-05, §5 v3 A0 gates

§13 v2 COMPLETE — WHAT IT ESTABLISHED

§13 is no longer a measurement system. It is the **Organizational Learning Governance Engine**: it pairs every KPI with an anti-gaming companion and rejects any proposal that improves a local metric while harming system health, learns from accounts never pursued as well as accounts won or lost, respects that different truths reveal themselves on different timelines, insists every KPI and every learned change name the decision it improves, hunts for the variables it isn't modeling at all rather than only tuning what it already models, tracks whether its models are still valid as the world moves rather than assuming correctness is permanent, reserves a safety-gated budget for genuine exploration, and — most importantly — tracks the cumulative drift of its own learning against a named, explicit, non-negotiable set of Constitutional Constraints that no metric improvement is ever allowed to override. 28 new rules.

This is the layer that makes every other section's discipline durable. Without it, the Bible's trust-first architecture is one quarter of well-intentioned, individually-reasonable optimizations away from becoming a different system than the one specified.

PLACEMENT MAP UPDATE

§13 status remains Complete-with-Opens (OPEN-Q13.1–13.3 from §13 v1 unchanged — this pass added architecture, not new undecided questions). Per both assessments, §13 is now ~98–99% architecturally exhausted; further iteration would generate philosophy rather than new architecture.

ON THE HORIZON

§14 State Management — per the standard assessment, likely the richest remaining section, since every module ultimately operates on state transitions and hidden state-machine failures are where complex ABM systems usually break in production. §17 KSA-Specific Business Rules remains the other major 'Not started' gap on the Placement Map.