

ABM BUSINESS LOGIC BIBLE

§18 v2 — EDGE CASE EPISTEMICS

How the Engine Handles Being Wrong About Its Own Edge Cases

PROVENANCE NOTE — READ BEFORE ANYTHING ELSE

Iterations 1–2 ('Unknown Unknown Problem', 'Multi-Edge Collision Problem') and the back half of Iteration 3 ('False Positive Edge Problem') were referenced by name/fragment in the source material but their full content was not provided. Per explicit instruction, Claude RECONSTRUCTED these three iterations in the established pattern (hidden discovery → new rule → deepest rewrite) rather than recovering original text. They are clearly marked as reconstructed below. Iterations 4 onward, where present, follow the same discipline applied fresh to this section. Anyone citing this document should know that Iterations 1–3's specific framing is Claude's inference about what the critique most likely intended, not a verbatim transcript of work already done.

§18 v1 reconciled two real catalogues — V2's local edge cases and the EDGE-01–16 global scenarios — and called itself Complete. The critique is sharper than 'incomplete': §18 v1 **makes confident mistakes at scale**. A catalogue of known scenarios, each with a definite response, projects more certainty than the engine should have about (1) whether it has even seen all the scenarios that exist, (2) what happens when several fire at once, and (3) whether a given trigger was ever real in the first place. §18 v2 is not a fourth edge-case list. It is the layer that governs the catalogue's own epistemic honesty.

ITERATION 1 (RECONSTRUCTED) — THE UNKNOWN UNKNOWN PROBLEM

§18 v1's catalogue, by construction, contains only scenarios someone already thought of. EDGE-16 ('review the catalogue at every update, add new scenarios before they recur') already implicitly admits the catalogue is incomplete — but it treats incompleteness passively, as something resolved AFTER the fact. The deeper discovery: the engine needs a defined behavior for the FIRST occurrence of a genuinely unclassified scenario, not just a promise to eventually classify it once it has recurred.

Rule ID	Phase	Specification
EDGE-UNK-01	P1	Every event that does not match an existing V2 local edge case (§18.0–18.3) AND does not match an existing EDGE-01–16 detector is classified UNCLASSIFIED_ANOMALY — a first-class outcome, not a fallback error. The system never silently processes an unmatched event as if it were ordinary.
EDGE-UNK-02	P1	An UNCLASSIFIED_ANOMALY default-routes to GOV A3 (human review) regardless of which account or module it occurred in — the engine's default posture toward a genuinely novel scenario is to ask a human, not to guess at a response by analogy to an existing entry that may not actually fit.
EDGE-UNK-03	P1	Every UNCLASSIFIED_ANOMALY is logged with full context and reviewed against §15 v2's FAIL-ANTI-01 Capability Improvement Proposal discipline: if it recurs (per §15 v2 FAIL-REPEAT-01's recurrence counter, generalized here to edge scenarios), a Root Cause Review evaluates whether it should become a new §18.4 EDGE catalogue entry — this is EDGE-16 made concrete: the FIRST occurrence gets a safe, human-gated default; only a RECURRING unclassified pattern triggers formal catalogue promotion.

ITERATION 2 (RECONSTRUCTED) — THE MULTI-EDGE COLLISION PROBLEM

§18 v1's EDGE-RECON-02 noted ONE specific pair (competitor-win + vendor-merger) that can co-fire without conflict. That was a special case, not a general solution. The deeper problem: what happens when THREE OR MORE EDGE scenarios are active on the same account simultaneously, especially when their RESPONSES actively conflict — one demands an immediate pause, another demands escalation, a third demands demoting a path that the first two don't reference at all?

Rule ID	Phase	Specification
EDGE-COL-01	P1	Every account maintains a visible set of CURRENTLY ACTIVE EDGE scenarios (not just the most recent one) — §18 v1 implicitly assumed one scenario at a time; this rule makes multiplicity an explicit, queryable state rather than something the engine discovers by accident when two fire close together.
EDGE-COL-02	P1	When multiple active EDGE scenarios' responses do NOT conflict (e.g. E4 vendor-merger's re-analysis and E8 departed-exec's path demotion), both responses execute independently — confirming §18 v1 EDGE-RECON-02's principle generalized beyond its one named pair.
EDGE-COL-03	P1	When multiple active EDGE scenarios' responses DO conflict (e.g. one demands immediate pause, another demands continued monitoring), the MOST RESTRICTIVE response wins — specifically, the response associated with the LOWEST GOV Autonomy Tier (A0 strictest) governs, and all less-restrictive responses are suspended until every active scenario resolves. An account is never partially-paused and partially-active at once.
EDGE-COL-04	P1	A multi-edge collision (3+ simultaneous scenarios, or any 2+ scenario collision resolved by EDGE-COL-03's restrictive-wins rule) is itself logged as a distinct event type and reviewed at the same cadence as §15 v2's quarterly Resilience Investment review — frequent collisions on the same account or same scenario pair may indicate the catalogue's scenarios are not as independent as §18 v1 assumed.

ITERATION 3 — THE FALSE POSITIVE EDGE PROBLEM (COMPLETED)

The source fragment: a source reports 'Bank under regulatory investigation.' EDGE-04 fires. Outreach pauses; governance review triggers. Completing the scenario: the source is later found to be wrong — a misreported story, a rumor, or an investigation that actually named a different entity. The engine has now paused a real opportunity, consumed a human reviewer's attention, and logged a false escalation, entirely on the strength of an unverified trigger. §18 v1's entire catalogue describes how to ENTER an edge state. None of it describes how to EXIT one when the trigger itself turns out to have been wrong.

Rule ID	Phase	Specification
EDGE-FP-01	P1	Every EDGE scenario trigger (§18.4 EDGE-01's detector) carries its EPIS-RCM confidence stamp (per §4 v2 EPIS-RCM-01) at the moment it fires — a HIGH-confidence trigger (e.g. corroborated across multiple independent sources) and a LOW-confidence trigger (a single unverified source) produce the SAME initial response (per §18 v1's existing rules, since the downside of under-reacting to a real regulatory action is severe) but are tagged differently for what follows.
EDGE-FP-02	P1	A LOW-confidence EDGE trigger carries a mandatory Verification Window (e.g. 48–72 hours depending on scenario severity) during which the triggering claim is actively checked against corroborating sources (§3 SIG-NEWS/SIG-REG cross-referencing) — the response (pause, escalation) remains in effect throughout the window; the window's purpose is determining whether it should remain in effect AFTER.
EDGE-FP-03	P1	If a LOW-confidence trigger fails to corroborate within its Verification Window, the EDGE scenario is marked RETRACTED (not silently closed) — the account/path/opportunity returns to its pre-trigger state, the retraction is logged with the original trigger's source and confidence, and this becomes an input to that SOURCE's reliability score (§4 v2 EPIS-SRC-01) — a source whose claims repeatedly fail to corroborate degrades in trust exactly like any other source in this Bible.
EDGE-FP-04	P1	A RETRACTED scenario does NOT automatically restore any GOV-A3-approved human decisions made during the active window (e.g. if a human reviewer, during the pause, separately decided the account should be deprioritized for unrelated reasons) — retraction undoes the AUTOMATIC system response; it never silently overrides a human's independent judgment made in the interim.
EDGE-FP-05	P2	§13 v2's Failure Cost discipline (§15 v2 FAIL-COST-01's Expected-Value family) is applied to false-positive EDGE triggers specifically: the Business Impact of a paused real opportunity (lost time, reviewer attention spent, relationship momentum interrupted) is tracked as a cost category distinct from a true-positive's cost — a catalogue that fires confidently and wrongly at scale has a real, measurable price, and §18 v2 makes it visible rather than treating every pause as costless caution.

ITERATION 4 — THE STALE TRIGGER PROBLEM

A trigger that was TRUE when it fired can become false later without ever being retracted — a different failure mode from Iteration 3's false positive. A regulator action (E2) genuinely occurred; six months later the investigation closes with no findings. §18 v1 has no mechanism for an EDGE scenario that was correctly entered to correctly EXPIRE.

Rule ID	Phase	Specification
EDGE-STALE-01	P1	Every active EDGE scenario (§18.4) carries an Expected Resolution Window (per §14 v2 STATE-TIME-01's pattern, generalized to EDGE state) — e.g. E2 regulator action defaults to a 90-day expectation, reviewed per §16 v1's governance-cadence discipline.
EDGE-STALE-02	P1	An EDGE scenario that exceeds its Expected Resolution Window without explicit GOV resolution (closed, escalated further, or extended with stated reason) triggers a mandatory re-review — NOT an automatic clearing. A regulator action that has simply gone quiet for 6 months is not assumed resolved; a human confirms its actual status before the pause lifts.
EDGE-STALE-03	P2	Stale-trigger re-reviews are tracked in the SAME Learning Change Ledger pattern as §13 v2 ANL-LCHANGE-01 — a scenario type that is frequently stale (resolved long before anyone re-reviews it) is a signal the Expected Resolution Window for that scenario type is miscalibrated, feeding back into EDGE-STALE-01's defaults.

ITERATION 5 (DEEPEST) — THE CATALOGUE'S OWN CONFIDENCE

The bedrock question underneath all four prior iterations: §18 v1 presented every entry with equal apparent authority — EDGE-01 through EDGE-16 read with the same confident, declarative voice as the V2 local cases. But the catalogue's own COVERAGE of reality is uneven, exactly the way any single account's intelligence coverage is uneven (EPIS-COV-01). The deepest discovery: the catalogue itself needs an honesty layer, the same way every other belief in this Bible does.

Rule ID	Phase	Specification
EDGE-EPIS-01	P1	The EDGE catalogue (§18.4) carries an aggregate Catalogue Coverage Score, reviewed at §15 v2's quarterly cadence: the rate of UNCLASSIFIED_ANOMALY events (Iteration 1) relative to total edge-triggering events is the catalogue's own coverage metric, exactly as EPIS-COV-01 measures an account's intelligence coverage. A rising UNCLASSIFIED_ANOMALY rate means the catalogue's coverage is DECLINING, even though no single existing entry changed.
EDGE-EPIS-02	P1	§18's own architectural test, going forward: a new EDGE-catalogue entry, or a new §18 v2 mechanism, is incomplete unless it states what happens when the entry itself turns out to be WRONG — this generalizes Iteration 3's retraction logic into a standing requirement, not a one-off fix for the regulatory-investigation example.
EDGE-EPIS-03	P1	§13 v2's Constitutional Constraints discipline (ANL-CONST-01) is extended explicitly to §18: NO EDGE catalogue entry, however well-evidenced, may be used to justify bypassing an A0 hard-gate (sanctioned-entity, opt-out, PDPL deletion) — the catalogue's job is to add caution and process, never to manufacture a path around the Bible's hardest constraints.

§18 v2 RULE ROLL-UP & PLACEMENT STATUS

Iteration	Status	Rules	Cross-section binding
1 Unknown Unknown	RECONSTRUCTED	EDGE-UNK-01..03	GOV A3 default, §15 v2 FAIL-ANTI/FAIL-REPEAT
2 Multi-Edge Collision	RECONSTRUCTED	EDGE-COL-01..04	GOV Autonomy Ladder (restrictive-wins), §15 v2 quarterly review
3 False Positive (completed)	Fragment completed	EDGE-FP-01..05	EPIS-RCM-01, EPIS-SRC-01, §15 v2 FAIL-COST-01
4 Stale Trigger	New, this pass	EDGE-STA-01..03	§14 v2 STATE-TIME-01, §13 v2 ANL-LCHANGE-01
5 Catalogue Confidence	New, this pass	EDGE-EPIS-01..03	EPIS-COV-01, §13 v2 ANL-CONST-01

§18 v2 COMPLETE — WHY 'COMPLETE' NO LONGER MEANS WHAT §18 v1 THOUGHT IT MEANT

The critique was right: §18 v1's two reconciled catalogues are genuinely complete AS CATALOGUES, but a catalogue that presents every entry with equal confidence, no exit mechanism, and no model of its own incompleteness makes confident mistakes at scale precisely BECAUSE it looks finished. §18 v2 adds: a first-class outcome for events that match nothing (Unknown Unknowns), a restrictive-wins resolution rule for scenarios that collide (Multi-Edge Collision), a retraction mechanism for triggers that were never real (False Positives), a re-review mechanism for triggers that WERE real but have gone stale (Stale Trigger), and — the deepest layer — a coverage metric for the catalogue's own confidence, exactly mirroring EPIS-COV-01's treatment of every other belief in this Bible. 14 new rules, with Iterations 1–2 and part of 3 explicitly marked as Claude's reconstruction rather than recovered original text.

PLACEMENT MAP UPDATE

§18 status: Complete-with-Opens is not required — all five iterations resolved to concrete rules. But the PROVENANCE of Iterations 1–3 (reconstructed, not recovered) should be tracked alongside §18's status so a future review knows to verify this reconstruction against the reviewer's actual original intent if that text is ever recovered.

ON THE HORIZON

§19 Glossary of Business Terms remains the only fully unbuilt V2-numbered item. Separately, this same epistemic-honesty pattern (Unknown Unknowns, Collision, False Positive, Stale Trigger, Catalogue Confidence) may be worth applying as a general audit lens across OTHER catalogues in this Bible (the Authority Registry, §15's Failure-type list, §3's signal-source list) if a similar deep pass is wanted there.