

ABM ENGINE — FINAL COHERENCE AUDIT

Full 1,058-Rule Set | Conflicts Flagged and Resolved | Decimal Technologies | June 2026

AUDIT METHOD AND HEADLINE FINDING

This audit was run against the complete 1,058-rule set (all 19 sections, all 5 Section-1 principles including the now-integrated Principle 3). The scan covered: duplicate rule IDs, dangling cross-references, numeric threshold contradictions, autonomy/review-gating conflicts, timing/coordination overlaps, scoring-formula consistency, and cultural/compliance handling.

Headline: the rule set is highly coherent. There are NO duplicate IDs and NO contradictory scoring formulas. The single Effective Opportunity / Decision Score equation (M2-EQUATION-001) is used consistently. ICS action thresholds are consistent (40 floor everywhere). The conflicts that exist are **9 real issues** — mostly overlapping rules that describe the same transition with slightly different numbers, plus two genuine gaps. All are resolvable without structural change.

Severity: 2 HIGH (must reconcile before build), 5 MEDIUM (reconcile during engineering), 2 LOW (documentation/cleanup).

FLAGGED CONFLICTS AND RESOLUTIONS

[HIGH] FIN-C01: Three Overlapping Review-Reduction Timelines With Different Anchors

Rules involved: HUMAN-AUTO-001, TRUST-001, FAIL-HUMAN-004 (+ AUTO-MVP-001)

The tension: Three separate rules each define when human review can be reduced, but they anchor the same transition to different timelines. HUMAN-AUTO-001 says Phase 2 begins Day 91. TRUST-001 says ESTABLISHED TRUST begins Day 90. FAIL-HUMAN-004 says Phase 2 is 'months 4-6' (~Day 120-180). An engineer reading all three cannot tell whether P3/P4 non-C-suite review becomes optional at Day 90, Day 91, or Day 120. They also each carry their own QC and override gates that mostly agree (QC>85%, override<10%) but are stated independently.

Resolution: Designate TRUST-001 (the 6-level Trust Progression Framework) as the SINGLE SOURCE OF TRUTH for autonomy progression. HUMAN-AUTO-001 and FAIL-HUMAN-004 become VIEWS of TRUST-001, not independent schedules. Add reconciliation rule TRUST-AUTHORITY-001: 'Review-reduction eligibility is governed solely by the Trust Progression level (TRUST-001). Day-based phases in HUMAN-AUTO-001 and FAIL-HUMAN-004 are minimums that cannot fire before the corresponding Trust level is reached. Where they disagree, the LATER date and the stricter gate win.' This makes the conservative interpretation always correct.

[HIGH] FIN-C02: Missing Governing Rule: GOV-001 Is Cited But Not Defined

Rules involved: EXPLAIN-003 cites 'GOV-001'; no GOV-001 exists in the set

The tension: EXPLAIN-003 logs negative decisions with the example 'Account X not promoted to HOT because ICS = 35 (below 40 per GOV-001).' This cites GOV-001 as the rule that sets the ICS>=40 promotion floor. But there is no rule with ID GOV-001 in the set. The ICS>=40 floor IS used consistently across TEMPORAL-DECISION-004 (enrichment sufficient at ICS>=40) and AIHUMAN-DISAGREE-002 (force-promote below ICS<40 triggers a warning) — so the threshold is real and consistent, but its defining rule is unnamed.

Resolution: Create the missing anchor rule GOV-ICS-001: 'An account may not be promoted to HOT/TIER_1 unless its Intelligence Confidence Score (ICS) is >= 40. Below 40, the account remains in WATCH state and enrichment continues. Human force-promotion below ICS 40 is permitted but triggers a logged SOFT WARNING (per AIHUMAN-DISAGREE-002).' Update EXPLAIN-003's reference from GOV-001 to GOV-ICS-001. This is a naming fix; the logic already exists and is consistent.

[MEDIUM] FIN-C03: Contact-Stagger Rule and Same-Contact Channel-Gap Rule Have a Coverage Seam

Rules involved: ACCT-COORD-001, TIMING-FAIL-003, INT-LIN-03

The tension: ACCT-COORD-001 governs spacing between DIFFERENT contacts at the same account (48h stagger). TIMING-FAIL-003 and INT-LIN-03 govern spacing between channels to the SAME contact (no email+LinkedIn within minutes; coordinated not concurrent). These are complementary, not contradictory — but neither states an explicit minimum gap for the same-contact cross-channel case. TIMING-FAIL-003 only gives the bad example ('9am + 9:05'); INT-LIN-03 says 'spaced per sequence design' without a floor. An engineer has no hard number for same-contact channel spacing.

Resolution: Add CHANNEL-GAP-002 to close the seam: 'Minimum gap between two different-channel touches to the SAME contact is 24 hours, unless the contact engaged on the first channel (open, click, accept), in which case the second channel may follow after 4 hours. This is distinct from ACCT-COORD-001, which governs spacing across different contacts at one account (48h).' Cross-reference both rules so the scopes are explicit.

[MEDIUM] FIN-C04: Two Minimum-Sample Standards for Learning

Rules involved: AILEARN-001 (Principle 3) vs GOODHART-SAMPLE-001 (audit resolution)

The tension: AILEARN-001 sets per-parameter minimum samples before any learning change: scoring weights need 10 meetings across 3+ signal types; persona prompts need 5 sequences; signal prioritization needs 15 signals; etc. The earlier coherence-audit resolution GOODHART-SAMPLE-001 set a flat '50 engagements across 15 accounts' before any weight recalibration. These overlap for the scoring-weights case and give different numbers (10 meetings vs 50 engagements).

Resolution: AILEARN-001 is the more granular, authoritative standard (it is per-parameter and comes from the deep Principle 3 excavation). Demote GOODHART-SAMPLE-001 to a PORTFOLIO-LEVEL guard that sits ON TOP of AILEARN-001: 'No GLOBAL weight recalibration affecting all accounts may fire until both (a) the per-parameter AILEARN-001 minimums are met, AND (b) at least 50 scored engagements across 15+ accounts exist (anti-Goodhart portfolio floor).' The two now stack rather than compete: AILEARN-001 gates individual parameter changes; GOODHART-SAMPLE-001 gates system-wide recalibration.

[MEDIUM] FIN-C05: P1-Signal 'Immediate Response' vs Mandatory Human Review Window

Rules involved: EVENT-RESP-001 vs HUMAN-AUTO-002 / GOVERNANCE-004

The tension: EVENT-RESP-001 requires the full cascade (scoring, tier, enrichment) to complete within 2 hours for P1 signals, and implies fast outreach. HUMAN-AUTO-002 requires a human review window (0-4h primary reviewer, up to 24h before expiry). For a P1 signal arriving at 11pm Thursday (KSA weekend starts), 'immediate' processing collides with 'mandatory human review' plus KSA send-window rules. GOVERNANCE-004 already noticed this exact tension and resolved it conceptually, but the resolution isn't encoded as an operational rule the other two point to.

Resolution: Promote GOVERNANCE-004's resolution into an explicit operational rule, QUEUE-ARCH-001: 'P1 processing (scoring, enrichment, message generation, QC) is immediate and runs 24/7 regardless of business hours. Human review and SEND are decoupled from processing: the QC-passed message enters the Ready-to-Fire Queue and waits for (a) human review per the active Trust level and (b) the next valid KSA send window. A P1 signal at 11pm Thursday is fully processed by ~1am, reviewed Sunday morning, sent in Sunday's 9am-12pm window.' EVENT-RESP-001 and HUMAN-AUTO-002 both reference QUEUE-ARCH-001. No rule changes — the queue satisfies all three.

[MEDIUM] FIN-C06: C-Suite 'Always Human Review' vs Full-Autonomy End State

Rules involved: HUMAN-AUTO-001 / HUMANREVIEW-002 / TRUST-001 (FULL PARTNERSHIP) vs Autonomy Ladder A3-A5

The tension: Multiple rules state C-suite messages are ALWAYS human-reviewed, permanently. The Autonomy Ladder (Document 3) defines A3-A5 supervised-to-full autonomy. A naive reader could think A5 eventually automates C-suite sends. It does not — but the rules don't explicitly carve C-suite out of the autonomy ceiling.

Resolution: Add AUTO-CEILING-001 as a constitutional carve-out: 'No autonomy level, including A5, ever removes human review from C-suite (and P1 Domino Account) messages. The RELATIONSHIP domain is permanently human (per AIHUMAN-FOUNDATION-001). Autonomy progression applies only to the data/execution domain below the C-suite relationship tier.' This aligns the Autonomy Ladder with the permanent-review rules and Meta-Principle MP-19.

[MEDIUM] FIN-C07: Trust Regression vs Progressive-Automation Forward-Only Phasing

Rules involved: TRUST-001 (trust can regress) vs FAIL-HUMAN-004 (phase roadmap reads forward-only)

The tension: TRUST-001 explicitly allows trust to REGRESS (a single hallucination drops one level; reputation damage drops to INITIAL). FAIL-HUMAN-004's 'Progressive Automation Roadmap' reads as a one-way march through Phase 1>2>3 by calendar month, with no defined reverse gear. If an incident occurs in month 7, do we revert automation or not?

Resolution: Make FAIL-HUMAN-004 explicitly subordinate to TRUST-001's regression mechanism: 'Automation phase is a function of current Trust level, not elapsed calendar time. Any trust regression event (TRUST-001) immediately reverts the corresponding automation phase by one level for the affected message categories, regardless of how many months have elapsed.' Calendar months become the earliest-possible schedule, not a guarantee.

[LOW] FIN-C08: Parser/Reference Truncation Artifacts (Documentation Cleanup)

Rules involved: DEP-RECORD-001, DNA-RECORD-001, INIT-RECORD-001, TIMING-MISS-001, TRUST-DEBIT-001

The tension: Five cross-references appear truncated in extraction (e.g. 'DEP-RECORD-0' instead of 'DEP-RECORD-001'). The target rules DO exist in the set with their full IDs; this is an artifact of how the source PDF wrapped lines, not a missing rule. Listed here only so it is not mistaken for a dangling reference during engineering review.

Resolution: No logic change. During the master-document cleanup pass, normalize all in-text rule references to their full 3-digit IDs so automated reference-checkers do not raise false positives. Confirmed: all five target rules are present and correctly defined in the set.

[LOW] FIN-C09: ICS Threshold Appears at 40, 60, and 35 — Confirm Not a Contradiction

Rules involved: TEMPORAL-DECISION-004 (≥ 40), HUMANREVIEW-003 (< 60), EXPLAIN-003 ($= 35$ example)

The tension: A numeric scan flagged ICS values of 35, 40, and 60 in different rules. On inspection these are NOT in conflict: 40 is the promotion/action floor (consistent everywhere); 60 is a separate, higher gate that triggers MANDATORY expansion of the reviewer's reasoning-chain view ($ICS < 60$ = look closer); 35 is merely an EXAMPLE of a value below the 40 floor used in a logging illustration. Flagged for confirmation, resolved as coherent.

Resolution: No change required. Documented here to confirm the audit examined it. Recommendation: add a one-line glossary note that ICS has two distinct gates — 40 (action floor) and 60 (review-scrutiny trigger) — so the two numbers are never conflated by implementers.

NEW RULES INTRODUCED BY THIS AUDIT

RECONCILIATION RULES TO ADD

TRUST-AUTHORITY-001 (FIN-C01): Trust Progression Framework is the single source of truth for autonomy. Day-based phases are minimums; later date and stricter gate win on disagreement.

GOV-ICS-001 (FIN-C02): Defines the $ICS \geq 40$ HOT-promotion floor that EXPLAIN-003 already references. Below 40 stays in WATCH; force-promotion triggers a logged warning.

CHANNEL-GAP-002 (FIN-C03): Minimum 24h between different-channel touches to the SAME contact (4h if engaged). Distinct from ACCT-COORD-001's 48h across-contact stagger.

Revised GOODHART-SAMPLE-001 (FIN-C04): Now a portfolio-level recalibration guard stacking on top of AILEARN-001's per-parameter minimums.

QUEUE-ARCH-001 (FIN-C05): Encodes GOVERNANCE-004's resolution — processing is immediate and 24/7; review and send are decoupled via the Ready-to-Fire Queue and KSA send windows.

AUTO-CEILING-001 (FIN-C06): No autonomy level ever removes human review from C-suite / P1 Domino messages. Constitutional carve-out aligned with MP-19.

Revised FAIL-HUMAN-004 (FIN-C07): Automation phase is a function of Trust level, not calendar time. Trust regression reverts the phase.

WHAT THE AUDIT CONFIRMED IS ALREADY COHERENT

No duplicate rule IDs across all 1,058 rules.

Single consistent scoring formula: Effective Opportunity and Decision Score (M2-EQUATION-001) used everywhere; no competing equation survived into the final set.

ICS action floor (40) consistent across all gating rules.

Constraint Hierarchy (Legal > Safety > Relationship > Timing > Revenue > Optimization) consistently referenced as the conflict-resolution spine, with CONFLICT_RESOLUTION entity and precedent library backing it.

Constitutional constraints (25 Meta-Principles, opt-out, anti-fabrication, vendor exclusivity, A0 hard-gates) consistently marked immutable across §13 and the entity model.

KSA cultural and timing rules (Ramadan, prayer windows, hierarchy, formality) are consistent and complementary across §16, §17, AIBOUNDARY-003, and the timing rules.