

BUILD ARTIFACT 3: BUILD ORDER + ACCEPTANCE CRITERIA

Dependency-ordered build sequence with pass/fail tests | Decimal Technologies

WHY THIS EXISTS

The 1,065 rules are organized by domain, not by what must be built first. This artifact sequences the build into 12 dependency-ordered stages and gives each stage a concrete acceptance test. An engineering lead can run this top to bottom. Each stage is only 'done' when its acceptance test passes. Stages 1-7 deliver the MVP; 8-10 the Production engine; 11-12 the Perfect engine.

BUILD SEQUENCE — MVP (Stages 1-7)

#	Stage	Depends On	Acceptance Test
1	Data layer: PostgreSQL schema for 12 core entities + Rule_Registry	1	12 tables exist; Rule_Registry seeded from Artifact 1; a signal row inserts and r
2	Signal ingestion (M1): poll, normalize, dedup, store	1	A polled news item stores once; a near-duplicate (freshness>0.8) is rejected; trace
3	Scoring engine (M2): Dynamic Score + ICS + modifier chain	2	BUDGET-005 worked example returns 46.8 exactly; decomposition matches Artifa
4	Tier assignment + state machine	3	Score 78 -> TIER_1 + ENRICHING state; score drop mid-sequence pauses per ST
5	Enrichment (M3): Apollo + LinkedIn, EPIS tagging	4	TIER_1 account gets 3+ contacts; each enrichment row carries source EPIS profile
6	AI chain (M4+M7): 7 agents + QC + Reasoning Passport	5	Generated msg carries full passport; QC fails a planted hallucination; max 3 regen
7	Outreach + engagement + handoff (M4/M8/M9)	6	Msg queues, waits for approval token + KSA window; reply classified; POSITIVE tr

BUILD SEQUENCE — PRODUCTION (Stages 8-10)

#	Stage	Depends On	Acceptance Test
8	Opportunity entity + dual-score + multi-channel	7	Two opportunities on one account route independently; LinkedIn+email respect 24
9	EPIS spine + Autonomy Ladder A0-A2 + Trust Progression	8	A2 promotion requires 100 decisions <5% override; trust regresses on a planted in
10	Learning loops + CRM bidirectional + \$5/\$9/\$94 modules	9	No weight recalibration before AILEARN-001 minimums AND 50/15 portfolio floor;

BUILD SEQUENCE — PERFECT (Stages 11-12)

#	Stage	Depends On	Acceptance Test
11	Cognition M13-M18 + Temporal Intelligence + Strategic Memory	10	Champion Score computes from 5 dims; Opportunity Death flags >70% prob; deca
12	Learning Governance + Antifragility + A3-A5 + Arabic	11	Kill switch halts all sends <5min; constitutional constraint blocks a learning propos

CRITICAL-PATH ACCEPTANCE TESTS (the ones that catch the dangerous bugs)

Test	Validates	Pass Condition
T-SCORE-1	Modifier chain correctness	Dynamic 90, ICS 80, EVALUATING, UNFUNDED = 46.8 (+/- 0.1)
T-SCORE-2	ICS gating	Account with ICS 35 cannot reach TIER_1 even with Dynamic 95 (GOV-ICS-001)
T-AUTO-1	Autonomy floor	No /v1/send succeeds without human_approval_token during MVP
T-AUTO-2	C-suite ceiling	C-suite message blocked from auto-send at every autonomy level (AUTO-CEILING-001)
T-TIME-1	P1 queue decoupling	P1 signal at 11pm Thu: processed by 1am, sent Sun 9am window (QUEUE-ARCH-001)
T-TIME-2	KSA blackout	Send attempt Friday or during Ramadan blackout returns 403
T-DEDUP-1	Signal dedup	Same signal within freshness window stored once, not twice
T-HALLUC-1	Hallucination catch	Planted fake funding claim flagged RED by fact-checker, blocked at QC
T-TRUST-1	Trust regression	One planted hallucination drops trust one level; automation reverts
T-STATE-1	No duplicate outreach	Reply mid-sequence pauses; next scheduled touch does not send
T-CONST-1	Constitutional block	Learning proposal to zero a scoring dimension auto-rejected (WEIGHT-FLOOR/CONST-WEIGHT)
T-PDPL-1	Compliance	Enrichment record without legal_basis_code flagged, excluded from use

HOW TO USE THIS WITH THE THREE RULE DOCUMENTS

For each stage, the engineer pulls the relevant rules from the FSD documents (Document 1 for stages 1-7, Document 2 for 8-10, Document 3 for 11-12), implements them, then runs the stage's acceptance test plus any critical-path tests that touch that stage. A stage is 'done' only when its test passes. Do not begin a stage until its dependencies are green. This converts 1,065 rules from a reference document into an executable plan.

Recommended team note: Stage 3 (scoring engine) is the highest-risk for silent divergence — two developers can implement it differently and both believe they followed the spec. T-SCORE-1 and the Artifact 1 lookup table exist specifically to prevent that. Treat Stage 3's tests as non-negotiable gates.